

PROF. DR.-ING. MARK SCHUTERA

FULL STACK HANDWERK

UNFINISHED LECTURE NOTES

Copyright © 2026 Prof. Dr.-Ing. Mark Schutera

PUBLISHED BY UNFINISHED LECTURE NOTES

Combobulated with the help of multiple large language model driven tools. Licensed under the Creative Commons Attribution-NonCommercial 4.0 International License ("CC BY-NC-SA 4.0"). You may not use this file for commercial purposes. If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original. You must obtain explicit permission from the author for uses beyond those permitted by this license. To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc-sa/4.0/>. Unless required by applicable law or agreed to in writing, distributed material is provided on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the license for details.

These notes are, by their very nature, unfinished, and they improve with every reader. If you spot an error, disagree with a framing, or want to add a source, a question, or a position, open a pull request at https://github.com/Quillstacks/LectureMaterial/tree/main/lecturenotes/notes_missingsemester. Every contribution is welcome.



*“SEINE WORT’ UND WERKE MERKT ICH UND DEN BRAUCH,
UND MIT GEISTESSTÄRKE TU ICH WUNDER AUCH.”*

DER ZAUBERLEHRLING

Contents

<i>Projektbericht</i>	9
<i>Development and Production</i>	11
<i>Version Everything: Git & GitHub</i>	23
<i>Dependencies & Structure</i>	39
<i>Integrate the Model</i>	57
<i>The Backend: APIs & Services</i>	75
<i>Databases & SQL</i>	93
<i>Reverse Proxy & Frontend</i>	109
<i>Continuous Integration Continuous Deployment</i>	125
<i>Ship on something that runs</i>	139
<i>Index</i>	141

Introduction

THIS IS AN ADAPTATION of the MIT Missing Semester course material. And oh boy, we need it here in Germany as well! In the Länd of the "Ingenieure", "Handwerk", and "Lehrlings" - we forgot the importance of practice and mastering our tooling skills.

Original course URL <https://missing.csail.mit.edu/>

AS A MECHANICAL ENGINEER BY TRAINING, I have always had this nagging feeling that something important was missing from my education. The more I transitioned into software engineering roles, the more this feeling intensified. Looking at your curriculum, there is no shortage of advanced concepts and theories to learn. This one fills the gap that many computer science programs overlook: practical proficiency with essential tools and workflows that every software engineer relies on. And truth be told, Professors rely on students having this basic proficiency already acquired, or to figure this stuff out on their own - which is fair in academia. For me, software engineering is also a practical discipline, and ought to be taught as such. I need you to know this stuff, so it is also fair to teach it.

THIS COURSE WILL TEACH YOU everything you need to know to become proficient with the command line, version control systems like Git, text editors, shell scripting, and other essential tools that are the backbone of modern software development - in the end you will be able to develop and deploy your software solutions. This will make you a better software engineer, and it will make your life easier latest when you need to go full ML Lifecycle.

Projektbericht

2026-06-09 · lively cherimoya Mufflon

REAL ARTISTS SHIP. - STEVE JOBS

The Why

A lecture that teaches tools and workflows must also assess tools and workflows. Written exams test recall; a project report tests whether you can actually build and deploy something. Enter the Projektbericht.

PICK ANY BLOCK FROM THE COURSE and build a meaningful extension onto it. The extension should add functionality, improve reliability, or strengthen security in a way that goes beyond what the lecture covered. The choice is yours, but each extension must be unique across the class: submit your proposal for approval on a first come, first served basis with me. If your proposal strikes home with me, we can also discuss small full stack projects, which live outside the scope of the lecture material.

SUBMIT A WRITTEN REPORT AS PDF and a link to your GitHub repository.

The report must cover:

- Which block you chose and what it covers in the course.
- What you built on top of it, which methods, tools and techstack you used, and why it adds value.
- Open questions and problems you encountered and how you resolved them.
- What works, what remains incomplete, and what would you do next.

THE PROJEKTBERICHT AS IN THE PRÜFUNGSORDNUNG. The official frame for the Projektbericht is set by the DHBW Studien- und Prüfungsordnung Wirtschaft, available [here](#):

“Die zu prüfende Person soll zeigen, dass sie in der Lage ist, Projekte oder Studien mit wissenschaftlicher oder praktischer Problemstellung selbstständig zu bearbeiten sowie deren Ergebnisse schriftlich zu dokumentieren. Der Projektbericht soll je zu prüfender Person in Modulen mit fünf bis sechs ECTS-LP zehn bis zwölf Seiten, (bei uns entsprechend 5 Seiten)

...

umfassen.”

HOW TO WRITE THE REPORT. A practical guide for writing reports and theses lives [here](#).



For deeper druidic knowledge on writing a thesis, consult [here](#).



EVALUATION considers three dimensions: *depth* of the extension beyond the lecture material, *understanding* demonstrated through explanation and a working GitHub repository, and *quality* of the writing.

LOOKING FOR INSPIRATION? Explore what tools modern teams actually use [here](#) or browse the ThoughtWorks Technology Radar [here](#). Pick something that excites you; the best reports come from genuine curiosity.

Development and Production

2026-06-12 · rosy tamarind Habicht


DIVIDE AND CONQUER.

The Why

Most software engineering curricula focus on algorithms, data structures, and the theory of computation. Writing a Python script that runs once on a single laptop is a skill; making that script available as a service, keeping it running, and recovering when it fails is a different one. This course addresses the second set of skills, the tools and workflows that sit between working code and the full stack of setting up and maintaining a running system.

ACROSS TEN LECTURE BLOCKS WE BUILD PIXELWISE, a hand-written digit recognition web application. The starting point is two empty virtual machines; the end point is a deployed service including a frontend, and a backend with a trained model, an API, and a database, all running on Linux servers and connected over a network, with version control, automated deployment, persistent storage, and monitoring. Each tool we introduce is motivated by a specific need that arises in the construction of a fullstack application: Pixelwise.

DEVELOPMENT AND PRODUCTION.



The figure shows the starting arrangement: two virtual machines that play distinct roles throughout the course. The dev VM is the

PIXELWISE is the narrative thread of the entire course. Users draw a digit on a pixelated 28×28 canvas in the browser, the backend runs an image classifier trained on MNIST, the model predicts class and confidence, and results are stored in a database.

THIS ENABLES YOU to deploy your own products and services in the future.

Figure 1: The two-machine setup. The dev VM is where code is written and tested; the prod VM is where the deployed application runs. Code travels from dev to prod through a deliberate deployment step over SSH.

development environment in which code is written, modified, and tested. The prod VM is the production environment in which the deployed application runs. Later blocks add the backend stack, the frontend, persistence, and automation between the two, but the separation between development and production stays intact from this block onward.

Hands On Experience

CONSIDER A SMALL WEB APPLICATION running on your laptop. A classmate two time zones away wants to try it; you send them the address. It works, until you close the lid to catch a train. That evening the Wi-Fi has changed, the address is different, and the process you started in the terminal is gone.

A SERVER is a machine whose job is to be available when no one is watching. It stays powered on, holds a stable address, and answers requests around the clock. Your laptop is where code is written and changed; the server is where code meets users. Keeping these two roles on separate machines is the first design decision of the course.

THIS FIRST BLOCK focuses on machines and environments rather than on code. Before introducing a programming language, a version control system, or a framework, we need a place which is always up, to run programs, isolated from experiments, and a place for development. By the end of the block you will have provisioned both virtual machines,

- set up a server and a development machine as virtual machines on your laptop,
- learned the subset of Linux needed to navigate a server from the command line,
- and configured remote access over SSH using key-based authentication.

COUNT THE HOURS your laptop is actually reachable at a stable address. Subtract the time it is closed, asleep, on battery, behind a café router, or re-booting after an update. What remains is the service level a laptop alone can offer.

TWO MACHINES, ONE MENTAL MODEL. dev at 192.168.56.10 stands in for your laptop. prod at 192.168.56.11 stands in for the always-on host. Later lectures fill in the bridge.

Virtual Machines

A VIRTUAL MACHINE IS A SIMULATED COMPUTER running inside your real computer. The software that creates and runs virtual machines is called a hypervisor. The hypervisor partitions the host's CPU, memory, and disk, and presents each virtual machine with what appears to be its own hardware. From the guest operating system's perspective, it is running on a physical machine.

HYPERVISORS are often categorised as type 1 or type 2. A type 1 hypervisor runs directly on the hardware, with the host operating system built into the hypervisor itself. Examples include VMWare ESXi and Microsoft Hyper-V in its server role; cloud providers use type 1 hypervisors to host the virtual machines you rent. A type 2 hypervisor runs as an application on top of a regular operating system. Oracle VirtualBox Manager is a type 2 hypervisor, which is why we can install it as a normal program on Windows, macOS, or Linux. The conceptual model is the same in both cases; the differences are in performance and deployment context.

ISOLATION is the primary reason to use a virtual machine. A misconfiguration, an uninstalled package, or a destructive command inside the guest does not affect the host operating system. Because the hypervisor abstracts away the hardware, the same VM runs identically on Windows, macOS, and Linux hosts, which is the basis of the reproducibility argument.

THERE ARE COSTS. A virtual machine executes a full operating system on top of the host, which carries a performance overhead that is typically small for I/O-bound work and larger for CPU-bound work. Each VM reserves some amount of RAM, disk space, and CPU time from the host, so the number of concurrent VMs is limited by available resources. Boot time is longer than launching an application but shorter than booting physical hardware. These costs are the price paid for isolation, reproducibility, and rollback.

HOST-ONLY NETWORKING puts dev and prod on a private subnet, typically 192.168.56.0/24, that only the host and the guests can reach. The guests get internet through a second, NAT-mode adapter, so they can install packages without being exposed to the outside world. We assign static IPs, 192.168.56.10 for dev and 192.168.56.11 for prod, so that every command in the course refers to the same addresses.

A SNAPSHOT captures the entire machine state at a point in time, including disk contents, memory, and running processes. Snapshots are inexpensive to create and well-suited to exploratory work; making one before a risky change and restoring it afterwards is a standard pattern.

OUTSIDE THE CLASSROOM, a production host is usually a rented or purchased server in a data centre, and the development machine is the laptop or workstation in front of you. Running both as VMs inside a single laptop is a teaching shortcut: it preserves the two-machine separation without asking you to rent hardware, and everything we learn here transfers directly when prod is later replaced by a cloud instance reached over the public internet.

Exercise: Provisioning Virtual Machines

INSTALL THE ORACLE VIRTUALBOX MANAGER for your host OS with default settings, and download two Ubuntu LTS ISOs (go for version 24.x.x if you find one): the Desktop edition for dev and the Server edition for prod. The asymmetry is deliberate, dev carries a graphical environment so you can run an editor and a browser locally, while prod stays headless because a production server has no business running a desktop. You can also decide to run dev on your machine directly, but this is not further supported in the exercises.

CREATE THE dev VM in the Oracle VirtualBox Manager. Click New and configure:

- Name dev, Type Linux, Version Ubuntu (64-bit)
- Memory at least 4096 MB, CPUs at least 2
- Hard disk at least 25 GB, VDI, dynamically allocated

CREATE THE prod VM with the same procedure but tighter resources, since it runs headless:

- Name prod, Type Linux, Version Ubuntu (64-bit)
- Memory at least 2048 MB, CPUs at least 2
- Hard disk at least 10 GB, VDI, dynamically allocated

ATTACH THE MATCHING ISO to each VM. VirtualBox offers an unattended installation checkbox during creation, don't select it, you want the interactive installer. Once a VM exists, select it in the manager and open Settings, Storage. Under Storage Devices you will see a controller, IDE or SATA, with an Empty optical drive beneath it. Click that Empty entry, then on the right under Attributes click the small disc icon next to Optical Drive, choose Choose a disk file, and pick the Desktop ISO for dev or the Server ISO for prod. Confirm with OK, start the VM, follow the installer, when in doubt go with the defaults and just forward-enter.

AFTER THE FIRST BOOT OF EACH VM (BOTH DEV AND PROD), INSTALL THE OPENSSSH SERVER TO ENABLE REMOTE ACCESS:

```
sudo apt update
sudo apt install openssh-server
```

This step is required for both Desktop and Server editions. Reboot after installation.

EXERCISES are for practice and reinforcing concepts. Work through them yourself first, break things, discuss, this is not a time trial. If your VM ends up in a state you cannot recover, restore the B1-fresh snapshot and start over. At the end of the session, take a snapshot named B1-complete, your safety net for the next block.

WHERE TO GET THEM. Virtual-Box: <https://www.oracle.com/virtualization/virtualbox/>. Ubuntu Desktop LTS ISO: <https://ubuntu.com/download/desktop>. Ubuntu Server LTS ISO: <https://ubuntu.com/download/server>. Pick the LTS release in both cases, not the interim one, so support extends across the semester. On an Ubuntu host you can use GNOME Boxes instead: <https://apps.gnome.org/Boxes/>. On macOS, UTM is the closest equivalent: <https://mac.getutm.app/>.

INSTALLER CHOICES that matter: use distinct credentials per VM, user/password on dev and produser/prodpassword on prod. The asymmetry is deliberate, you should feel that logging into prod is a different act from working on dev.

WHY THE SPREAD. Ubuntu Desktop with GNOME wants roughly 4 GB of RAM and 25 GB of disk to feel responsive. Ubuntu Server has no GUI and runs comfortably in a quarter of that. Sized this way the pair fits inside a 6 GB / 35 GB envelope on the host.

GERMAN KEYBOARD? The interactive installer often defaults to a US layout, which turns every hunt for a slash, tilde, or brace into a guessing game. Switch to German with:

```
sudo dpkg-reconfigure
keyboard-configuration
```

Pick Generic 105-key when prompted, then German, and accept the defaults for the rest.

CONFIGURE THE NETWORK ADAPTERS FOR EACH VM: Select the VM in VirtualBox Manager and open Settings, then go to the Network section. For Adapter 1, check "Enable Network Adapter" and set it to NAT (for internet access). For Adapter 2, select the Adapter 2 tab, check "Enable Network Adapter," set it to Host-only Adapter, and choose the host-only network you created.

ASSIGN STATIC IPs INSIDE EACH VM:

1. Boot the VM and log in.
2. Run `ip a` to list network interfaces. The host-only adapter is usually `enp0s8`, but confirm whether the name shows up in the output.
3. Edit the Netplan config with `sudo nano /etc/netplan/00-installer-config.yaml` Save and exit nano with (Ctrl+X, Y, Enter).
4. For dev, set the config as:

```
network:
  version: 2
  ethernets:
    enp0s8:
      addresses:
        - 192.168.56.10/24
```

5. Apply the config with `sudo netplan apply`
6. Check the new address with `ip a` again; you should see the static IP assigned to `enp0s8`.

VERIFY CONNECTIVITY: On dev, run `ping 192.168.56.11` to test if prod responds. If you get no reply, double-check the Netplan configuration (correct interface and address), ensure Adapter 2 is set to Host-only in both VMs, and confirm that both VMs are running and attached to the same host-only network.

```
user@dev:~$ ping 192.168.56.11
PING 192.168.56.11 (192.168.56.11) 56(84) bytes of data.
64 bytes from 192.168.56.11: icmp_seq=1 ttl=64 time=1.63 ms
```

MISSING ADAPTER 2? Power-off the VM, make sure you are in Expert-mode, go to Settings, Network, select Adapter 2, check "Enable Network Adapter," and set it to Host-only Adapter.

OPEN A TERMINAL (DESKTOP). On Ubuntu Desktop press Ctrl+Alt+T, or open Activities (top-left) and type "terminal"; press Enter. You can also right-click the desktop and select "Open Terminal" if the menu is present.

For prod, use 192.168.56.11/24 instead.

NETPLAN PERMISSION WARNING. Netplan will warn if a configuration file is writable by group/others or not owned by root. This is a security check to avoid loading untrusted configuration files; expected ownership is root:root and modes such as 644 for the YAML files.

TROUBLESHOOTING. If the interface name is not `enp0s8`, use the correct one from `ip a`.

Figure 2: ping

Linux on the Server

LINUX is a family of Unix-like operating systems built around a common kernel, first released by Linus Torvalds in 1991. A distribution combines the kernel with a package manager, system libraries, and user-space tools into a coherent product. Ubuntu, Debian, Fedora, and Arch are distributions that differ in release cadence and defaults but share most of what matters at the command line.

A SERVER IN THIS SECTION is a computer configured to run services and accept connections over a network, typically without a graphical session. You interact with it through a shell, a program that reads commands and prints their output; the shell we use is Bash, the default on Ubuntu.

FIVE COMMANDS are sufficient for the first session. `ls` lists a directory, with `-la` showing permissions, ownership, and hidden files. `cd` changes directory; `cd ..` moves up one level and `cd ~` returns home. `pwd` prints the current path, `mkdir` creates a directory, and `cat` prints a file to the terminal.

LINUX IS A MULTI-USER OPERATING SYSTEM. Every file belongs to a user and a group and carries three permission sets, for owner, group, and everyone else. Each set independently grants read, write, or execute access. `whoami` reports the current user, and `ls -la` shows ownership and permissions:

```
-rw-r--r-- 1 user user 1234 Jan 15 10:00 config.txt
drwxr-xr-x 2 user user 4096 Jan 15 10:00 mydir/
```

TWO COMMANDS modify these. `chmod` changes permissions and `chown` changes ownership:

```
chmod 600 ~/.ssh/id_ed25519          # owner-only read and write
sudo chown user:user file.txt        # assign user and group
```

The principle of least privilege applies throughout the course: grant each user, process, or service only the access it needs, so that narrow defaults limit the blast radius of a compromise or a mistake.

UBUNTU at <https://ubuntu.com/> is the distribution we use on both VMs. Long-term support releases receive five years of security updates, which is why we pick the LTS ISO rather than the latest release.




Figure 3: A running Ubuntu guest inside the Oracle VirtualBox Manager. The guest believes it owns a physical machine; the hypervisor is what makes that illusion hold.

THE UNIX PHILOSOPHY shapes the design of these tools. Each program does one thing well, programs pass streams of text to each other, and configuration lives in plain files. The result is a small vocabulary of composable commands that combine into pipelines no single command was designed for.

READING THE MODE STRING from left to right, the first character tells you whether it is a file or directory, then three groups of `rwx` for owner, group, others. A dash means the permission is not granted. The leading `d` marks a directory.

Exercise: Navigating Linux

THIS EXERCISE turns the five commands and the permission model into something your hands remember. You walk the directories common to every Linux system, identify yourself to the machine, and read and change a file's mode string. By the end, the output of `ls -la` should read as a description you can speak aloud rather than a blob of symbols.

EXPLORE THE FILESYSTEM by visiting `/home`, `/etc`, and `/opt` in turn with `cd` and `ls -la`, and read two files with `cat`:

```
cat /etc/hostname
cat /etc/os-release
```

CREATE A WORKSPACE and check that you are where you think you are:

```
mkdir ~/workshop && cd ~/workshop
pwd
touch notes.txt
rm notes.txt
```

`pwd` is a reassurance command. When something does not behave as you expect, asking the shell where it thinks it is usually explains the problem faster than re-reading the error.

UNDERSTAND WHO YOU ARE and look at the permissions on your new file:

```
whoami
id
ls -la notes.txt
```

Decode the first column of `ls -la` against the margin note on mode strings from the theory above. Name aloud what each of the ten characters means; if you can do that without looking, permissions have stopped being mysterious.

CHANGE THE MODE AND WATCH IT:

```
chmod 600 notes.txt
ls -la notes.txt
```

The next exercise will require `600` on your SSH private key, or OpenSSH refuses to use it; you have now seen what `600` looks like and why ownership matters.

TEN-FINGER TYPING is part of the toolchain once the terminal is your main tool. A server has no menus, so your throughput is bounded by how quickly you can speak to the shell without looking at the keys; every hunt for a slash, tilde, or brace is attention spent on the keyboard instead of the problem. Ten minutes a day at <https://www.keybr.com/>, <https://typing.io/>, or <https://monkeytype.com/> is enough to remove that tax over a few weeks.

THE FILESYSTEM AS MAP. `/home` is where users live, `/etc` holds configuration, and `/opt` is where PixelWise will be installed in a later block. These locations reflect decades of convention you will meet on every Linux machine; knowing them lets you locate things by instinct.

A THOUGHT EXPERIMENT ON LEAST PRIVILEGE. Your user account can read `/etc/os-release` but cannot edit `/etc/ssh/sshd_config` without `sudo`. Why does that separation exist, and what would change if the web service we deploy in a later block ran as `root`? Name one file an attacker who compromised such a service could reach that a non-root service could not; the mental habit matters more than the exact file.

SSH, the Only Door

SSH, THE SECURE SHELL, is a protocol for operating a remote machine from a terminal over an encrypted channel. In the setup used in this course, SSH is the only means of access to the prod VM: there is no graphical login and no console attached. Misconfiguring SSH can lock a user out of the machine, which is why key material and the daemon's configuration are treated carefully in the exercises.

THE BASIC COMMAND takes the form `ssh <username>@<host>`; for example, `ssh produser@192.168.56.11` opens an encrypted connection to the server as the `produser` account. All input and output travel through this channel until the session is closed with `exit` or `Ctrl+D`.

PASSWORD-BASED AUTHENTICATION is vulnerable to guessing, brute-force attacks, and phishing. SSH supports an alternative in which the user presents a cryptographic key pair. The private key is kept on the client and never transmitted; the public key is placed on the server. During authentication the server issues a challenge that only the holder of the private key can answer, proving identity without the key itself ever leaving the client.

GENERATE AN ED25519 KEY PAIR ON dev:

```
ssh-keygen -t ed25519 -C "you@example.com"
```

The command writes the private key `id_ed25519` and the public key `id_ed25519.pub` to `~/.ssh/`. The private key must have permissions `600`, owner-only read and write, or SSH refuses to use it.

COPY THE PUBLIC KEY TO THE SERVER:

```
ssh-copy-id produser@192.168.56.11
```

This appends the key to `~/.ssh/authorized_keys`. Subsequent connections authenticate using the key pair and do not prompt for a password.

ONCE KEY AUTHENTICATION IS VERIFIED, password authentication can be disabled on the server. Edit `/etc/ssh/sshd_config` to set `PasswordAuthentication no`, then restart the daemon:

```
sudo systemctl restart ssh
```

OPENSSH at <https://www.openssh.com/manual.html> is the implementation shipped with Ubuntu and most other Linux distributions. The client on `dev` and the daemon on `prod` are both part of this project.

PuTTY at <https://www.putty.org/> is a legacy SSH client for Windows, from the years before OpenSSH shipped with the operating system. Windows 10 and later include `ssh.exe` natively, so PuTTY is rarely necessary; you may still encounter it in older tutorials and on locked-down corporate machines.

VS CODE REMOTE-SSH at <https://code.visualstudio.com/docs/remote/ssh> runs the editor locally while the language server, terminal, and file system live on the remote host. It reads `~/.ssh/config` and reuses the same keys as the `ssh` command, so any host you can reach with `ssh user@host` is one click away inside the editor.

ED25519 keys use elliptic-curve cryptography and produce shorter keys than the older RSA algorithm while providing comparable or stronger security. They are the default choice on modern SSH clients.

MANAGING PRIVATE KEYS. Private keys live in `~/.ssh/` on the machine that initiates connections, protected by filesystem permissions and, when generated with a passphrase, by symmetric encryption. `ssh-agent` holds the decrypted key in memory for the session so the passphrase is typed once per login. Password managers such as KeePass at <https://keepass.info/> or 1Password can store the passphrase and, with their SSH-agent integrations, supply the key to `ssh` without writing it to disk in plaintext. Keys are never synced to cloud storage in the clear, and a separate key pair per machine keeps the blast radius of a compromised laptop local.

After this change, password-based login attempts are rejected before credentials are even evaluated, which eliminates brute-force and credential-stuffing as viable attacks against the machine. The exercise introduces nano as a terminal editor and `systemctl` as the interface to the `systemd` service manager, both of which return throughout the course.

Exercise: SSH and Key-Based Access

THIS EXERCISE walks through three states of remote access in order: password login, key-based login, and password login disabled. The point is not the commands but to recognise which state you are in at any moment and why each is stronger than the last.

FIRST, LOG IN WITH A PASSWORD. From dev, open a session to prod:

```
ssh produser@192.168.56.11
```

Type your password, look around briefly, and log out with `exit`. This is the baseline. Every later step is interpretable as an upgrade over this state.

GENERATE A KEY PAIR on dev and install it on prod:

```
ssh-keygen -t ed25519 -C "you@example.com"
ssh-copy-id produser@192.168.56.11
```

Reconnect with `ssh produser@192.168.56.11` and notice the password prompt is gone. Log in and inspect what `ssh-copy-id` wrote for you:

```
cat ~/.ssh/authorized_keys
```

DISABLE PASSWORD AUTHENTICATION. Edit `/etc/ssh/sshd_config` with `sudo nano`, set `PasswordAuthentication no`, and restart the daemon:

```
sudo systemctl restart ssh
sudo systemctl status ssh
```

The status output should show `active (running)`. If it does not, keep your current session open and fix the configuration before logging out; a locked-out VM is rescued by restoring the snapshot, not by wishful thinking. From a machine without the key, the daemon now rejects the attempt without even asking for a password.

TAKE A SNAPSHOT of both VMs once everything works, naming them B1; this is your known good state for the rest of the course. Then deliberately break something inside a VM, watch the fallout, and restore the snapshot. You now have physical evidence that the safety net holds.

WHAT JUST HAPPENED.

`ssh-copy-id` appended the contents of `id_ed25519.pub` on dev to `~/.ssh/authorized_keys` on prod. At the next connection the daemon issued a challenge that only the private key on dev could answer; no secret material ever left the client. Seeing your public key in a readable file anchors the whole exchange to something concrete.

NUKE THE MACHINE. From prod, run `sudo rm -rf /` to delete all files on the machine. Restore from the snapshot.

RUN PROD WITHOUT GUI. In future run prod headless without GUI from the VirtualBox Manager. You can still log in with SSH and run commands, but there is no desktop to interact with. This is how a real production server works.

A NOTE ON WHAT COMES NEXT. The commands you ran in this session are reproducible but not yet reproduced. You ran them in a terminal, they worked, and they are gone. In the next block we place them under version control and turn the sequence into a script that any future VM can run end to end, without you remembering anything.

Self-Reflection and Recap

SELF-REFLECTION Questions which can guide your thoughts during and after the exercises:

- Why do we use two separate machines instead of running everything on one?
- Why is disabling password authentication a stronger protection than picking a longer password?
- What does the principle of least privilege mean in the context of file ownership, and where did it show up in today's workshop?
- If your laptop died tomorrow, how much of today's setup could you reproduce, and what would you lose?
- Which of the five Linux commands do you still have to look up, and which one surprised you?
- How does the two-machine model map to a real cloud deployment with staging and production? How is our development machine different from staging?

RECAP of Key Concepts:

- Virtual machines isolate experiments from the host and make snapshots a first-class safety net.
- The two-machine architecture separates development from production at the hardware level.
- Files on Linux carry owners, groups, and three sets of permissions, and `chmod` and `chown` manage both.
- SSH is the only door into the server, and key-based authentication replaces passwords with cryptography that cannot be guessed.

MILESTONE. So far, we own the machine. The next step is to own the way we build software on it. In the next chapter we introduce Git and GitHub: not just a version history, but the collaboration layer that lets a team branch, review, and merge work in parallel. It is the tool that turns two developers editing the same file from a disaster into a workflow. ,

WITHOUT SHIPPING ANYTHING YET, we already have two machines that talk to each other, a secure door, and a known good state to fall back on.

TEASER. You and a classmate both want to improve PixelWise at the same time. How do you work on the same codebase without overwriting each other's changes, review what the other person did, and merge the results safely?

Version Everything: Git & GitHub

2026-06-12 · plucky mango Eichhörnchen

GIT A LIFE.

The Why

Block 1 gave us two machines and a way to reach one from the other. We can write code on dev and, eventually, deploy it to prod. But nothing yet records what changed between deployments, who changed it, or how to undo a mistake. The moment two people touch the same file, the same function, or one experiment needs to coexist with another, the workflow breaks down entirely.

VERSION CONTROL solves an entire family of problems at once. It records every change with a message, a timestamp, and an author, so the full history of a project is always available. It lets you return to any previous state in seconds, turning risky experiments into cheap, reversible ones.

PARALLEL DEVELOPMENT becomes natural. Multiple people work on the same codebase without overwriting each other's changes. A new feature, a bug fix, and a refactor can all proceed simultaneously on isolated branches and be integrated later with explicit, reviewable merge steps.

DOWNTIME SHRINKS as a direct consequence. Developers merge through a guided review step that catches conflicts and errors before they reach main, deploying only code that has already been integrated and verified. When something still slips through, you can roll back to the last known-good state in a single command instead of debugging under pressure while users wait.

COMBINED WITH A HOSTING PLATFORM like GitHub, version control becomes the backbone of code review, continuous integration,

PIXELWISE needs to live on two machines. Right now there is no reliable way to move code from dev to prod and keep both in sync. Git gives us exactly that: a shared history that either machine can pull from, so deployment becomes a deliberate, testable, repeatable step instead of manual file copying.

GitHub: <https://github.com>

issue tracking, and project management, built around the software development process.

Hands On Experience

CONSIDER THIS SCENARIO: you and a classmate are both improving the PixelWise classifier. You rewrite the data loader; your classmate changes the model architecture. You each work on your own laptop for a few hours, then try to combine the results. Without a system for tracking changes, combining means copying files back and forth, comparing line by line, and hoping nothing was lost. Worse still, if both of you would have worked on the same machine, overwriting each others files as you go.

GIT IS THAT SYSTEM. It tracks exactly what changed in every file, when, and by whom. It lets you branch off to try an idea, or build a feature without disturbing stable code, merge branches back together with full visibility into conflicts, and review each other's contributions before they reach the shared codebase. The design principle is simple: never lose work, and never wonder what happened.

THIS SECOND BLOCK introduces Git and GitHub as the version control layer for PixelWise. By the end you will have

- understood the Git data model and why it makes history tamper-proof,
- practised the staging workflow of editing, staging, and committing changes,
- created branches, merged them, and resolved conflicts,
- pushed code to a remote repository on GitHub and collaborated through pull requests,
- and configured SSH-based authentication so Git operations are seamless from both VMs.

EVERY TEAM that ships software uses version control. It is not an optional workflow preference; it is the infrastructure that makes collaborative, iterative development possible at any scale. As projects grows, interface contracts, clear agreements on inputs, outputs, and responsibilities between components, will become increasingly important. Version control lays the foundation by making every change visible and reviewable, so contracts can be enforced rather than assumed.

Optional Exercise: Catch Up to Block 1's End State

This block builds on the state at the end of Block 1: Two virtual machines, dev and prod, with key-based SSH from dev to prod and password authentication disabled. If you joined late or your VMs are in an unknown state, restore the snapshots rather than redoing every prior exercise.

RESTORE THE SNAPSHOT. Open VirtualBox Manager, select the dev VM, and restore the B1-complete snapshot. Then select the prod VM and restore its B1-complete snapshot as well. Both machines now sit at the exact state in which Block 1 ended.

VERIFY THE NETWORK AND SSH. From a terminal on dev, confirm that prod is reachable and that the key-based login still works:

```
ping -c 3 192.168.56.11
ssh produser@192.168.56.11
```

The ping should succeed and the SSH login should complete without a password prompt. If either fails, replay Block 1's networking and SSH exercises before continuing.

You are now at the state where Block 1 ended, with both VMs aligned, ready to start Block 2.

SNAPSHOT MAP. B1-complete is the snapshot taken at the end of Block 1 on both VMs. From Block 2 onwards, repository state is also versioned with Git tags, the first being v0.1 at the end of this block.

WITHOUT A SNAPSHOT. If you never took B1-complete, replay the Block 1 exercises end to end: Provisioning the two VMs, configuring host-only networking with static IPs, generating an Ed25519 key on dev, copying it to prod with ssh-copy-id, and disabling password authentication on prod's sshd.

The Git Data Model

GIT IS NOT A BACKUP TOOL. It is a content-addressable filesystem with a version history layered on top. Understanding the data model makes every command intuitive.

Blob A file's contents, stored by its SHA-1 hash. The name is irrelevant; only the content matters.

Tree A directory listing: maps filenames to blobs (files) or other trees (subdirectories).

Commit A snapshot of the entire project at a point in time. Contains: a pointer to the root tree, the author, a timestamp, a message, and a pointer to the parent commit(s).

Branch A lightweight, movable pointer to a commit. `main` is a branch. Creating a branch is instantaneous; it just creates a new pointer.

HEAD A pointer to the branch you are currently on. When you commit, `HEAD` moves forward.

EVERY COMMIT is identified by a SHA-1 hash, a 40-character hexadecimal string like `a3f2b7c...`. This hash is computed from the commit's contents. If anything changes, even a single character in a single file, the hash changes. This makes Git history tamper-proof.

DATA HYGIENE follows directly from the data model. Because Git stores every version of every blob permanently, committing a large dataset means carrying it in the repository history forever, even if you delete the file later. A 500 MB training set committed once, modified twice, and then removed still occupies 1.5 GB of history. Repositories should contain code, configuration, and small metadata files. Large or frequently changing binary assets, datasets, and model weights do not belong in Git.

The `.gitignore` file, covered in its own subsection below, is the practical mechanism for keeping these artefacts out of the repository.

Git was created by Linus Torvalds in 2005 to manage the Linux kernel source code. It is now the de facto standard for version control in software development.

Git: <https://git-scm.com/>

Pro Git book (free): <https://git-scm.com/book/en/v2>

DEDICATED TOOLS fill the gap. Git LFS replaces large files with lightweight pointers. Hugging Face Hub hosts datasets and model weights with built-in Git LFS. DVC tracks data pipelines and links datasets to experiments. Weights & Biases versions datasets alongside experiment logs.

Git LFS: <https://git-lfs.com>

Hugging Face Hub: <https://huggingface.co>

DVC: <https://dvc.org>

W&B: <https://wandb.ai>

Exercise: Setting Up Git and the Repository

CREATE A GITHUB ACCOUNT if you do not have one. Then open a terminal on the dev VM.

INSTALL GIT. Ubuntu does not ship with Git by default, so install it through the package manager. It does not matter which directory you run this in, apt installs system-wide:

```
sudo apt update
sudo apt install -y git
git --version          # confirm the install succeeded
```

CONFIGURE YOUR GIT IDENTITY:

```
git config --global user.name "Your Name"
git config --global user.email "your@email.com"
```

GENERATE AN SSH KEY. SSH keys authenticate you to GitHub from the command line, replacing the password prompt on every push and pull. Before generating, check whether a key already exists so you do not overwrite one you rely on for another server:

```
ls ~/.ssh/
```

The listing includes `id_ed25519`, the key you generated in Block 1 for the prod VM. Leave it alone and add a second key for GitHub under a distinct filename:

```
ssh-keygen -t ed25519 -C "your@email.com" \
-f ~/.ssh/id_ed25519_github
```

Passing `-f` writes the key straight to the named path and skips the interactive prompt where one stray Enter would have overwritten the production key. Set a passphrase when prompted; it encrypts the private key on disk so a stolen file is not immediately usable.

TELL SSH WHICH KEY TO USE FOR GITHUB. A second key in `~/.ssh/` is invisible to SSH unless you name it. Pin `id_ed25519_github` to the `github.com` host:

```
cat >> ~/.ssh/config <<'EOF'
Host github.com
    HostName github.com
    User git
    IdentityFile ~/.ssh/id_ed25519_github
```

AT THE END OF THE SESSIONS, take a snapshot named `B2-complete`, your safety net for the next block.

OPENING A TERMINAL ON UBUNTU. The fastest way is the keyboard shortcut `Ctrl+Alt+T`, which launches the default terminal emulator. Alternatives: press the Super key (the Windows key) and type "terminal", or right-click on the desktop and choose "Open Terminal". Inside a terminal you can open another tab with `Ctrl+Shift+T` and a fresh window with `Ctrl+Alt+T` again.

`sudo` runs the command with administrator privileges, which apt needs to write to system directories. You will be asked for your password the first time. `apt update` refreshes the package index so you get the current version; `apt install -y` installs without prompting for confirmation.

RINGS A BELL? You generated an `id_ed25519` key in Block 1 to log in to the prod VM without a password. That same key is sitting in `~/.ssh/` now, which is exactly why `ls` matters before the next step: accepting the default path would overwrite the production key and lock you out of the server.

IF `~/` MISBEHAVES in your shell, pass the absolute path instead: `/home/<user>/.ssh/id_ed25519_github`.

CHECK THE PERMISSIONS. `ssh-keygen` writes the private key with mode `600` by default, which is what SSH wants. If anything looks off, `ls -l ~/.ssh/` confirms `-rw-----` for `id_ed25519_github`; `chmod 600` restores it if it is not.

WHY IdentitiesOnly? Without it SSH offers every key it can find, starting with `id_ed25519`. GitHub counts each as a failed attempt and slams the door after a handful, even if the correct key sits later in the queue.

```
IdentitiesOnly yes
EOF
chmod 600 ~/.ssh/config
```

ADD THE PUBLIC KEY TO GITHUB. Print it to the terminal and select the entire output, beginning with `ssh-ed25519` and ending with the comment:

```
cat ~/.ssh/id_ed25519_github.pub
```

In GitHub, navigate to Settings, then SSH and GPG keys, then New SSH key. Give it a recognisable title like `pixelwise-dev-vm` so you know which machine it belongs to, paste the public key into the Key field, and save.

TEST THE CONNECTION:

```
ssh -T git@github.com
```

The first time, SSH asks whether to trust GitHub's host fingerprint. Type `yes`. If the key is set up correctly, GitHub responds with `Hi <username>!` You've successfully authenticated, but GitHub does not provide shell access. The message about shell access sounds like an error but is the intended response: GitHub only allows Git operations over SSH, not interactive shells.

INITIALISE THE PIXELWISE REPO on the dev VM. Pick a stable working directory, your home folder is the natural place, so the project does not get lost in `/tmp` or buried under Downloads:

```
cd ~ # start from your home directory
mkdir pixelwise && cd pixelwise
pwd # confirm: /home/<user>/pixelwise
git init
```

Write a `README.md` with a one-line project description.

WHAT DID JUST HAPPENED?

`ssh-keygen` created two files in `~/.ssh/`. `id_ed25519_github` is the private key: never share it, never commit it, never copy it to another machine without good reason. `id_ed25519_github.pub` is the public key, safe to share, this is the file you paste into GitHub. The `-t ed25519` flag picks a modern signature algorithm; the `-C` comment labels the key so you can recognise it later when you have several.

`cd ~` jumps to your home directory from anywhere. `pwd` prints the current path so you can confirm where you are before running `git init`. Avoid initialising a repo inside `/tmp`, on a mounted USB drive, or inside another existing Git repository.

The Staging Workflow

GIT HAS THREE AREAS:

- *Working directory* The files on disk. What you edit.
- *Staging area* What you have marked for the next commit. A deliberate selection.
- *Repository* The committed history. Permanent, immutable snapshots.

THE CORE COMMANDS:

```
git status          # what changed? what is staged?
git add file.py     # stage a file for the next commit
git add .           # stage everything (use with care)
git commit -m "message" # commit the staged changes
git diff            # unstaged changes vs last commit
git diff --staged   # staged changes vs last commit
git log             # show commit history
git log --oneline    # compact view: one line per commit
```

The .gitignore File

NOT EVERYTHING BELONGS IN VERSION CONTROL. The .gitignore file tells Git which files and directories to exclude from tracking:

```
# Python
__pycache__/*
*.pyc
.venv/

# Secrets
.env

# Data and models (too large, reproducible)
data/
models/*.pkl

# Editor files
.vscode/
*.swp
```

THE NAPKIN CYCLE as easy as it gets:

```
git pull
git add <files>
git commit -m "..."
```

git push
Pull the latest changes, stage your work, commit it, push it. Every other command is a variation on this loop.

Write commit messages in the imperative mood: "Add classifier module" not "Added classifier module." The message completes the sentence "This commit will ...". Keep the first line under 72 characters. For a gallery of how not to do it, visit <https://whatthecommit.com>.

WHEN TO COMMIT? Commit when you have completed one logical change: a bug fix, a new function, a config update. If you struggle to write a short commit message, you probably changed too many things at once. If you go hours without committing, you are accumulating risk. Small, frequent commits make history useful; large, rare commits make it a wall of text.

```
(\ /)
(0.0)
(> <) Bunny approves these changes.
```

THIS INCLUDES ALL CREDENTIALS: API keys, database passwords, webhook secrets, tokens for AI services. These live on the server in environment variables or protected configuration files, never in the repository. A leaked production database password or API key can be exploited within minutes by automated scanners.

Git history is permanent. A secret committed once is leaked forever, even if you delete the file in a later commit. The old commit still contains it. `.gitignore` is your first line of defence.

Exercise: Staging, Committing, and .gitignore

WRITE A .gitignore. Create a .gitignore file that excludes Python caches, virtual environments, secret files, dataset directories, model checkpoints, and editor files. Verify it works: create a file called scratch.pyc, run `git status`, and confirm Git does not list it. Then remove the test file.

MAKE YOUR FIRST COMMIT. Stage and commit `README.md` and `.gitignore`:

```
git add README.md .gitignore
git commit -m "Add project README and .gitignore"
```

Run `git log` to see your first commit. Run `git status` to confirm the working directory is clean.

THE HABIT TO BUILD: after every commit, run `git status` and `git log --oneline` to confirm what just happened. It takes two seconds and prevents surprises.

Branching and Merging

BRANCHES LET YOU WORK ON IDEAS IN ISOLATION. The main branch is the stable baseline. Feature branches diverge, evolve, and merge back when ready.

```

git branch                # list branches
git branch feature-x      # create a new branch
git checkout feature-x    # switch to it
git checkout -b feature-y # create and switch in one step
git branch -d feature-x   # delete after merge

```

Delete branches after merging. Stale branches accumulate fast and make `git branch` unreadable. If the branch is merged, it is safe to delete; the history lives on in main.

WHEN THE FEATURE IS READY, merge it back and clean up:

```

git checkout main
git merge feature-x
git branch -d feature-x

```

Merge Conflicts

A MERGE CONFLICT occurs when two branches changed the same line in the same file. Git cannot decide which version to keep, so you must resolve it manually.

GIT MARKS THE CONFLICT in the file. The block between «««< HEAD and ===== is what the current branch has; HEAD always means “the commit you are on right now.” The block between ===== and »»»> feature-x is what the incoming branch has.

Edit the file to keep the correct version and remove the markers. Then stage the resolved file with `git add file.py` and commit with `git commit -m "Resolve merge conflict in file.py"`.

```

<<<<<< HEAD
result = model.predict(x)
=====
result = classifier.classify(x)
>>>>>> feature-x

```


Exercise: Branching, Merging, and Conflicts

CREATE A FEATURE BRANCH and make a change to README.md:

```
git checkout -b feature-readme
nano README.md          # add a project description line
git add README.md
git commit -m "Add project description to README"
```

SWITCH BACK TO main and edit the same line differently:

```
git checkout main
nano README.md          # edit the same line, but differently
git add README.md
git commit -m "Add alternative description to README"
```

MERGE AND RESOLVE THE CONFLICT. Now merge the feature branch and observe the conflict:

```
git merge feature-readme
```

Git reports a conflict and pauses the merge. Open README.md with nano README.md. You will see a block that looks like this:

```
<<<<<<< HEAD
... your text on main
=====
... text from feature-readme
>>>>>>> feature-readme
```

HEAD marks the version sitting on the branch you are currently on, in this case main. The lines after ===== are what feature-readme is bringing in. Choosing the correct version means deciding what the file should look like once both changes are reconciled. You have three options: keep only the HEAD side, keep only the incoming side, or write a merged version that combines parts of both. Then delete all three marker lines, <<<<< HEAD, =====, and >>>>>> feature-readme, so the file contains only the resolved content. Save and exit nano with Ctrl+X, Y, Enter, then complete the merge:

```
git add README.md
git commit -m "Resolve merge conflict in README"
```

Verify with git log --oneline --graph that both branches are now integrated.

EDITING FILES IN THE TERMINAL. nano README.md opens the file in a simple terminal editor. The bar at the bottom of the screen lists the shortcuts, where ^ stands for the Ctrl key: Ctrl+O writes the file to disk, Ctrl+X exits, and Ctrl+G shows full help. Vim and VS Code over Remote SSH work just as well if you prefer them.

main OR master? Older Git versions called the default branch master. Since 2020, the convention has shifted to main, and GitHub creates new repositories with main by default. If git checkout main reports that the branch does not exist, run git branch to see what your repo actually calls it. To bring an existing repository in line with the new convention, rename in place with git branch -m master main, then set the default for future git init runs with git config --global init.defaultBranch main.

git status DURING A MERGE. While the merge is paused, git status reports "You have unmerged paths" and lists every file with conflict markers still in it. Run it any time to remind yourself which files still need to be resolved before you can finish the merge.

Tags

A TAG MARKS A POINT IN HISTORY THAT MATTERS: a release, a milestone, a known-good state.

```
git tag v0.1                # lightweight tag
git tag -a v0.1 -m "initial setup" # annotated tag (preferred)
git push --tags             # push tags to remote
```

SEMANTIC VERSIONING gives tag names a shared meaning. A version number follows the pattern MAJOR.MINOR.PATCH. Increment PATCH for bug fixes that do not change behaviour, MINOR for new features that remain backwards-compatible, and MAJOR when you introduce breaking changes. A jump from v1.2.3 to v1.3.0 tells users that something was added but nothing they depend on was removed. A jump to v2.0.0 warns them to check for incompatibilities.

Lightweight tags are just a name pointing to a commit. Annotated tags store the tagger, date, and message; use them for releases. We tag v0.1 today; the tag becomes meaningful when you look back at it in Block 8.

tig is a terminal-based Git browser that makes it easy to navigate commits, branches, and tags on a headless Ubuntu machine. Install with `sudo apt install tig`.

<https://jonas.github.io/tig/>

Semantic Versioning: <https://semver.org>

Exercise: Tags and History Browsing

TAG THE RELEASE. Mark this point in history:

```
git tag -a v0.1 -m "initial server setup"
```

Verify the tag exists with `git tag -l` and inspect it with `git show v0.1`.

INSTALL tig, a terminal-based Git browser:

```
sudo apt install -y tig
tig
```

Navigate the commit history and find your tag. Press q to quit.

WHY NO `git push` YET? Tags live locally until you push them. We have not configured a remote yet, that happens in the next exercise. Once the remote exists, `git push --tags` ships every local tag to GitHub, where it appears under the repository's "Releases" tab.

VS CODE can also browse Git history visually. Open the PixelWise folder via `code .` or Remote SSH, then use the Source Control panel and the GitLens extension to explore commits, branches, and tags with a graphical interface.

Remotes and GitHub

```
git clone <url>           # copy a remote repo locally
# fork: copy repo to your GitHub account (GitHub UI, not a git command)
git remote add origin git@github.com:user/pixelwise.git
```

A REMOTE is a copy of your repository on another machine, usually GitHub.

CLONE AND FORK are two ways to start from an existing repository.

A fork is a GitHub concept, not a Git command. It creates a full copy of someone else's repository under your own GitHub account. You clone your fork locally, work on it freely, and propose changes back to the original via a pull request. Forking is the standard workflow for contributing to open-source projects where you do not have write access to the original repository.

`git clone` copies a remote repository to your local machine, keeping the original as the `origin` remote. You can pull updates and, if you have permission, push changes back.

GitHub is not Git. Git is the version control system; GitHub is a hosting platform that adds pull requests, issues, CI/CD, and collaboration features on top of Git. Alternatives include GitLab and Bitbucket.

GitHub: <https://github.com/>
 GitHub Docs: <https://docs.github.com/>

Pull Requests

PUSH sends your local commits to the remote. Until you push, your work exists only on your machine.

FETCH does the opposite: it downloads new commits from the remote but leaves your working directory untouched, so you can inspect what changed before merging.

PULL combines fetch and merge in one step, updating your local branch to include the remote's latest commits.

```
git push -u origin main    # push and set upstream
git pull                  # fetch + merge from remote
git fetch                  # fetch without merging
```

THE `-u` FLAG links your local branch to the remote branch. After running `git push -u origin main` once, Git remembers the connection and future `git push` and `git pull` work without specifying the remote or branch.

A PULL REQUEST is a proposal to merge a branch into another branch, typically a feature branch into `main`. It is the standard mech-

anism for code review: your changes are visible, reviewable, and discussable before they become part of the main codebase.

PULL REQUESTS AND CI. Because a PR represents a well-defined set of changes that has not yet reached `main`, it is the ideal moment to run a test suite, a linter, or any other check automatically. If the tests fail, the merge is blocked and `main` stays healthy. We will wire up this automation in a later block; for now, recognise that the PR is where human review and machine verification meet.

Exercise: Remotes, Push, and Clone

CONNECT TO GITHUB. Create a repository on GitHub, add it as a remote, push the commits, then push the tag you created in the previous exercise:

```
git remote add origin git@github.com:user/pixelwise.git
git push -u origin main
git push --tags
```

Verify on GitHub that your commits, .gitignore, and the v0.1 tag under the “Releases” tab are all visible.

CLONE TO THE SERVER. SSH into the server and clone PixelWise to its permanent home:

```
sudo mkdir -p /opt/pixelwise
sudo chown produser:produser /opt/pixelwise
git clone https://github.com/user/pixelwise.git /opt/pixelwise
```

PixelWise now lives on both machines. Run `git log --oneline` on the server to confirm the full history arrived.

REPLACE user WITH YOUR GITHUB USERNAME. The user placeholder in the remote URL is just a stand-in. Substitute your own GitHub handle, the one shown in your profile URL at `github.com/<your-handle>`, so the URL points at the repository under your account. The same goes for the `git clone` URL on the server below. Pushing to a user/pixelwise that you do not own will fail.

SSH INTO prod AGAIN. From the dev VM, open a session to prod the same way you did in Block 1:

```
ssh produser@192.168.56.11
```

The key-based login from Block 1 is still in place, no password should be required. Once you are on prod, the prompt changes to `produser@prod`, and the commands below run on the server. Type `exit` when you are done to return to dev.

WHY HTTPS ON prod, SSH ON dev? dev pushes commits back to GitHub, so it needs the SSH key you registered earlier. prod only ever pulls, it is a deployment target, never an authoring machine. HTTPS clone of a public repository requires no key, no token, no extra setup, which matches how real deployment servers usually have read-only access to the source repository. If your repository is private or you later need to push from prod, generate a second SSH key on prod and register it on GitHub the same way you did on dev.

Self-Reflection and Recap

REFLECTION QUESTIONS:

- Why does changing a single character in a file produce a completely different commit hash?
- What is the difference between `git add` and `git commit`, and why does Git separate the two steps?
- When would you use a branch instead of committing directly to `main`, and what happens when a merge produces a conflict?
- Why should you never commit `.env` files or API keys, and what is the role of pull requests in catching mistakes before they reach `main`?
- Why does PixelWise use SSH keys rather than passwords for Git operations across both VMs?

MILESTONE: Your project lives on GitHub with a README, a `.gitignore`, and a tagged release. The server has the repo cloned at `/opt/pixelwise`. Both machines share the same history. In the next block we turn the manual package installs into a reproducible setup script and tackle dependency management.

Dependencies & Structure

2026-06-12 · plucky mango Eichhörnchen

WORKS ON MY MACHINE.

The Why

Block 2 left us with a PixelWise repository on GitHub and a server that has Git and Python installed. Both machines share the same history, both can pull from origin, and the workflow that moves code between them is in place. What is still missing is the layer beneath the code: the system packages, the language runtime, the libraries, and the secrets your application reads (installed or provisioned by `setup-server.sh` and recorded in `.env/.env.example`).

KALTSTART CAPABILITIES, Your colleague clones the repository, runs the code, and immediately hits an Error. You installed a package with a specific version months ago and forgot. They did not. The remainder of the day is spent debugging an artefact of the install history rather than contributing real value to the program itself. At the end of this third block we want any fresh VM to go from empty to ready to code with a small, fixed sequence of commands. The two halves of that sequence are the load-bearing idea of the chapter: system packages on one side, Python packages on the other, version controlled in different files.

PIXELWISE now lives on GitHub and on both machines, but the repository contains only `README.md` and `.gitignore`. Before we can write application code, we need a reproducible way to install everything the project depends on, so that dev and prod stay in sync without human memory. Use `setup-server.sh` to provision system packages and the runtime, and keep secrets and runtime configuration in `.env.example` (copied to `.env`).




Figure 4: Two layers, two files. `setup-server.sh` provisions the operating system; `requirements.txt` provisions Python inside an isolated virtual environment. Together they reconstruct the environment from scratch on any fresh VM. System packages, the things you install with `apt`, are captured in a shell script called `setup-server.sh`. Python packages, the libraries your application imports, are captured in `requirements.txt` and installed inside an isolated virtual environment. The two files are committed to Git, run on either machine, and produce identical environments without any further intervention.

Hands On Experience

CONSIDER THIS SCENARIO: you hand your PixelWise repository to a classmate at the end of the day. They clone it, type `python3 train.py` without running the provisioning script `setup-server.sh`, and immediately hit `ModuleNotFoundError: No module named 'sklearn'`. You tell them to install scikit-learn. They do, but they get version 2.0 while you wrote the code against 1.4. The training script crashes with a different error. By the time the original problem is fixed, a few hours are gone and neither of you has touched the actual code nor added any value.

REPRODUCIBLE PROJECTS do not require remembering, asking, or guessing. The exact set of packages, the exact versions, the exact configuration contract, are written down once, committed to Git, and replayed on any machine that has the script and the file.

THIS THIRD BLOCK introduces dependency management and project structure for PixelWise. By the end you will have

- captured system-level setup in a reproducible shell script,
- created and activated a Python virtual environment,
- installed packages with pip and pinned them in `requirements.txt`,
- separated secrets from code using `.env` and `.env.example`,
- scaffolded the PixelWise project directory,
- and audited your dependencies for known security vulnerabilities.

THE TWO-COMMAND PROMISE. A well-managed project lets any new developer go from clone to running code with two commands: create the virtual environment and install from a pinned requirements file. No guessing, no mails or messages asking “which version of numpy do I need?” This block is what makes that promise hold.

OUTSIDE THE CLASSROOM, the same files are how cloud deployment works in practice. A continuous-integration runner clones the repository, executes the setup script (for example, `setup-server.sh`), installs from `requirements.txt`, and runs the tests. The reproducibility you build for two VMs scales without modification to a hosted server, a colleague’s laptop, or a CI runner you have never seen.

Optional Exercise: Catch Up to v0.1

CATCHING UP VIA THE REFERENCE TAG. From this block onwards, the GitHub repository is the safety net. The course maintains a public reference repository at <https://github.com/schutera/pixelwise> with a tag at the end of every block. v0.1 marks the end of Block 2; v0.2 will mark the end of this block. If you joined late, missed Block 2, or your repository has drifted from a known good state, reset to the previous tag rather than redoing every prior exercise.

THE MODEL: pull from schutera/pixelwise, push to your own <you>/pixelwise. The reference repository is read-only for you; the place you commit to is your personal copy on GitHub. The steps below assume nothing from Block 2 is in place, so the exercise is fully self-contained.

TAG MAP. v0.1 is the end of Block 2, v0.2 is the end of this block, v0.3 the end of Block 4, and so on. Every tag is a complete, runnable state of PixelWise. Browse them at <https://github.com/schutera/pixelwise/tags>.

CREATE AN EMPTY pixelwise REPOSITORY ON GITHUB under your own account. Log in to GitHub, click New repository, name it pixelwise, and leave it empty, no README, no .gitignore, no licence, since the push step below will populate it. Without this repository the git remote set-url URL points at a path GitHub cannot find, and the first push reports does not appear to be a git repository.

RESTORE B1-complete ON dev so you start from the same baseline as everyone else, with Ubuntu installed and the user account in place but nothing on top.

INSTALL GIT, CONFIGURE YOUR IDENTITY, AND GENERATE A SEPARATE SSH KEY FOR GITHUB on the fresh VM. B1-complete predates Block 2, so the new VM has no git binary and no Git config. It does already have ~/.ssh/id_ed25519 and ~/.ssh/id_ed25519.pub from Block 1, the key pair you use to log in to prod. The -f flag below writes the new key to a distinct path, so the existing prod-login key files stay byte-for-byte intact on disk and your password-less SSH into prod keeps working:

```
sudo apt update && sudo apt install -y git
git config --global user.name "Your Name"
git config --global user.email "your@email.com"
ssh-keygen -t ed25519 -C "your@email.com" \
    -f ~/.ssh/id_ed25519_github
ls -l ~/.ssh/id_ed25519*
cat ~/.ssh/id_ed25519_github.pub
```

TELL SSH WHICH KEY TO USE FOR GITHUB.COM. With two keys in `~/.ssh/`, SSH may offer the prod-login key first and GitHub authenticates you as the wrong account, or fails outright. Worse, if `ssh-agent` has already cached the prod-login key, it gets offered before any IdentityFile on disk, and the symptom is a misleading `git@github.com:<you>/pixelwise.git` does not appear to be a git repository. Open (or create) `~/.ssh/config` in nano and add an entry that pins the GitHub key to github.com and bypasses the agent for that host:

```
nano ~/.ssh/config
```

Paste the following block at the end of the file, then save and exit nano with Ctrl+X, Y, Enter:

```
Host github.com
  HostName github.com
  User git
  IdentityFile ~/.ssh/id_ed25519_github
  IdentitiesOnly yes
  IdentityAgent none
```

LOCK DOWN THE FILE PERMISSIONS. SSH refuses to read config if it is world- or group-readable, so tighten it to the owner only:

```
chmod 600 ~/.ssh/config
```

RUN THE VERIFICATION HANDSHAKE. The first SSH connection to GitHub asks you to confirm GitHub's host fingerprint and writes it into `~/.ssh/known_hosts`:

```
ssh -T git@github.com
```

Type yes when prompted to trust the fingerprint. GitHub then answers Hi <you>! You've successfully authenticated, but GitHub does not provide shell access. The shell-access line sounds like an error but is the intended response: GitHub allows Git operations over SSH, not interactive shells.

CATCH-UP PROCEDURE ON dev. With the empty repository in place and the key registered, clone the reference repository, rewind main to the v0.1 tag, point origin at your own empty repository, and push:

SANITY CHECK. `ls -l` should now list four files under `~/.ssh/`: the `id_ed25519` pair from Block 1 and the `id_ed25519_github` pair you just generated. If the Block 1 files are missing, you accepted the default path during `ssh-keygen` and overwrote the prod key. Restore B1-complete and rerun the step with the `-f` flag.

ADD THE KEY TO GITHUB. Copy the entire `ssh-ed25519 ...` line printed by `cat ~/.ssh/id_ed25519_github.pub`, go to Settings, then SSH and GPG keys, then New SSH key on GitHub, give it a title like `pixelwise-dev-vm`, and paste.

WHY IdentityAgent none? IdentitiesOnly yes alone tells ssh to ignore extra agent keys, but in practice some agent setups still offer them first. IdentityAgent none is the belt-and-braces fix: for connections to github.com, skip the agent entirely and read the key from the file you named, period. Other hosts, including prod, are unaffected and still use the agent normally, so password-less SSH into prod keeps working.

ALREADY POLLUTED THE AGENT? If you ran `ssh-add` earlier in the session and `ssh-add -l` now lists more than one key, drop them all from the agent and reload only the GitHub key:

```
ssh-add -D
ssh-add ~/.ssh/id_ed25519_github
```

`ssh-add -D` only forgets the keys held in agent memory; the files in `~/.ssh/` are not touched. Re-add the prod key later with `ssh-add ~/.ssh/id_ed25519` when you next log in to prod if you prefer not to retype the passphrase.

```

cd ~
git clone https://github.com/schutera/pixelwise.git
cd pixelwise
git reset --hard v0.1
git remote set-url origin \
    git@github.com:<you>/pixelwise.git
git push -u origin main
git push --tags

```

MIRROR IT ON prod. By the end of Block 2, PixelWise also lived on the production VM at /opt/pixelwise. Restore B1-complete on prod as well, install git, then clone the reference repository to its permanent home and rewind main to v0.1:

```

sudo apt update && sudo apt install -y git
sudo mkdir -p /opt/pixelwise
sudo chown produser:produser /opt/pixelwise
git clone https://github.com/schutera/pixelwise.git \
    /opt/pixelwise
cd /opt/pixelwise
git reset --hard v0.1

```

You are now at the state where Block 2 ended, with both VMs aligned, ready to start Block 3.

WHY reset --hard? The clone lands you on the reference repository's current main, which is ahead of v0.1. git reset --hard v0.1 moves your local main back to the tagged commit, discarding the later ones, so the working tree matches the end-of-Block 2 state exactly. There is no work to lose because you have not committed anything yet.

NO SSH KEY ON prod. prod only ever pulls; it never pushes, so an HTTPS clone of a public repository is enough and no SSH key needs to be registered with GitHub. Skip the ssh-keygen ceremony you just did on dev. If your repository is private, or you later need to push from prod, generate a separate key on prod and add it to GitHub the same way.

ALREADY HAVE A FORK? If your own fork already carries the v0.1 tag, swap the URL for git@github.com:<you>/pixelwise.git on dev and the HTTPS equivalent on prod, so push and pull continue to target your account rather than the reference repo.

The Setup Script

PROVISIONING A SERVER is the act of turning a freshly installed operating system into one that can run your application. So far it has been done by typing `sudo apt install` commands at the prompt and reading their output. That works for one-time setups, it does not scale. Software is meant to be scaled.

THE SETUP SCRIPT is a plain Bash file that captures every `apt install` the project needs. It lives in the repository alongside the code and is run on any fresh VM to bring it up to speed:

```
#!/bin/bash
# setup-server.sh: provision a fresh server VM
set -euo pipefail

sudo apt update
sudo apt install -y git python3 python3-pip \
    python3-venv curl
```

THE FIRST LINE OF THE BODY, `set -euo pipefail`, makes the script fail loudly. `-e` aborts on the first error, `-u` aborts on an undefined variable, and `-o pipefail` propagates a failed command in the middle of a pipeline rather than letting a successful command later on mask it. Half-configured machines are a common cause of mysterious bugs; the four-character preamble eliminates most of them.

MARK IT EXECUTABLE and run it:

```
chmod +x setup-server.sh
./setup-server.sh
```

SEE HOW VERSION CONTROL ALREADY PAYS OFF. The script is committed to Git, so when the VM is rebuilt or a new team member joins, `bash setup-server.sh` reproduces the exact same environment. The commands you typed once are written down where the next person can find them.

CONFIGURATION MANAGEMENT is a whole discipline that grew from this small itch. Ansible, Salt, and Puppet declare the desired state of a machine in a structured format and reconcile the live system to match. `setup-server.sh` is the entry-level version: a script you can read end to end. The principle is the same.

Ansible: <https://docs.ansible.com/>

THE SHEBANG LINE `#!/bin/bash` tells the kernel which interpreter to use when you run the file directly. Without it, the file is read as a sequence of commands by whichever shell happens to invoke it, which on Ubuntu is usually `dash`, not `Bash`, and the two differ in subtle ways that surface only when something breaks.

IDEMPOTENCE is the property that running a script twice yields the same result as running it once. `apt install` is idempotent by default: an already-installed package is a no-op. So re-running is always safe, and the script can be used to fix a machine that has drifted.

Exercise: The Setup Script

WITH THE SETUP SCRIPT MOTIVATED ABOVE, put it on the dev VM. At the root of the PixelWise repository, that is ~/pixelwise/ where you ran `git init` in Block 2, create a file called `setup-server.sh` that captures the system packages installed manually in Block 2:

```
#!/bin/bash
set -euo pipefail

sudo apt update
sudo apt install -y git python3 python3-pip \
    python3-venv curl
```

MARK IT EXECUTABLE AND RUN IT. On a machine that already has the packages, every `apt install` is a no-op; the script reports nothing to do, which is the desired outcome of an idempotent provisioning step:

```
chmod +x setup-server.sh
./setup-server.sh
```

COMMIT IT. This is the first file that makes provisioning reproducible:

```
git add setup-server.sh
git commit -m "Add server provisioning script"
git push
```

WHY AT THE REPO ROOT? Provisioning is a project-wide concern, not the responsibility of any single subdirectory, so `setup-server.sh` sits next to `README.md` and `.gitignore` where anyone cloning the repository will spot it immediately. Open the file with `nano ~/pixelwise/setup-server.sh` or your editor of choice and paste the contents below.

(OPTIONAL) TRY IT ON A FRESH VM. The real test of the script is on a machine that has not seen the commands yet. A second prod VM, restored from the B1-complete image, clones the repository, runs `bash setup-server.sh`, and provisions end to end without you typing a single `apt install`.

Virtual Environments

A VIRTUAL ENVIRONMENT is an isolated Python installation. Each project gets its own set of packages, independent of every other project and of the system Python. The mechanism is local: a directory, by convention called `.venv/`, that holds an interpreter, a copy of `pip`, and every library the project depends on. Activating the environment rewrites your shell's `PATH` so that `python` and `pip` resolve to the local copies.

```
python3 -m venv .venv
source .venv/bin/activate
```

THE DIRECTORY IS LARGE and machine specific, often 50 to 200 MB. Thus it is listed in `.gitignore`: never commit it. A teammate clones the repository, creates their own `.venv`, installs from `requirements.txt`, and ends up with the same packages without ever copying the binary contents.

VERIFY THE SWITCH. After activation, which `python` should report a path inside `.venv/bin/`, and `pip list` should show only `pip` and `setuptools`:

```
which python
# /home/user/pixelwise/.venv/bin/python

pip list
# Package      Version
# -----
# pip          24.0
# setuptools 69.5.1
```

A fresh environment starts empty. Every package you install from this point is tracked and intentional. `deactivate` returns the shell to the system Python.

THE RULE IS SIMPLE: every project gets its own `.venv`. Activate it before you work; deactivate when you are done. The one-time cost is a directory and a command; the ongoing benefit is never debugging a version conflict between two unrelated projects on the same laptop.

WHY ISOLATION? Project A needs `numpy==1.24`; project B needs `numpy==2.0`. Without per-project environments, one of them breaks the moment the other is installed. With virtual environments, both coexist on the same machine and make it easy for you to switch between projects seamlessly.

THE DATA SCIENCE ECOSYSTEM often uses `conda` instead of `venv`. `Conda` manages both Python packages and system-level libraries like `CUDA` or `MKL`, which makes it popular for GPU workflows. For web projects, `venv` is the lighter and more standard choice.

Conda: <https://docs.conda.io/>

Exercise: The Virtual Environment

WITH THE RATIONALE FOR ISOLATION IN MIND, create one on the dev VM, inside the PixelWise repository:

```
cd ~/pixelwise
python3 -m venv .venv
source .venv/bin/activate
```

VERIFY THE SWITCH. which python should print a path inside .venv/bin/. pip list should show only pip and setuptools:

```
which python
pip list
```

A fresh environment is empty by design. The next exercise fills it with intent.

ADD THE ENVIRONMENT TO .gitignore. The .venv/ directory must never be committed. Open the .gitignore from Block 2 and confirm .venv/ is listed; if not, add it now. You will notice that currently we ignore *.py python files, but in future you do not want to ignore those, so remove that line:

```
.venv/
```

Run git status. The only change shown should be the .gitignore edit itself, even though .venv/ sits right next to it on disk with hundreds of files inside. That silence is .gitignore doing its job: the absence of .venv/ in the output is the demonstration, not a missing piece.

COMMIT THE CHANGE. With the diff read and git status clean of surprises, record the update so the next exercise starts from a known state:

```
git add .gitignore
git commit -m "Ignore .venv/ and stop ignoring *.py"
git push
```

NOTICE THE PROMPT CHANGE. After activation, your shell prompt usually gains a (.venv) prefix to remind you which environment is active. If you ever wonder whether you are inside the project's environment, the prompt is the first place to look; which python is the second.

READ THE DIFF. git diff .gitignore shows what you added compared to the Block 2 version. Reading the diff before committing is the habit that catches accidental changes early; running it once now is muscle memory you will reuse every day for the rest of the course.

Managing Packages with pip

PIP IS THE PACKAGE INSTALLER for Python. It downloads packages from PyPI, the Python Package Index at <https://pypi.org/>, and installs them into the active environment. The core cycle has three commands:

```
pip install scikit-learn joblib python-dotenv
pip freeze > requirements.txt
pip install -r requirements.txt
```

`pip install` downloads each named package and its dependencies and places them in `.venv/`. `pip freeze` prints every installed package with its exact version, including transitive dependencies you never asked for directly. `pip install -r` reads such a file and installs everything in it, the step that turns a captured environment back into a live one.

Anatomy of `requirements.txt`

ONE PACKAGE PER LINE, pinned to an exact version:

```
scikit-learn==1.4.0
joblib==1.3.2
python-dotenv==1.0.1
```

PINNING MATTERS because packages are a moving target. A package you install today as `numpy>=1.24` may resolve to 1.27 tomorrow, with new features added, new behaviour, new bugs, and possibly new security advisories. Packages are software, and they are subject to change; just the same as your own project. Pinning to `==1.24.3` freezes the resolution and makes switching to newer versions a deliberate act.

COUNT WHAT `pip freeze` WRITES. Three direct installs typically produce twenty or more lines, the rest are transitive dependencies pulled in automatically. The flatness is the point: every package that has to be present for the project to work is listed, even the ones you never typed. The cost is that the file is hard to read; the benefit is that it captures exactly what is on disk.

THE THREE PACKAGES. `scikit-learn` provides the classifier `PixelWise` will train and serve. `joblib` saves and loads trained models to disk as `.pkl` files. `python-dotenv` reads `.env` files into environment variables so secrets stay out of code.

READ THESE THREE LINES AS A CONTRACT. The first installs into the environment. The second writes a snapshot of it to a text file. The third reconstructs the same environment from the file on any machine. Together they are the reproducibility layer for your programming environment.

VERSION SPECIFIERS. `==` pins exactly, `>=` sets a floor, `~` allows patch updates: `~=1.4.0` means `>=1.4.0`, `<1.5.0`. For maximum reproducibility, pin everything with `==`. Looser specifiers are appropriate for libraries that other projects depend on, which is not where most of your work lives.

Semantic versioning: <https://semver.org>

SECURITY AUDITING. Your dependencies have their own dependencies, and any of them can harbour known vulnerabilities. `pip-audit` checks every installed package against the CVE database: `pip install pip-audit` && `pip-audit`. Run it after every `pip freeze`. In Block 8 we wire it into a cron job so vulnerabilities surface without human labor.

`pip-audit`: <https://github.com/pypa/pip-audit>

Dependabot: <https://docs.github.com/en/code-security/dependabot>

Exercise: Pinning Dependencies and pip-audit

WITH `pip` AS THE PACKAGE INSTALLER AND `requirements.txt` AS THE CAPTURED ENVIRONMENT, put both to work. With the virtual environment active, install the three direct dependencies PixelWise needs in this block:

```
pip install scikit-learn joblib python-dotenv
```

FREEZE THE ENVIRONMENT into a pinned requirements file:

```
pip freeze > requirements.txt
```

Open `requirements.txt` and inspect. Note the pinned versions and the transitive dependencies that were pulled in automatically. Count how many lines `pip freeze` produced compared to the three packages you installed directly; the gap is the dependency graph working invisibly.

AUDIT THE DEPENDENCY GRAPH. Install `pip-audit` and run it against the active environment:

```
pip install pip-audit
pip-audit
```

Inspect the output. What is a CVE? What would you do if a critical vulnerability were reported in one of your dependencies?

COMMIT THE REQUIREMENTS:

```
git add requirements.txt
git commit -m "Pin core Python dependencies"
```

THE HABIT TO FREEZE AND COMMIT: after installing or upgrading any package, immediately run `pip freeze > requirements.txt` and commit the change. Stale requirements files are the most common cause of “works on my machine” incidents.

(OPTIONAL) READ ONE CVE. Pick any advisory `pip-audit` reports, click through to the CVE record, and read the description, the affected versions, and the fixed version. Practising this once now makes the production response straightforward when an alert fires under time pressure.

Configuration: .env and python-dotenv

SEPARATING CONFIGURATION FROM CODE is a foundational principle. Configuration is anything that varies between development, staging, and production: API keys, database URLs, feature flags, debug settings. Code is the part that does not change between deployments.

A `.env` FILE holds key-value pairs that your application reads at startup:

```
SECRET_API_KEY=my-actual-secret-key-here
DEBUG=true
```

In Python, `python-dotenv` loads the file into process environment variables, after which the standard `os.getenv` reads them:

```
from dotenv import load_dotenv
import os

load_dotenv()
api_key = os.getenv("SECRET_API_KEY")
debug = os.getenv("DEBUG", "false").lower() == "true"
```

The .env.example Pattern

THE REAL `.env` contains real secrets, so it must never be committed. `.env.example` is the contract: it lists which variables exist, explains what each one is for, and provides a safe placeholder. The example file is committed and public. The filled contracts live on the respective machines and stays out of Git.

```
# API key callers must send in the X-API-Key header
# (a secret, never a real value here)
SECRET_API_KEY=replace-me

# Debug mode: when true, FastAPI shows /docs and full
# error tracebacks. Never true in production.
DEBUG=true
```

ON FIRST SETUP CREATE YOUR ACTUAL `.env` by copying `.env.example` and replacing each placeholder with a real value. The contrast is immediate: one is the template, the other is a vault for your secrets.

The Twelve-Factor App: <https://12factor.net/>
 python-dotenv: <https://github.com/theskumar/python-dotenv>

A FILE FOR CONVENIENCE. Environment variables can be exported manually, set inline before a command, or supplied by the `systemd` service file we install in Block 5. A `.env` file is the development-time convenience that mirrors the production-time mechanism: same names, same shape, different source.

The pattern grows block by block. Every new service we add, the API key in Block 5, the database URL in Block 6, the model path in Block 4, lands in `.env.example` as a new line with a comment explaining its purpose. A teammate cloning the repository six weeks from now sees the full surface of the application's configuration in one file.

Exercise: Configuration and Secrets

WITH `.env` AND `.env.example` AS THE CONTRACT for separating configuration from code, draft both files. Create `.env.example` at the repository root with two entries:

```
# API key (secret, never a real value here)
SECRET_API_KEY=replace-me

# Debug mode (never true in production)
DEBUG=true
```

CREATE YOUR REAL `.env` by copying the example and replacing the placeholder with a value you choose:

```
cp .env.example .env
```

Edit `.env` so `SECRET_API_KEY` holds a value of your own choosing. Make sure to add `.env` to `.gitignore` if it is not already there, then Run `git status` and confirm that `.env` does not appear.

COMMIT THE EXAMPLE ONLY:

```
git add .env.example .gitignore
git commit -m "Add .env.example and update .gitignore"
```

The contrast is immediate. The example is now public, the real values stay on the machine.

THE NAMING CONVENTION is upper-case identifiers with underscores. Environment variables are case sensitive on Linux but case insensitive on parts of Windows; convention picks the form that works everywhere.

IF `.env` APPEARS IN `git status`, the `.gitignore` entry is missing. Add it now and verify again before committing anything. A secret committed once is in the history forever, even if the next commit deletes it.

Exercise: Replicate on the Server

THE REPRODUCIBILITY PROMISE is verifiable only by running the recipe on a different machine. SSH into the prod VM, pull the latest code, and recreate the environment from the two files you committed:

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git pull origin main
bash setup-server.sh
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
```

VERIFY THE AUDIT ON prod:

```
pip-audit
```

The same advisories should appear as on dev. Identical inputs, identical outputs. This is what reproducibility looks like in practice.

WHY origin main? After the catch-up `reset --hard v0.1`, prod's local main may have lost its upstream tracking, in which case bare `git pull` asks you to specify a branch. Naming origin main explicitly sidesteps that. To restore the tracking once and let plain `git pull` work afterwards, run `git branch --set-upstream-to=origin/main`.

WATCH THE INSTALL OUTPUT. `pip` resolves the dependency graph from the lock-shaped `requirements.txt` and downloads each package at the pinned version. The output you see on prod should match the output you saw on dev the first time you ran `pip install`; if it does not, something has drifted.

FROM THIS BLOCK ONWARDS, the GitHub repository is the single source of truth. prod only ever pulls; it never authors. A `git pull` on prod should always be the deployment step, never a place to make edits.

Optional: When requirements.txt Starts to Hurt

`pip freeze > requirements.txt` captures everything: your direct dependencies and all their transitive dependencies, flattened into one list. Six months later you want to upgrade one package: which of those forty-seven packages did you deliberately install, and which were pulled into that list automatically? Without that distinction, every upgrade is a guessing game. Solving dependency conflicts becomes a nightmare and quickly feels like brute-force.

`pyproject.toml` solves this by separating “what you depend on”, your direct deps loosely pinned, from “what gets installed”, a generated lock file with the full resolved graph. Upgrades become surgical: bump one line, regenerate the lock, done.

```
[project]
name = "pixelwise"
version = "0.1.0"
dependencies = [
    "scikit-learn>=1.4",
    "joblib>=1.3",
    "python-dotenv>=1.0",
]
```

COMPARE this to `requirements.txt`: the project file lists only what you chose to install, with flexible version bounds. The tool resolves the full dependency graph and writes a lock file that pins every transitive dependency exactly. You commit both files; the lock file is your reproducibility guarantee and the project file is your editable surface.

```
# to upgrade a single package and update lock:
poetry update <package>
```

TOOLS LIKE `uv` and `poetry` use the project-file model and are fast becoming the standard. Once the `requirements.txt` workflow is second nature, this is the upgrade that makes dependency management sane at scale.

Python Packaging User Guide: <https://packaging.python.org/en/latest/>
`uv`: <https://docs.astral.sh/uv/>
`Poetry`: <https://python-poetry.org/>

WE STAY WITH `requirements.txt` through the rest of the course because it is the format every Python developer recognises and the only one some hosting environments still accept. The reasoning generalises: when an upgrade is available, learn the workflow beneath it before adopting the tool that abstracts it away.

Optional: The Container Alternative

VIRTUAL ENVIRONMENTS ISOLATE PYTHON PACKAGES. Containers isolate the entire operating system: libraries, system packages, Python version, environment variables, even the filesystem layout. Where `requirements.txt` specifies which Python packages to install, a `Dockerfile` does the whole stack in layers: This Linux image, install these system packages, then install these Python packages.

A MINIMAL `DOCKERFILE` for a Python service looks much like the two scripts you are about to write, compressed into a single declarative file:

```
FROM python:3.12-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install -r requirements.txt
COPY . .
CMD ["python", "main.py"]
```

The result is a self-contained unit that runs identically on any machine with Docker installed. The trade is opacity: when something fails inside a container, the layers between your code and the kernel multiply, and debugging across them requires familiarity with the layer above and below.

WE DO NOT USE DOCKER IN THIS COURSE, but you should know it exists and why teams adopt it. The vocabulary, image, container, registry, layer, volume, recurs in any later production work; the next-block decisions in this course translate to it directly.

A NOTE ON WHAT COMES NEXT. The repository now has a scaffold, pinned dependencies, and a provisioned server, but no application code. `app/` is an empty Python package waiting for content. In the next block we drop a trained model into `models/`, write the inference layer in `app/classifier.py`, and verify end to end that `PixelWise` can classify a digit before any API or frontend exists.

DOCKER and Docker Compose solve the reproducible environment problem at the OS level rather than at the Python level. This course teaches the `virtualenv` path to keep the OS layer visible and learnable. Later Docker enables you to run Linux containers under Windows as well. Docker is the natural next step after this course; it does not replace any of what you learn here, it builds on it.

Docker: <https://www.docker.com/>

Docker docs: <https://docs.docker.com/get-started/>

WHY WE HOLD OFF. The OS layer is the layer most students are least exposed to. Containers hide it before students have seen it. By the end of this course you will have provisioned, deployed, and monitored a service on bare Ubuntu; lifting that into a container at that point is mechanical, and you will know exactly what each line of the `Dockerfile` means.

Self-Reflection and Recap

SELF-REFLECTION questions to guide your thoughts during and after the exercises:

- Why do we use a virtual environment instead of installing packages system-wide?
- What is the difference between `.env` and `.env.example`, and why do we need both?
- What does `pip freeze` capture that you did not explicitly install, and why does it matter?
- If a colleague clones your repository, which two commands do they run to get a working environment?
- What is the risk of committing `.env` to Git, even briefly, and what would the recovery look like?
- Why does the course teach virtual environments before containers, even though many production systems use containers?

RECAP of key concepts:

- `setup-server.sh` captures system-level provisioning so any fresh VM can be brought up with a single command.
- Virtual environments isolate Python packages per project; never install globally.
- `requirements.txt` pins every dependency so both machines run identical code.
- `.env.example` documents the configuration contract; `.env` holds the real values and never enters Git.
- `pip-audit` catches known vulnerabilities in the dependency graph and runs anywhere a Python environment lives.

MILESTONE. PixelWise has `app/`, `data/`, `models/`, a virtual environment, pinned dependencies, a setup script, and a minimal `.env.example`. The `.gitignore` is doing real work. Anyone who clones the repository can recreate the environment in a handful of commands. The next chapter integrates a trained machine learning model so the project can actually classify a digit.

WITHOUT WRITING APPLICATION CODE YET, we already have a reproducible system on both machines. Anyone with the repository and the script can reach the same starting line.

TEASER. What does it take to turn a trained model file into something the rest of the application can call?

Integrate the Model


2026-06-12 · plucky mango Eichhörnchen

TRAIN HARD, SHIP EASY.

The Why

Block 3 left us with a reproducible environment on both machines. The virtual environment is ready, the dependencies are pinned, the project foundation is in place. A foundation that would work for all sorts of different projects by the way. Let's get our project some specific functionality.

WE BUILD PIXELWISE. This block is the first that produces application logic in the form of a neural network, in the following referred to as model. We assume the model is already trained and available, and our task is to integrate it into the codebase and make it callable.



THE GAP BETWEEN A TRAINED MODEL AND A WORKING SYSTEM is larger than most students expect. A .pkl file sitting in a folder is not a product. It needs to be loaded safely, called with the exact input format it was trained on, and wrapped in an interface the rest of the application can depend on, suffice to say validated. The perfect mdoel however, is not useful if it is not integrated into the system, so it can be exposed and put to work.

Figure 5:

The inference pipeline for one image: the caller binarises a raw (28,28) uint8 array at threshold 128 and flattens it into a (1,784) row vector before feeding it to the loaded .pkl model. The model's internal Binarizer is a redundant idempotent guard on already-binarised input, and predict_proba turns the result into a dictionary of class scores.

Hands On Experience

CONSIDER THIS SCENARIO: a colleague hands you a trained model file and a one-line note, "ship it". This block teaches you how to, and why you should, ask the right questions to your colleague.

THIS FOURTH BLOCK integrates a trained model into PixelWise and builds the inference layer. By the end you will have

- understood what a model file is and why it is not source code,
- loaded the model safely, learned about and verified its class contract at startup,
- implemented a batch-first inference interface in `app/classifier.py`,
- written a smoke test that proves the integration honours the contract,
- and read the canonical model card alongside your introspection output to verify the contract.

VERIFICATION AND VALIDATION are two different questions. Verification asks whether the system was built right: does the integration return what the model is trained to produce, with the expected shape, dtype, and class set? Validation asks whether the right system was built: is this model the right choice for the problem in the first place and is the performance sufficient? This block is squarely about verification. Validation belongs to pre-integration.

Optional Exercise: Catch Up to v0.2

CATCHING UP VIA TAG. This catch-up extends the one in Block 3. If you have not yet done the Block 3 catch-up, do that first: it sets up your GitHub repository, SSH key, Git identity, and remote URL, and lands you at v0.1, the end-of-Block 2 state. The exercise here picks up from there and rewinds to v0.2, the end-of-Block 3 state, so you arrive at the starting point this block expects: a provisioned VM with the virtual environment, pinned dependencies, and .env.example in place.

TWO LAYERS TO RESTORE. Block 3 produced both Git-tracked files (setup-server.sh, requirements.txt, .env.example, plus an updated .gitignore) and unversioned state (system apt packages, the .venv/ directory, real .env values). `git reset --hard v0.2` restores the first; the second has to be reproduced by replaying Block 3's setup commands. Both VMs need both layers, so run the sequence below on each.

RUN THE SAME SEQUENCE ON BOTH VMs. Rewind the repository to v0.2, reinstall system packages, recreate the virtual environment, reinstall pinned dependencies, and copy .env.example to .env. The dev sequence runs in ~/pixelwise; the prod sequence is identical except the directory is /opt/pixelwise and you reach it via `ssh produser@192.168.56.11` first. The exact commands are in the margin.

YOU ARE NOW at the state where Block 3 ended, with both machines aligned, ready to start Block 4. Did you realize how seamless this was, that is thanks to our solid git practices, and our dependency and setup scripts.

TAG MAP. v0.2 marks the end of Block 3, v0.3 will mark the end of this block. Browse the full set at <https://github.com/schutera/pixelwise/tags>.

SEQUENCE ON dev.

```
cd ~/pixelwise
git fetch origin --tags
git reset --hard v0.2
bash setup-server.sh
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
cp .env.example .env
Then edit .env with real values.
```

SEQUENCE ON prod.

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git fetch origin --tags
git reset --hard v0.2
bash setup-server.sh
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
cp .env.example .env
```

What a Model File Actually Is

A `.pkl` FILE IS A BUILD ARTEFACT, not source code. You distribute it; you do not commit it to your repo. Just as a compiled binary is produced by a compiler, a model file is produced by a training script. The training script and its configuration is the source of truth: It takes data, runs an algorithm, and writes weights, hyperparameters, and any preprocessing into a single serialised object.

VERSIONING MAKES THAT CONTRACT DURABLE. Files carry a version in the filename or in metadata shipped alongside. A retrain on more data bumps the patch version; an architecture change or a different class set bumps the major version. Integration code pins to a specific version and updates deliberately, the same discipline you would apply to any external API. Open source tooling that automates this includes `DVC` for git-style data and model tracking, `MLflow` and `ClearML` as self-hostable experiment registries, `Aim` as a lighter-weight tracker, and the Hugging Face Hub itself, which is free for public repositories and versions every push as a git revision.

HUGGING FACE IS THE GITHUB OF MODELS. The Hugging Face Hub at <https://huggingface.co> hosts hundreds of thousands of public model repositories, each with versioned weights, a model card describing it. Downloading a specific revision is a one-line call against the `huggingface_hub` library. The same workflow runs against private repositories with a token, which is how many companies share proprietary models internally. For this course the artefact lives in the course-maintained model repository at <https://github.com/schutera/pixelwise-model>, with a tagged release per model version and the model card sitting next to it.

LOADING A `.pkl` FROM AN UNTRUSTED SOURCE CAN EXECUTE ARBITRARY CODE. `pickle` and `joblib` deserialise objects by reconstructing them, including any code embedded in the file. This is convenient when you produced the file yourself; it is dangerous when the file came from somewhere you do not control. Treat `.pkl` files the way you treat executables: only run what you trust.

TRAINING FROM SCRATCH EVERY TIME IS UNREALISTIC. Even our small classifier takes seconds; a production scale model takes days of GPU time and terabytes of data. The training team tunes, the inference team integrates, and the artefact is the contract between them.

PULLING A MODEL FROM THE HUB.

```
from huggingface_hub import
hf_hub_download
path = hf_hub_download(
    repo_id="distilbert-base-uncased",
    filename="model.safetensors",
    revision="v1.0",
)
```

THE TRAINING SCRIPT `train.py` lives in the course `pixelwise-model` repo alongside the artefact it produces, not in `PixelWise`. The artefact and the script that produced it travel together, so any consumer of a tagged release can audit how that exact `.pkl` came to be. Students who want to understand how the artefact was produced can run it: MNIST downloads automatically via `sklearn.datasets.fetch_openml`, and training takes ~30 seconds on CPU. Further reading on training practices: Karpathy, A Recipe for Training Neural Networks: <https://karpathy.github.io/2019/04/25/recipe/> Google, Rules of Machine Learning: <https://developers.google.com/machine-learning/guides/rules-of-ml>

Exercise: Pull the Model into Dev

TREAT THE MODEL ARTEFACT AS SOMETHING YOU LOAD, not something you commit to PixelWise. The course hosts the model pkl as a tagged release in the central pixelwise-model repo; you pull it into models/ of the main project, you do not redistribute it.

PIN THE MODEL VERSION in PixelWise. Add to .env.example so others know which release the integration expects, and .env so your code knows:

```
MODEL_REPO=https://github.com/schutera/pixelwise-model.git
MODEL_VERSION=v1.0
```

TEACH setup-server.sh TO FETCH THE MODEL. Pulling the artefact by hand once is fine for understanding what happens; encoding the same steps in setup-server.sh means that re-running the script on prod, on a fresh VM, or after a snapshot restore brings the model back without any extra commands. Open setup-server.sh in nano and append a model-pull block after the existing apt install lines:

```
nano setup-server.sh
```

Add the following at the bottom of the file, then save and exit nano with Ctrl+X, Y, Enter:

```
# Pull the pinned model artefact
if [ -f .env ]; then
    set -a; source .env; set +a
    if [ -n "${MODEL_REPO:-}" ] && \
        [ -n "${MODEL_VERSION:-}" ]; then
        mkdir -p models/
        rm -rf /tmp/pixelwise-model
        git clone --depth 1 --branch "$MODEL_VERSION" \
            "$MODEL_REPO" /tmp/pixelwise-model
        cp /tmp/pixelwise-model/*.pkl models/
        cp /tmp/pixelwise-model/MODELCARD.md models/
        rm -rf /tmp/pixelwise-model
    fi
fi
```

RUN THE UPDATED SCRIPT ON dev:

```
bash setup-server.sh
ls models/
git status
```

THE COURSE ARTEFACT. at <https://github.com/schutera/pixelwise-model> digit_classifier_v1.pkl is a LogisticRegression pipeline trained on MNIST digits 1 to 9 (9 classes, ~54,000 training samples, public domain). Class 0 is intentionally held back; students can add it later as a v2 release without collecting any new data.

WHY PULL FROM A SEPARATE REPO? The model has its own release cadence and its own contract. Pinning a specific tag in PixelWise's .env makes the integration explicit: a v2 only takes effect when the integration code chooses to. The optional exercise at the end of the block walks through publishing your own v2 to a fork of the model repo.

WHY ALL THE GUARDS? The if [-f .env] check lets the script run on a VM that has not yet been configured (for example, immediately after the Block 3 catch-up). The rm -rf /tmp/pixelwise-model before git clone keeps the script idempotent: a leftover temp directory from a previous run does not block the next one. Together these turn the model pull into a re-runnable provisioning step rather than a one-shot command.

IF digit_classifier_v1.pkl APPEARS IN git status, your .gitignore from Block 3 needs an update. Add models/*.pkl on its own line and rerun git status. A model file accidentally committed once is in the history forever, even if the next commit deletes it.

CONFIRM THE ARTEFACT LOADS in a Python shell on dev:

```
source .venv/bin/activate
python3
>>> import joblib
>>> model = joblib.load("models/digit_classifier_v1.pkl")
>>> type(model)
```

If `joblib.load` returns a Pipeline object without raising, the artefact is sitting at `models/digit_classifier_v1.pkl` on your dev VM, ready to be scrutinised next.

COMMIT THE SCRIPT AND `.env.example`, then mirror the change to prod:

```
git add setup-server.sh .env.example
git commit -m "Pull model artefact in setup-server.sh"
git push
```

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git pull origin main
cp .env.example .env          # Don't forget to edit .env to fir production
bash setup-server.sh
ls models/
```

MACHINE SETUP AND MODEL PROVISIONING BECOMES A SINGLE COMMAND. After this block, both dev and prod are set up by one command, `bash setup-server.sh`, regardless of whether the machine setup or the model artefact, or both. Optionally also do that for the virtual environment and the system packages.

The Model Contract and Its Card

EVERY MODEL WAS TRAINED ON INPUTS IN A SPECIFIC FORMAT.

The format is part of the contract: a 28-by-28 greyscale image, pixel values in $[0, 1]$, flattened to 784 features. Most violations of this contract are loud: a wrong shape raises `ValueError` on the first call, a wrong dtype trips an assertion in the pipeline, a missing class never appears in the predictions at all. The integration layer's job is to catch these before the model reaches production.

SILENT FAILURES live in the gap between syntactic and semantic correctness. Same dtype, same dimensions, but pixel range $[0, 255]$ instead of $[0, 1]$, or raw greyscale instead of binarised. The model returns a class label and a confidence score with no way to tell that the input did not come from the training distribution. This is the train/serve skew bug class, and the only fix is an explicit preprocessing step that mirrors training.

THE MNIST MODEL EXPECTS:

- a 28×28 greyscale image,
- pixel values normalised to $[0, 1]$,
- flattened to a 784-element vector.

THE SKLEARN PIPELINE bundles preprocessing and model into one serialised object that is itself the contents of `digit_classifier_v1.pkl`. `joblib.load` returns the whole Pipeline, preprocessing steps included, so loading the file is loading the preprocessing, so we reduce steps a maintainer might forget; adhering the contract the training script established.

THE MODEL CARD IS WHERE YOU WRITE THE CONTRACT DOWN.

A model card answers four questions any consumer of the model should know the answer to before they trust it: where did the data come from, what can the model do, what does it fail at, and what is it intended for. The contract that runtime code enforces and the card that humans read are two views of the same thing.

FOR PIXELWISE the v1 model card lives in the central `pixelwise-model` repository alongside the `.pkl` it documents. The card travels with each tagged release, so anyone who pulls v1.0 also gets the card describing what v1.0 can and cannot do:




Figure 6: An MNIST sample: a hand-written digit on a 28×28 greyscale grid, the input shape every consumer of `digit_classifier_v1` must respect.

Documentation: [scikit-learn](#), [joblib](#), [NumPy](#).

Model cards as a concept were proposed by Mitchell et al. in 2018 and are now standard at Google, HuggingFace, and increasingly across the industry.

[Model Cards \(Google\)](#), [HuggingFace Model Cards](#).

```
# MODELCARD.md: digit_classifier_v1
```

```
## Training Data
```

```
MNIST digits 1-9, ~54k train / ~9k test, public domain.  
Class 0 withheld intentionally.
```

```
## Capabilities
```

```
Predict handwritten digits 1-9.  
Expected accuracy: ~92% (LogisticRegression baseline).
```

```
## Known Failures
```

```
6/9 confusion, 3/8 confusion, messy or rotated digits.
```

```
## Intended Use
```

```
28x28 canvas drawings.  
Out of scope: photos, non-digit characters.
```

THE CARD DOES NOT REPLACE THE RUNTIME CHECKS built into the loader. Documentation tells humans what to expect; the assertion at module load tells the program what to expect. Both are needed. Documentation drifts; assertions cannot.

THE CARD IS A CONTRACT, like any other. When v2 is published, the card is updated alongside it; the version number, the training data section, and the capabilities all change in lockstep. Traceability across versions is the entire reason the card is a separate file.

Exercise: Scrutinize the Pipeline against the Card

VERIFY WHAT THE PIPELINE DOES rather than rebuild it. The pre-processing for the model is part of the artefact, not something you re-derive on the caller side; your job here is to confirm what is inside the .pkl matches what the canonical model card promises.

PIN PANDAS AS A DEPENDENCY before the first fetch_openml call: Its default parser loads dense datasets such as MNIST through pandas, so the dependency has to be installed and pinned. Install it into the virtual environment, refreeze, and commit:

```
source .venv/bin/activate
pip install pandas
pip freeze > requirements.txt
git add requirements.txt
git commit -m "Pin pandas"
```

GET AN MNIST SAMPLE ONTO YOUR DEV MACHINE. The dataset is available through scikit-learn and downloads on first access. Activate the virtual environment and start an interactive python3 session from the repository root, then run the snippet below at the »> prompt:

```
source .venv/bin/activate
python3
>>> import joblib
>>> model = joblib.load("models/digit_classifier_v1.pkl")
>>> from sklearn.datasets import fetch_openml
>>> X, y = fetch_openml(
...     "mnist_784", version=1,
...     return_X_y=True, as_frame=False,
... )
>>> sample = X[0].reshape(28, 28).astype("uint8")
>>> label = y[0]
```

You now have one ground-truth example to scrutinise the pipeline against.

OPEN THE PIPELINE AND LOOK AT ITS PARTS. Still inside the same python3 session, the Pipeline object exposes its named steps as a dictionary, and the classifier exposes the classes it knows:

```
>>> for name, step in model.named_steps.items():
...     print(name, type(step).__name__)
>>> print("classes:", model.classes_)
```

```
>>> print("n_features expected:", model.n_features_in_)
```

NOW RUN THE MODEL ON THE SAMPLE you fetched and compare against the ground-truth label:

```
>>> import numpy as np
>>> arr = (sample > 128).astype(float).reshape(1, -1)
>>> probs = model.predict_proba(arr)
>>> pred = model.classes_[probs.argmax()]
>>> print(f"pred={pred} true={label} conf={probs.max():.2f}")
```

THE POINT IS VERIFICATION. You are not reproducing the model's preprocessing; you are confirming the artefact behaves the way the model card claims. A pipeline that contains an unexpected step, or is missing one you expected, is a contract violation you want to catch here, not later.

OPEN THE CANONICAL MODEL CARD at [schutera/pixelwise-model](#) and put it side by side with your introspection output. Answer each of these against both sources:

- Which preprocessing steps are inside the pipeline, and which are the caller's responsibility? Does the boundary match what the card claims is in scope?
- Does `model.classes_` match the class roster the card promises?
- Does `n_features_in_` match the input shape the card advertises?
- Does the accuracy you can reproduce on a handful of MNIST samples land near the figure the card cites?

A mismatch between what the pipeline contains and what the canonical card promises is exactly the kind of contract violation the integration layer exists to catch. The next exercise turns this manual check into a smoke test that runs on every commit.

Designing the Inference Interface

BEFORE WRITING ANY CODE, design the contract. Three decisions propagate to every later block, the API, the database schema, the frontend, so they are worth naming up front: the class roster, the function shape, and the loading strategy.

A NAMED CLASS CONSTANT makes the valid output set explicit and independent of the model file. `CLASSES` lives in the code as a hard-coded list, not derived from `model.classes_` at import time, so the integration has its own opinion about what labels it speaks. The frontend, the database schema, and the API documentation all reference it. If the model is swapped for one trained on different classes, an assertion described below fires loudly at startup rather than returning silently wrong labels for the rest of the service's lifetime.

BATCH-FIRST DESIGN. What happens when 50 students submit a drawing in the same second? If the server processes each request individually, the model is called 50 times with a (1,784) array. If it batches them, the model is called once with a (50,784) array. sklearn's `predict_proba` is natively vectorised, so the cost is nearly identical. Design around batch; single-user is a special case.

TWO FUNCTIONS, ONE BATCH-FIRST. The public surface of `app/classifier.py` is:

```
classify_batch(images) -> list[dict]
classify(image)         -> dict
```

`classify_batch` takes an (N, 28, 28) uint8 array and returns a list of N dictionaries, each with prediction, confidence, and scores. `classify` is a single-image convenience wrapper that delegates to `classify_batch` with `image[np.newaxis]`, so there is exactly one inference path. When a request queue is added later to amortise cost, the queue assembles a batch and the same function consumes it; no signature change.

LOAD THE PIPELINE ONCE AT MODULE LEVEL. Deserialising on every request kills throughput, so `joblib.load(os.getenv("MODEL_PATH"))` runs at import time and the resulting Pipeline stays in memory for the lifetime of the process. The path comes from `.env`, not a hard-coded string, so swapping the model in production is an env-var change plus a service restart, the same pattern you will see repeated in later blocks.

THE STARTUP ASSERTION IS THE HEADLINE. Right after `joblib.load`, one line: `assert list(_pipeline.classes_) == CLASSES`. The assertion fires at startup if the model's classes disagree with the code constant. A swapped model with the wrong classes is caught the moment the service starts, not the next time a user gets a wrong digit. The next exercise has you trip it on purpose to feel the difference.

Exercise: The Classifier Module

CREATE THE CLASSIFIER MODULE FILE on the dev VM, inside the PixelWise repository. The module belongs in the app/ package, so create the directory if it is not there yet and open the file in nano:

```
cd ~/pixelwise
mkdir -p app
nano app/classifier.py
```

Paste the implementation below into the empty file, then save and exit nano with Ctrl+X, Y, Enter. It includes the class constant, module-level loading, the startup assertion, and the two functions discussed in the previous section:

```
import os
import joblib
import numpy as np
from dotenv import load_dotenv

load_dotenv()

CLASSES = ["1", "2", "3", "4", "5", "6", "7", "8", "9"]

_pipeline = joblib.load(os.getenv("MODEL_PATH"))
assert list(_pipeline.classes_) == CLASSES, \
    f"Model/CLASSES mismatch: {_pipeline.classes_}"

def classify_batch(images: np.ndarray) -> list[dict]:
    if images.ndim != 3 or images.shape[1:] != (28, 28):
        raise ValueError(
            f"Expected (N,28,28), got {images.shape}")
    arr = (images > 128).astype(float).reshape(
        len(images), -1)
    probs = _pipeline.predict_proba(arr)
    return [
        {"prediction": CLASSES[p.argmax()],
         "confidence": float(p.max()),
         "scores": dict(zip(CLASSES, p.tolist()))}
        for p in probs
    ]

def classify(image: np.ndarray) -> dict:
    return classify_batch(image[np.newaxis])[0]
```

(OPTIONAL) THE ASSERTION IS THE HEADLINE. Comment it out, save, restart Python, and import the module. Now mutate CLASSES to add a fake "10" entry, restore the assertion, restart, and observe the assertion firing. The startup error message is exactly what you want to see when a model swap goes wrong.

UPDATE `.env.example` AND `.env` at the repo root so the loader knows where to find the artefact. Open each in nano and append the line below, save, exit:

```
nano .env.example # commit this; safe to share
nano .env         # local only; never commit
```

```
MODEL_PATH=models/digit_classifier_v1.pkl
```

WRITE `predict.py` AS A SMOKE TEST at the repo root, next to `setup-server.sh`, not inside `app/`: it is a script you run, not a module the application imports. Open it in nano:

```
nano predict.py
```

Paste the smoke test below, save, exit:

```
from app.classifier import classify_batch
from sklearn.datasets import fetch_openml
import numpy as np

X, y = fetch_openml("mnist_784", version=1,
                    return_X_y=True, as_frame=False)
images = X[:5].reshape(-1, 28, 28).astype(np.uint8)
truth = y[:5]
results = classify_batch(images)
for r, t in zip(results, truth):
    print(f"Pred: {r['prediction']} "
          f"(conf {r['confidence']:.2f}) True: {t}")
```

RUN THE SMOKE TEST ON DEV from the repo root with the virtual environment activated, then commit:

```
source .venv/bin/activate
python predict.py
git add app/classifier.py predict.py .env.example
git commit -m "Add classifier module pinned to model v1.0"
git push
```

TAG THE RELEASE AS `v0.3`. The end of every block is a tagged, runnable state. Mark it now so future catch-up exercises can reset back here:

```
git tag -a v0.3 -m "End of Block 4: classifier module"
git push origin v0.3
```

`__pycache__` IN YOUR COMMIT? Running `python predict.py` produces compiled bytecode in `app/__pycache__`. If those `.pyc` files made it into the commit, extend `.gitignore` from Block 3 with the entries below so they stay out of every future commit:

```
__pycache__/
*.pyc
```

Then untrack what was already added and rewrite the previous commit so the bytecode never reaches the remote:

```
git rm -r --cached
app/__pycache__
git add .gitignore
git commit --amend --no-edit
git push --force-with-lease
```

Only `--force-with-lease` this commit if you pushed it once and nobody else has pulled yet. Otherwise create a fresh follow-up commit.

The repository is now at the state v0.3 marks: a working classifier, a dev-side smoke test, and an updated configuration contract pinned to the model release. The .pkl stays on each developer's machine, the artefact path is in .env, and the rest of the team can reproduce the environment from the code in Git.

Optional: Add the Missing 0

THE DIGIT CLASS "0" WAS INTENTIONALLY EXCLUDED FROM v1. With v1 integrated and dev-verified, you can publish your own v2. Fork the course model repository so you have a place to push:

```
gh repo fork schutera/pixelwise-model --clone --remote
cd pixelwise-model
```

The fork already contains `train.py`, `digit_classifier_v1.pkl`, and the v1 `MODEL_CARD.md` from upstream. Add "0" to the class list, include MNIST's class-0 samples in the training split, re-run `python train.py`, and save the new artefact as `digit_classifier_v2.pkl`.

NOW WRITE A v2 MODEL CARD. This is the first time card authoring is load-bearing in the course: a consumer who pulls v2 will read this file to learn what changed. Rewrite `MODEL_CARD.md` so it documents v2 specifically:

- training data section now reflects the full ten-class MNIST split,
- capabilities list all ten digits and the new accuracy you measured,
- known failures section reflects what your retrained model actually struggles with (rotated digits, 4/9 confusion, whatever the confusion matrix says),
- intended use is the same as v1 unless you changed it.

Commit artefact, script changes, and card together, and cut a fresh tag:

```
git add digit_classifier_v2.pkl train.py MODEL_CARD.md
git commit -m "v2: adds class 0"
git tag -a v2.0 -m "Adds class 0"
git push --follow-tags
```

On the PixelWise side, point `.env` at your fork and the new tag by setting `MODEL_REPO=https://github.com/<you>/pixelwise-model.git` and `MODEL_VERSION=v2.0`, then re-run `bash setup-server.sh` to pull the v2 artefact and card into `models/`. The assertion will break the moment v2 is loaded against v1's `CLASSES`; fix it by updating the constant.

A NOTE ON WHAT COMES NEXT. `app/classifier.py` works on dev, but it is trapped in a Python script.

WHY FORK INSTEAD OF BRANCHING UPSTREAM? Forking is the standard pattern for publishing a derivative of someone else's repository. You get write access to your fork, the upstream stays canonical, and a pull request is the natural way to upstream improvements later.

This challenge sets up a thread we will revisit later, when feature flags decide which model version is served and analytics compare their performance side by side. No data collection is needed; MNIST has the os already.

Self-Reflection and Recap

SELF-REFLECTION questions to guide your thoughts during and after the exercises:

- What is train/serve skew, and why is it the most common production ML defect?
- Why do we load the model once at module level instead of on every request?
- What does the startup assertion protect against, and what would happen without it?
- Why does PixelWise pin a tagged model release rather than vendoring the .pkl into the repo?
- What does the canonical model card promise, and which of its claims did you verify by introspecting the pipeline?
- How would you detect at runtime that the model is missing a class it should predict?
- Why is batch-first design the right default, even when most requests carry a single image?

RECAP of key concepts:

- A .pkl file is a build artefact, not source code; never commit it to Git.
- Train/serve skew is silent by default; the startup assertion is the loudest defence.
- Batch-first design scales from one user to N concurrent users with no code changes.
- Module-level loading keeps the model in memory for the lifetime of the process.
- The model card travels with the tagged release in pixelwise-model; integration code reads it as part of contract verification.
- `predict.py` proves end-to-end correctness before any API exists.

MILESTONE. `app/classifier.py` exposes a batch-first inference interface that scales from one student to N concurrent users with no code changes. `predict.py` proves it works end to end on real MNIST samples. The canonical `MODEL_CARD.md` in `pixelwise-model` documents what the model is and is not, and the integration verifies against it.

SECURITY THREAD. Never load a .pkl from an untrusted source; `joblib` and `pickle` deserialisation execute arbitrary code. `MODEL_PATH` comes from `.env`, never hard-coded, so swapping the model in production requires only an env-var change, not a code change. Training data and model artefacts are not committed to Git.

TEASER. The classifier is callable from Python, but who, outside your terminal, can reach it?

WE DELIBERATELY STOP SHORT OF prod. The classifier works on dev, but prod has not been updated, on purpose. There is no service on prod yet that would call `classify_batch`, no HTTP endpoint, no managed background process, no entry point at all. Pulling the new code there now would just add an unused `app/classifier.py` to the production checkout. Deployment to prod happens once there is something to deploy.

The Backend: APIs & Services

2026-06-12 · tender coconut Iltis

API FIRST. ANYONE WHO DOESN'T DO THIS, WILL BE FIRED.

The Why

Block 4 left us with `app/classifier.py`, a function that turns a 28-by-28 array of pixels into a prediction dictionary. The function works. `predict.py` proves it on real MNIST samples. What it cannot do is be reached from anywhere except the same Python process that imported it.

THIS BLOCK. is the first that exposes the application to the outside world. Every client that follows, whether a browser frontend, an automated test runner, or a monitoring probe, talks to the service through the contract you design here.

HTTP IS THE PROTOCOL that every browser, every CLI tool, and every web service in this course speaks. Figure 7 shows the path a single request takes through the stack you will build in this Block.




Figure 7:
Request and response flow for a single classification. The client sends an HTTP request to uvicorn; FastAPI dispatches it to the `/classify` handler, which calls `classify_batch` from Block 4. The dictionary returned by `classify_batch` is serialised to JSON and travels back along the same path as the HTTP response. The whole stack runs under a systemd unit that restarts it if it crashes.

A SERVICE is what `predict.py` becomes when you wrap it in HTTP and run it under a supervisor. uvicorn is the HTTP server that hosts the FastAPI application; systemd is the supervisor that keeps uvicorn running, restarts it when it crashes, and captures its logs, turning it into something that survives.

Hands On Experience

THE CONTRACT BECOMES THE PRODUCT. Once the service ¹ is running, the request and response shapes are what every other component in the system depends on. Change them silently and the frontend breaks; document them clearly and new clients can integrate without ever reading your Python. Once an API is public, it will be digested, users will build on it, and it will be hard to change ².

THIS FIFTH BLOCK introduces the backend service layer for Pixel-Wise. By the end you will have

- understood HTTP and REST well enough to design and probe an endpoint,
- written request and response contracts as Pydantic models,
- wrapped `classify_batch` in a FastAPI application,
- added API key authentication and basic rate limiting,
- run the service under `systemd`, observed crashes, and read its logs with `journalctl`.

¹ T. Winters, T. Manshreck, and H. Wright. *Software Engineering at Google: Lessons Learned from Programming Over Time*. O'Reilly Media, Sebastopol, CA, 2020. URL <https://scholar.google.com/scholar?q=Software+Engineering+at+Google+Winters+Manshreck+Wright>

² F. P. B. Jr. *The Mythical Man-Month: Essays on Software Engineering, Anniversary Edition*. Addison-Wesley, Reading, MA, 2 edition, 1995. URL <https://scholar.google.com/scholar?q=The+Mythical+Man-Month+Brooks>

THE SCRIPT-TO-SERVICE JUMP. A script runs once, prints, exits. A service runs continuously, accepts requests, returns responses, survives crashes. The work in this block is what turns the first into the second; the classifier itself does not change.

Optional Exercise: Catch Up to v0.3

CATCHING UP VIA TAG. This catch-up extends the one in Block 4. If you have not done that one, do it first: It lands you at v0.2, the end-of-Block 3 state. The exercise here picks up from there and rewinds to v0.3, the end-of-Block 4 state, with the classifier module in place, the model artefact pulled into `models/`, and `.env` pinned to the model release.

DEV ONLY. Block 4 deliberately left prod behind: The classifier had no HTTP entry point, so pulling the code there would have added an unused `app/classifier.py` to the production checkout. This catch-up does the same. Bring dev to v0.3 and leave prod where Block 4 left it; the `systemd` exercise later in this block is what finally promotes the service to prod.

THREE LAYERS TO RESTORE ON dev. Block 4 produced Git-tracked changes, the new `app/classifier.py` and `predict.py`, an updated `setup-server.sh`, and a refreshed `.env.example`. It also produced unversioned state, the model artefact `digit_classifier_v1.pkl` in `models/`, plus three new `.env` entries, `MODEL_REPO`, `MODEL_VERSION`, and `MODEL_PATH`. `git reset --hard v0.3` restores the first layer, `cp .env.example .env` seeds the third, and the extended `setup-server.sh` pulls the artefact once `.env` is in place.

RUN THE CATCH-UP ON dev. Rewind the repository to v0.3, refresh `.env` from the updated example, run `setup-server.sh` to pull system packages and the model artefact, reinstall pinned dependencies if your virtual environment was wiped, and run `predict.py` to confirm the integration works end to end. The exact commands are in the margin.

YOU ARE NOW at the state where Block 4 ended: dev is at v0.3 and prod still trails behind, ready to start Block 5.

TAG MAP. v0.3 marks the end of Block 4, v0.4 will mark the end of this block. Browse the full set at <https://github.com/schutera/pixelwise/tags>.

SEQUENCE ON dev.

```
cd ~/pixelwise
git fetch origin --tags
git reset --hard v0.3
cp .env.example .env
bash setup-server.sh
source .venv/bin/activate
pip install -r requirements.txt
python predict.py

If .venv/ is missing, run python3 -m
venv .venv before the source step.
```

ALREADY HAVE A FORK? If your own fork carries the v0.3 tag, swap origin on dev with `git remote set-url origin git@github.com:<you>/pixelwise.git` so push and pull continue to target your account. The HTTPS clone on prod stays read-only.

HTTP and REST

EVERY WEB INTERACTION IS A REQUEST AND A RESPONSE. The client sends a request that names a method, a path, and an optional body; the server sends back a status code, headers, and an optional body. The protocol carries no state of its own: each request is interpretable on its own, without knowledge of any previous request from the same client.

METHODS AND QUERIES YOU WILL USE.

- GET retrieves,
- POST submits,
- PUT replaces,
- DELETE removes.

PixelWise uses POST for /classify when the client submits a drawing, and GET for /results and /health when the client reads state.

REST PRINCIPLES layer convention on top of the protocol. Every distinct thing the API exposes is a resource with a stable URL: /classify, /results, /health. Methods do what their names suggest. Responses are JSON. Two clients written by two teams, six months apart, agree on what the API does without ever speaking, because the conventions remove ambiguity.

THE INTERACTIVE DEFINITION OF AN API is the documentation that ships with the running service. FastAPI generates an OpenAPI specification from your code, exposes the raw JSON at /openapi.json, and renders it through two built-in interfaces: Swagger UI at /docs and ReDoc at /redoc. Exposing them changes the social contract: A teammate does not need to read your source to integrate, it's your responsibility to fulfill the stated contract. Block 3's DEBUG=true flag controls whether the documentation endpoints are exposed, so the same service is self-documenting in development and silent in production.

HTTP is the request and response protocol that sits underneath every browser call, every curl invocation, and every API in this course. A reference overview of the methods, status codes, and headers lives [here](#).

STATUS CODES YOU WILL SEE MOST. Three-digit numbers grouped by leading digit. 2xx means success, 4xx means the client made a mistake, 5xx means the server made a mistake.

- 200 OK,
- 201 Created,
- 400 Bad Request,
- 401 Unauthorized,
- 404 Not Found,
- 422 Unprocessable Entity,
- 429 Too Many Requests,
- 500 Internal Server Error.

STATELESS is the load-bearing property of HTTP. The server keeps no per-client memory between requests; everything the server needs to act is in the current request, including any authentication token. This is what lets you scale.

REST is the architectural style that maps the HTTP verbs onto application resources. Each resource has a stable URL; each verb has a clear meaning. We use REST because it is the convention every frontend, every tool, and every framework in this stack expects, not because it is the only valid choice. A reference introduction to the style and its conventions lives [here](#).

HTTP status codes reference:

<https://developer.mozilla.org/en-US/docs/Web/HTTP/Status>

REST overview:

<https://restfulapi.net/>

ONE SPEC, MANY RENDERERS.

/openapi.json is the contract; Swagger UI and ReDoc are two ways to look at it. The same spec format is what Flask, Django, Express, and Spring all emit, and the same renderers and SDK generators work against any of them.

OpenAPI: <https://www.openapis.org/>

Swagger Editor: <https://editor.swagger.io/>

Exercise: HTTP, REST, and curl

PROBE A PUBLIC API before writing your own. `httpbin.org` is a free request-and-response echo service: It returns a JSON document describing exactly what it just received, so you can read status codes, headers, and the body of your own request side by side with the response. On dev, with the virtual environment active:

```
curl -i https://httpbin.org/get
curl -i -X POST https://httpbin.org/post \
  -H "Content-Type: application/json" \
  -d '{"hello": "world"}'
```

THE BODY IS JSON. JSON is the universal serialisation format for web APIs: A text-shaped tree of objects, arrays, strings, numbers, booleans, and null. Your GET response body is short enough to read in full:

```
{
  "args": {},
  "headers": {
    "Accept": "*/*",
    "Host": "httpbin.org",
    "User-Agent": "curl/7.81.0"
  },
  "origin": "203.0.113.42",
  "url": "https://httpbin.org/get"
}
```

Curly braces enclose an object, each key is a quoted string, and each value is a primitive, an array, or another object. PixelWise's `/classify` response will use the same shape, with keys `prediction`, `confidence`, and `scores`, parsed identically by any client in any language.

PICK THE VERB for each PixelWise endpoint and say what success looks like:

- `/classify`: Submit a drawing for prediction.
- `/results`: Read the last 10 predictions.
- `/health`: Check whether the service is alive.

READ EACH SECTION OF THE RESPONSE. HTTP/2 200 is the status line. The lines that follow are headers; the blank line marks the boundary; the JSON below it is the body. Every HTTP response in this course has the same shape.

JSON SYNTAX Objects `{...}`, arrays `[...]`, strings in double quotes, numbers and booleans and null unquoted, commas between items but never trailing. Spec: <https://www.json.org/>

ANSWERS. POST `/classify`: Image data does not fit in a query string; POST carries it in the body and is not cached. 200 OK with the prediction.

GET `/results`: Safely cacheable, no side effect. 200 OK.

GET `/health`: Reads liveness, no side effect. 200 OK, or 503 Service Unavailable when unhealthy.

FastAPI in Practice

FASTAPI IS A MODERN PYTHON WEB FRAMEWORK built on top of Pydantic and the ASGI standard. It generates interactive API documentation automatically, validates inputs and outputs against the models you declared, and runs on uvicorn, an asynchronous HTTP server fast enough that your bottleneck will be the model, not the framework.

PYDANTIC. Python data-validation library. A class inherits from `BaseModel` and declares typed fields; Pydantic validates incoming data against the schema and raises a structured error on mismatch. FastAPI runs the same check at the HTTP boundary and returns `422 Unprocessable Entity` on failure.

<https://docs.pydantic.dev/>

A PATH OPERATION is the unit of routing in FastAPI: A Python function whose decorator binds it to one HTTP method and one URL path. When the application starts, FastAPI walks the decorated functions and builds a routing table that uvicorn consults on every incoming request. The type hints on the parameters are not just documentation: They are read at registration time and turned into a validator that runs before your function body executes. A request that fails the validator never reaches your code. It is rejected at the boundary with a 422 response whose body lists exactly which field was wrong and why, so the caller can fix the request without guessing at your implementation.

THE DECORATOR IS THE CONTRACT. `@app.post("/classify", response_model=ClassifyResponse)` encodes three things at once: The method clients must use, the path they must hit, and the schema the response must satisfy. The first two shape the routing table. The third shapes what leaves the server. FastAPI runs the value your function returns through `ClassifyResponse` on the way out, drops fields that are not declared on the model so internal state cannot leak by accident, and fills the OpenAPI schema served at `/docs` from the same class. A handler that returns the wrong shape fails inside the framework on the first request, long before a downstream client notices that the documented contract has drifted from the implementation.

CURL IS A COMMAND-LINE HTTP CLIENT. It writes the bytes of a request directly to a socket and prints the bytes of the response,

FastAPI: <https://fastapi.tiangolo.com/>
 Pydantic: <https://docs.pydantic.dev/>
 uvicorn: <https://www.uvicorn.org/>
 ASGI: <https://asgi.readthedocs.io/>

with no browser, no JavaScript, and no implicit defaults beyond what you pass on the command line. That is what makes it useful for debugging: The request you send is exactly the request you typed, so the only place a bug can hide is in your invocation or in the server. The interactive `/docs` page, by contrast, does a lot on your behalf. It JSON-encodes the body for you, sets the `Content-Type: application/json` header on every `POST`, and, once you click `Authorize`, attaches the API key header to subsequent calls. A request that succeeds from `/docs` and fails from `curl` therefore points at something the form was adding silently, almost always a header you forgot on the command line. A request that fails from both rules out the client and points back at the handler. The exercise below gives the exact `curl` invocations against the running service so you can compare the two side by side.

THREE ENDPOINTS cover the surface this block exposes: `POST /classify` performs inference, `GET /results` is a stub that a later persistence layer will fill in, and `GET /health` reports liveness so a monitor can tell whether the service is up.

```
curl: https://curl.se/
-X sets the HTTP method.
-H sets a header.
-d sets the request body.
-i prints the response headers alongside the body, useful when a status code surprises you.
```

Exercise: Building app/main.py

INSTALL THE RUNTIME DEPENDENCIES into the dev virtual environment and pin them.

```
cd ~/pixelwise
source .venv/bin/activate
pip install fastapi uvicorn[standard] \
    slowapi pydantic
pip freeze > requirements.txt
```

WIRE THE PATH OPERATIONS FROM THE THEORY ABOVE INTO A WORKING SERVICE. From the repository root on dev, create the file:

```
cd ~/pixelwise
nano app/main.py

from fastapi import FastAPI
from pydantic import BaseModel
import numpy as np
from app.classifier import classify_batch

class ClassifyRequest(BaseModel):
    pixels: list[list[int]]

class ClassifyResponse(BaseModel):
    prediction: str
    confidence: float
    scores: dict[str, float]

app = FastAPI()

@app.get("/health")
def health():
    return {"status": "ok", "model_version": "v1"}

@app.get("/results")
def results():
    return {"results": [], "note": "persistence not yet implemented"}

@app.post("/classify", response_model=ClassifyResponse)
def classify(req: ClassifyRequest):
    arr = np.array(req.pixels, dtype=np.uint8)[np.newaxis]
    return classify_batch(arr)[0]
```

THE [standard] EXTRA on uvicorn pulls in optional but recommended packages that ship alongside the core server: httptools for a faster HTTP parser, uvloop for a faster event loop on Linux, websockets for the WebSocket protocol, and watchfiles so --reload can detect code changes during development. Without the extra you still get a working server, just slower and without auto-reload.

CTRL + R IN THE TERMINAL starts a reverse search through your shell history. Hit it, type a few characters of a command you ran earlier, such as `pip freeze`, and the most recent match appears. Press Ctrl + R again to step further back, Enter to run the match, or the right arrow to edit it before running. This is much faster than retyping or looking up commands again or scrolling with the up arrow.

RUN IT LOCALLY from the repository root, so that uvicorn can resolve `app.main:app` as a Python import against the `app/` package on disk:

```
cd ~/pixelwise
uvicorn app.main:app --reload --port 8000
```

The `--reload` flag restarts on every code change, which is what you want in development and never want in production.

TEST IT FROM THE BROWSER. Open <http://localhost:8000/docs> to see the auto-generated OpenAPI page.

TRY GET `/health`. Click “Try it out” and “Execute”. A 200 with `{"status": "ok", "model_version": "v1"}` is the first end-to-end loop. The first real POST `/classify` round-trip lands in Block 7, when the frontend sends an actual digit drawn on a canvas.




Figure 8: The Swagger UI rendered at `/docs`. FastAPI generates this page from the type hints and Pydantic models in `app/main.py`. Each endpoint expands to show its request schema, an example payload, and an interactive button that issues a live request against the running service.

OPTIONAL: TRY POST `/classify` ANYWAY. A Python list comprehension is not valid JSON, which is why the default `[[0]]` returns 422. Print a real 28-by-28 body in a Python REPL with

```
import json
print(json.dumps(
    {"pixels":
     [[0]*28 for _ in range(28)]}))
```

and paste the printed line into Swagger’s request body field.

OPTIONAL: DESIGN A FOURTH END-POINT. The three endpoints above cover inference, retrieval, and liveness. What is missing? Before reading on, sketch one more endpoint on paper. Pick the HTTP method and path, write the Pydantic request and response models, and decide what status code each error case returns. Plausible candidates include GET `/version` so a deployment can report the running build and model hash, POST `/feedback` so a frontend can flag wrong predictions for later retraining, or POST `/classify/batch` so a client can submit many images in a single round trip. Additional endpoints can help where the current contract has gaps, and to feel how naming, shape, and status codes are decisions you make before any code runs.

Authentication and Rate Limiting

Who is calling, and how often?

API KEY AUTHENTICATION is the simplest form that fits this course. The client sends a secret in a header; the server compares it to the value in `.env` and rejects requests that do not match with 401 Unauthorized. FastAPI exposes the header through a small dependency function that runs before every protected handler, which the exercise wires up.

RATE LIMITING prevents both abuse and accident. `slowapi` attaches a per-route decorator that counts requests per client IP and rejects them with HTTP 429 when the limit is exceeded. A budget like “30 requests per minute per IP” is enough to stop a runaway client without inconveniencing a normal one.

PYDANTIC, THE API KEY MIDDLEWARE, AND THE RATE LIMITER together form a three-stage filter at the boundary of the application. A request that arrives must parse as JSON, match the schema, present a valid key, and stay within the rate budget before any inference code runs. Each stage rejects with a different status code, so a client failure tells you exactly which contract was violated.

WHY A HEADER? Headers travel separately from the body and the URL, so the secret never ends up in the path component of an access log. Putting the key in a query parameter would write it into every log line the request touched on the way to the server.

WHERE DOES RATE LIMITING BELONG? Application-level limits are easy to write and travel with the code. Reverse-proxy limits, a topic we will revisit later in the course, run in front of the application and protect it even when it is overwhelmed. Both layers are useful; the latter is harder to bypass.
`slowapi`: <https://github.com/laurentS/slowapi>

(Optional) Exercise: API Key Auth and Rate Limiting

ADD THE DEFENSIVE LAYER FROM THE THEORY ABOVE TO THE SERVICE. Reopen `app/main.py` from `~/pixelwise` on dev:

```
cd ~/pixelwise
nano app/main.py
```

Add the API key dependency alongside the existing imports:

```
from fastapi import Header, HTTPException, Depends
import os

def verify_api_key(x_api_key: str = Header(...)):
    if x_api_key != os.getenv("SECRET_API_KEY"):
        raise HTTPException(
            status_code=401, detail="Invalid API key")
```

Then attach `verify_api_key` to the `/classify` route by extending its existing decorator with `dependencies=[...]`; the handler body itself does not change:

```
@app.post("/classify",
           response_model=ClassifyResponse,
           dependencies=[Depends(verify_api_key)])
def classify(req: ClassifyRequest):
    arr = np.array(req.pixels,
                   dtype=np.uint8)[np.newaxis]
    return classify_batch(arr)[0]
```

ADD RATE LIMITING with `slowapi`. Add the imports and the limiter setup near the top of the file:

```
from fastapi import Request
from slowapi import Limiter
from slowapi.util import get_remote_address
from slowapi.middleware import SlowAPIMiddleware

limiter = Limiter(key_func=get_remote_address)
app.state.limiter = limiter
app.add_middleware(SlowAPIMiddleware)
```

Then stack the limiter decorator above the route decorator and add a request: `Request` parameter so `slowapi` can identify the caller:

THE HEADER TRAVELS IN EVERY REQUEST. `X-API-Key` is a custom header name; FastAPI's `Header(...)` maps it to the function parameter automatically. A missing header returns 422 from Pydantic, a wrong value returns 401 from your handler.

```

@app.post("/classify",
         response_model=ClassifyResponse,
         dependencies=[Depends(verify_api_key)])
@limiter.limit("30/minute")
def classify(request: Request,
            req: ClassifyRequest):
    arr = np.array(req.pixels,
                  dtype=np.uint8)[np.newaxis]
    return classify_batch(arr)[0]

```

TEST POST /classify FROM SWAGGER. Reload /docs, expand the endpoint, click “Try it out”, and fill the key from your .env into the x_api_key header field that now sits alongside the request body. Press “Execute”. Without a key the classify request returns 401; with the right key the same request goes through.

OPTIONAL: ADD A GLOBAL AUTHORIZE BUTTON. Swap Header(...) for APIKeyHeader from fastapi.security to register the API key as an OpenAPI security scheme. Swagger then renders a global “Authorize” button at the top of /docs; after clicking it once, every Try-it-out includes the header automatically.

OPTIONAL: WATCH A 429 HAPPEN. A small Bash loop firing 40 curl requests at the endpoint inside the same minute should see the last ten return 429 Too Many Requests. The remaining requests stay rejected until the window rolls over.

From Dev to Production: systemd

uvicorn RUNS FINE, but one Ctrl+C and it is gone. A closed terminal kills it. A reboot kills it. A panic in a worker kills it without restart. We got the interface right, but none of the above is the behaviour of a robust service.

systemd is the OS-level supervisor that comes with Ubuntu. It starts services on boot, restarts them when they crash, captures their output to a structured log, and exposes a uniform interface for operators. The contract between you and systemd is a unit file, an INI-style description with three sections: An [Unit] block for metadata and dependencies, a [Service] block naming the user, working directory, environment file, command to run, and restart policy, and an [Install] block declaring when the unit should be started. The exercise below has the full file PixelWise needs.

systemctl is the operator's interface to systemd. start, stop, and restart are the verbs you reach for during deployment; status reports the current state and the last few lines of output; enable marks the service as one that should start on boot; daemon-reload tells systemd to re-read its unit files after you edit one. The exercise runs the full sequence on prod.

journalctl is where you go first when something breaks. journalctl -u pixelwise prints every log line from your service since boot; -f follows new lines as they arrive, the way tail -f follows a file; --since "5 min ago" narrows the window when the bug is recent.

WE USE A FRACTION of what systemd can do. It also manages timers (cron with dependency awareness), socket activation (start on connection), resource limits (cap CPU/memory per service), and dependency ordering (start the database before the app). It is the init system of virtually every modern Linux server.

systemd: <https://www.freedesktop.org/wiki/Software/systemd/>

gunicorn is what production Python deployments put in front of uvicorn to spawn multiple worker processes, giving concurrency and crash isolation. For our classroom setup, a single uvicorn worker behind systemd is sufficient; students who deploy beyond this course will encounter gunicorn within the first week.

Exercise: Running as a systemd Service

PROMOTE THE SERVICE FROM A FOREGROUND uvicorn TO A SUPERVISED systemd UNIT. Author the unit file on dev alongside the app code first, then refer to the marginnote for exact commands.

AUTHOR THE UNIT FILE on dev at `deploy/pixelwise.service` in the repository. Keeping it under version control means anyone can rebuild prod from a clean VM by cloning, pulling, and copying, instead of remembering what was typed into nano six months ago.

```
# deploy/pixelwise.service
[Unit]
Description=PixelWise API
After=network.target

[Service]
User=produser
WorkingDirectory=/opt/pixelwise
EnvironmentFile=/opt/pixelwise/.env
ExecStart=/opt/pixelwise/.venv/bin/uvicorn \
    app.main:app --host 0.0.0.0 --port 8000
Restart=always
RestartSec=3

[Install]
WantedBy=multi-user.target
```

TEACH `setup-server.sh` TO INSTALL THE UNIT. Open the file in nano on dev and append a small block at the bottom so any future catch-up or fresh provision picks it up:

```
nano setup-server.sh
```

```
# Install, start, and report the systemd unit on prod
if [ -f deploy/pixelwise.service ] && \
    command -v systemctl >/dev/null 2>&1 && \
    id produser >/dev/null 2>&1; then
    sudo cp deploy/pixelwise.service \
        /etc/systemd/system/pixelwise.service
    sudo systemctl daemon-reload
    sudo systemctl enable pixelwise
    sudo systemctl restart pixelwise
```

PRODUCTION HAS SOME CATCHING UP TO DO. FIRST DEV PUSH THEN PROD PULL.

```
git add deploy/pixelwise.service
app/main.py requirements.txt
.env.example

git commit -m "Add systemd
service file and FastAPI service
with auth and limits"

git push

ssh produser@192.168.56.11
cd /opt/pixelwise
git pull
source .venv/bin/activate
pip install -r requirements.txt
```

WHY SHIP THE UNIT FILE IN THE REPO? The unit file is part of how the app runs, so it belongs next to the code. Diffs are reviewable, deploys are reproducible, and there are no secrets inside, since `EnvironmentFile` only points at `/opt/pixelwise/.env`, which stays on prod.

PROD-ONLY BY DESIGN. The block runs only when `produser` exists, which marks the prod VM. On dev the user is missing, so `setup-server.sh` skips the whole step instead of copying a unit that references a user who does not exist there. On prod, the script installs, enables, restarts, and reports the service status on every run, picking up any change to `deploy/pixelwise.service` from the latest git pull. The `--no-pager` flag keeps the output inline so the script does not stop.


```
sudo systemctl status pixelwise --no-pager
fi
```

INSTALL THE UNIT ON PROD after the pull. Run `setup-server.sh`; it copies the unit, reloads systemd, enables the service, restarts it, and prints the status block in one shot:

```
bash setup-server.sh
```

KILL IT ON PURPOSE. A service that has never crashed is a service whose recovery has never been tested.

```
sudo systemctl stop pixelwise
sudo systemctl status pixelwise --no-pager
sudo systemctl start pixelwise
journalctl -u pixelwise --since "5 min ago"
```

READ THE STATUS OUTPUT. Active: active (running) is what you want. Active: failed means the unit started and exited; the last few log lines below the status block usually say why. Active: inactive (dead) means the unit was never started.

FIRST END-TO-END LOOP ON THE NETWORK. From dev, hit the unprotected health endpoint on prod:

```
curl http://192.168.56.11:8000/health
```

A response of `{"status": "ok", "model_version": "v1"}` comes back from the systemd-supervised uvicorn on prod. No body, no key, just a cross-machine round-trip.

THE SYSTEM IS REACHABLE ACROSS MACHINES ON A SHARED NETWORK FOR THE FIRST TIME.

Self-Reflection and Recap

SELF-REFLECTION questions to guide your thoughts during and after the exercises:

- What is the difference between running `uvicorn` manually and running it under `systemd`?
- Why does Pydantic validation matter for security as well as correctness?
- What happens when the `systemd` service crashes, and how does `Restart=always` help?
- Why is `GET /results` a stub right now, and what would have to arrive to fill it in?
- How would you debug a failing endpoint using only `journalctl`?
- Which of the three rejection points (Pydantic, API key, rate limiter) is hardest to bypass, and why?

RECAP of key concepts:

- HTTP is a stateless request-response protocol; REST layers conventions on top so APIs are predictable across teams.
- Pydantic models are the API contract; FastAPI enforces them at the boundary and renders them as live documentation.
- `POST /classify`, `GET /results`, `GET /health` cover the surface this block exposes; `/results` is a deliberate stub awaiting a persistence layer.
- API keys answer “who is calling”; rate limits answer “how often”.
- `systemd` turns a script into a service: starts on boot, restarts on crash, logs to `journalctl`.

MILESTONE. PixelWise runs as a managed `systemd` service on the prod VM, listening on `http://192.168.56.11:8000`. The endpoints answer requests from dev and from your host machine over the host-only adapter; the machines outside the network cannot reach the service yet. `POST /classify` returns live digit predictions, `GET /results` is stubbed pending persistence, and `GET /health` reports liveness. You have killed and restarted the backend at least once and read the logs with `journalctl`. The contract is now the load-bearing surface of your backend: Every component that follows will speak through it.

SECURITY THREAD. The API key lives in `.env`, never in code. Pydantic models are the input validation boundary; rate limits cap any single caller; the `X-API-Key` header keeps the secret out of access logs. Transport-layer encryption is a thread we will revisit later in the course.

TEASER. Where do all those predictions go once the response has been sent?

WHAT IS MISSING IS MEMORY. The service answers each request and forgets it the moment the response goes out. That is fine for a smoke test and broken for everything else: A user looking back at their last ten drawings has nothing to show, an operator investigating a wrong prediction cannot see what was sent, and a retraining loop cannot use real-world inputs because none of them survived. Block 6 puts PostgreSQL behind the service so every classification becomes a row, GET /results starts returning real history, and the service grows the audit trail the rest of the course depends on.

Databases & SQL


2026-06-12 · tender grape Hase

IN COMPARISON, DUMBLEDORE'S PENSIEVE IS A JOKE.

The Why

We have a service that classifies, but it does not remember anything. The moment the response leaves the server, the prediction is gone. This makes it hard to track the model's performance over time, compare versions, or build any type of features and functionality that revolve around past predictions.

WE NEED TO WIRE THE API ONTO A DATABASE. This block is the first that adds persistent state. Every later block depends on the database schema you design today.



WHY PERSIST AT ALL? A response that vanishes the moment it is sent is a response that cannot be counted, compared, or queried. A database makes the prediction permanent: Every drawing becomes a row, the class label becomes a column, the confidence becomes another, and the model version becomes a third.

BEYOND THIS COURSE, the same table is where production teams build retraining datasets, let users correct wrong predictions for labelled data, and detect distribution drift. We do not do any of those here, but it is trivial to see how every one of them starts with having the database.

Figure 9: The persistence flow. POST /classify computes a prediction with `classify_batch`, the handler maps the result through the Prediction SQLAlchemy model, and a row lands in the predictions table inside the `predictions.db` SQLite file before the response goes out.

Hands On Experience

CONSIDER THIS SCENARIO: A teammate asks how confident the model has been on average across the last hundred predictions. You realise nobody can answer that question. The API returned the confidence to every caller, but nothing recorded it. The only way back to those numbers is to ask the users to draw their digits again, which is not an option. In an increasingly data-driven world, data defines your product, your moat, and your competitive advantage. Not having the data is like throwing away your most valuable asset, and until someone comes up with a solution for time-travel, it's a non-recoverable loss.

THE SCHEMA IS A CONTRACT, the same way the API contract is a contract. Get it right and every later question, distribution per digit, average confidence per model version, error rate over time, becomes a query. Get it wrong and you find out later, with a hundred thousand rows already written, that the column you need does not exist. Schemas evolve through migrations, which we cover at the end of the block as an optional path.

THIS SIXTH BLOCK introduces persistence for PixelWise. By the end you will have

- written your first SQL queries against a SQLite database,
- modelled the predictions table as a SQLAlchemy class,
- wired POST /classify so every prediction becomes a row,
- replaced the GET /results stub with a real query,
- configured database credentials through .env,
- and, optionally, run your first schema migration with Alembic.

DESIGNING SCHEMAS is its own discipline. We design the predictions table once and keep it that way for the mainline of the course. The optional Alembic exercise at the end lets you feel the cycle of design, deploy, regret, migrate on a real database.

Optional Exercise: Catch Up to v0.4

CATCHING UP VIA TAG. This catch-up extends the one in Block 5. If you have not done that one, do it first: It lands you at v0.3, the end-of-Block 4 state. The exercise here picks up from there and rewinds to v0.4, the end-of-Block 5 state, with the FastAPI service exposing `POST /classify` and `GET /health`, API key authentication, basic rate limiting, and a systemd-supervised uvicorn running on prod.

BOTH VMs THIS TIME. Block 5 ended with the service actually running on prod, so unlike the previous catch-up, this one brings both machines up to v0.4. The Git-tracked changes since v0.3 include `app/main.py`, `deploy/pixelwise.service`, an extended `setup-server.sh` that installs the systemd unit on prod, an updated `requirements.txt` adding `fastapi`, `uvicorn`, `slowapi`, and `pydantic`, and a refreshed `.env.example` with `SECRET_API_KEY`. The local `.env` on prod needs the new key written by hand; `setup-server.sh` then pulls the model, copies the unit, enables, and restarts the service in one shot.

RUN THE CATCH-UP ON dev. Rewind the repository to v0.4, refresh `.env` from the updated example and edit it to set `SECRET_API_KEY`, re-install pinned dependencies if your virtual environment was wiped, then run `setup-server.sh` to pull the model artefact. The exact commands are in the margin.

RUN THE CATCH-UP ON prod. Same sequence on prod, where `setup-server.sh` also copies the systemd unit into place and restarts the service. After the script finishes, verify the API is reachable from dev.

YOU ARE READY PLAYER ONE.

TAG MAP. v0.4 marks the end of Block 5, v0.5 will mark the end of this block. Browse the full set at <https://github.com/schutera/pixelwise/tags>.

SEQUENCE ON dev.

```
cd ~/pixelwise
git fetch origin --tags
git reset --hard v0.4
cp .env.example .env
nano .env
source .venv/bin/activate
pip install -r requirements.txt
bash setup-server.sh
```

Edit `.env` to set `SECRET_API_KEY` to any string of your choosing; `openssl rand -hex 32` prints a clean random value if you want one. If you already have a `.env` from earlier work, skip the `cp` and just edit the file in place. If `.venv/` is missing, run `python3 -m venv .venv` before the source step.

SEQUENCE ON prod.

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git fetch origin --tags
git reset --hard v0.4
cp .env.example .env
nano .env
source .venv/bin/activate
pip install -r requirements.txt
bash setup-server.sh
```

Use the same `SECRET_API_KEY` value as on dev. If you already have a `.env` on prod, skip the `cp` and edit the file in place. Then from dev: `curl http://192.168.56.11:8000/health` should return `{"status": "ok" ...}`.

Relational Databases and SQL

A RELATIONAL DATABASE STORES DATA IN TABLES. Each table has a fixed schema: a list of named columns, each with a type. Each row is a single record. The schema is the contract that every record must follow; the database refuses inserts or queries that violate it.

STRUCTURED QUERY LANGUAGE (SQL) is the language you use to talk to a relational database. The database is a single file, there is no server to install, no users, no network configuration. For a single-developer prototype or a desktop application, this is exactly what you want. Four operations cover almost everything:

```
-- Insert a row
INSERT INTO predictions (prediction, confidence,
                        model_version)
VALUES ('7', 0.94, 'v1');

-- Read rows
SELECT prediction, confidence FROM predictions
    WHERE model_version = 'v1'
    ORDER BY created_at DESC
    LIMIT 10;

-- Delete a row
DELETE FROM predictions WHERE id = 42;
```

PostgreSQL for Production

POSTGRESQL IS A PRODUCTION DATABASE. It runs as a separate process, accepts connections over the network, has real users with real permissions, and supports concurrent reads and writes without serialising the whole file. `psql` is the command line client.

WHY A DEDICATED USER? LEAST PRIVILEGE The postgres superuser can drop any database, create any user, and read any table. A dedicated pixelwise user can only do what its privileges allow, which by default is read and write its own tables. The principle is the same as the non-root produser on your server. You provision a pixelwise role and a pixelwise database owned by that role inside PostgreSQL, then verify the credentials with `psql` before any Python is involved.

VOCABULARY TO KEEP STRAIGHT. A table is a named set of rows; a row is one record; a column is a named field with a type; a primary key is the column or columns that uniquely identify each row. Most tables use an auto-incrementing integer or a UUID as their key.

id	prediction	confidence	model_version
1	7	0.94	v1
2	3	0.81	v1
3	5	0.99	v1

WHAT YOU WILL MEET LATER.

- JOIN combines rows from multiple tables.
- UPDATE adjusts existing values without changing the schema.
- GROUP BY, COUNT, and AVG aggregate.
- ORDER BY and LIMIT shape the output.
- CREATE INDEX speeds up queries that scan large tables.

The full treatment belongs in a dedicated database course; the four operations above are enough to get PixelWise running.

WHERE SQLITE STOPS SHORT? SQLite has no user authentication, no password protection, and limited concurrency - due to file locking.

PostgreSQL: <https://www.postgresql.org/>
 SQLite: <https://www.sqlite.org/>

Exercise: SQLite on the Server

WHILE YOU COULD SET THINGS UP MANUALLY ON prod, it is a best practice to automate every step of the server configuration through `setup-server.sh`; you edit it on dev, commit, push, and `bash setup-server.sh` carries the work on both VMs. The one manual step is writing the secret `DB_PASSWORD` into the local `.env` on each VM, since secrets never travel through git. Provisioning on dev too means you can write `models.py`, run `init_db.py`, and inspect rows with `psql` without ever ssh'ing to prod.

EXTEND `setup-server.sh` ON dev. Add `postgresql` to the `apt install` line, then append a stanza that provisions the pixelwise role and database wherever `.env` is present:

```
cd ~/pixelwise
nano setup-server.sh
```

RESOLVE THE SCRIPT DIRECTORY AT THE TOP. Add this line to `setup-server.sh` right after the shebang, so every later stanza can locate the repo root regardless of the caller's working directory:

```
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" \
&& pwd)"
```

APPEND THE PROVISIONING STANZA after the `systemd` block:

```
# Provision the pixelwise role and database on every VM
if command -v psql >/dev/null 2>&1 && \
[ -f "$SCRIPT_DIR/.env" ]; then
set -a; source "$SCRIPT_DIR/.env"; set +a
sudo -u postgres psql -tAc \
    "SELECT 1 FROM pg_roles WHERE rolname='pixelwise'" \
    | grep -q 1 || \
sudo -u postgres psql -c \
    "CREATE USER pixelwise \
    WITH PASSWORD '$DB_PASSWORD';"
sudo -u postgres psql -tAc \
    "SELECT 1 FROM pg_database WHERE datname='pixelwise'" \
    | grep -q 1 || \
sudo -u postgres createdb -O pixelwise pixelwise
fi
```

If `psql` is missing after the script runs, the provisioning stanza fails. The `postgresql` meta-package normally pulls `postgresql-client-common` in transitively; on older Ubuntu releases, minimal cloud images, or when the `apt` cache is stale, that chain breaks and the package has to be installed by hand:

```
sudo apt install
postgresql-client-common
bash setup-server.sh
```

THE CONNECTION STRING is a single URL:

```
sqlite:///absolute/path/to/file.db.
```

Note the four slashes after the colon: three for the URL scheme, one for the absolute path. A relative path uses three slashes (`sqlite:///./pixelwise.db`). The application reads this from `.env`, so the same code works on dev and prod with different paths.

IDEMPOTENT BY DESIGN. Each `psql -tAc` probe asks PostgreSQL whether the role or database already exists; the `||` short-circuits the create when the answer is yes. Run the script ten times, the role and database are still there exactly once. `SCRIPT_DIR` resolves to whichever checkout invoked the script, so the same stanza runs on dev and prod; the `systemd` block from Block 5 stays gated behind `id produser` since the service only ever runs on prod.

DOCUMENT THE NEW SECRET in `.env.example` on dev, then commit and push:

```
echo "DB_PASSWORD=" >> .env.example
git add setup-server.sh .env.example
git commit -m "Provision PostgreSQL"
git push
```

APPLY ON dev FIRST. Write the password into your local `.env` and run the script:

```
echo "DB_PASSWORD=database" >> .env
bash setup-server.sh
```

REPEAT ON prod. Same shape, plus the `git pull` to bring in the committed script:

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git pull
echo "DB_PASSWORD=database" >> .env
bash setup-server.sh
```

The first run on each VM installs `postgresql` and provisions the `pixelwise` role and database; subsequent runs change nothing because the provisioning stanza is idempotent.

VERIFY THE CONNECTION on prod with the application user:

```
psql -U pixelwise -d pixelwise -h localhost
```

`pixelwise=>` prompt means the role and database are wired up. If `psql` fails to connect, fix the credentials before going any further. A working `psql` is the foundation of everything the rest of the block builds on. If it worked you know

DID IT TAKE ON dev?

```
psql -U pixelwise -d pixelwise -h localhost
```

Enter the password from `.env`. A `pixelwise=>` prompt means the role exists, the database is yours, and the password matches. If the connection is refused, re-read the script output before moving on to prod.

WIPE AND START FRESH. The provisioning stanza skips creation when the role or database already exists, so editing `DB_PASSWORD` after a first run does not propagate. If `psql` rejects the password or reports a missing database, drop both and re-run:

```
sudo -u postgres dropdb pixelwise
sudo -u postgres dropuser pixelwise
bash setup-server.sh
```

WHAT EACH FLAG DOES.

- `-U pixelwise` selects the role to log in as.
- `-d pixelwise` selects the database to connect to.
- `-h localhost` forces a TCP connection to the local machine, which makes `psql` prompt for the password from `.env` instead of trusting the Unix socket's peer authentication on the `postgres` system user.

ROLE AND DATABASE SHARE A NAME. The setup script creates both:

- A role `pixelwise` with a password.
- A database `pixelwise` owned by that role.

PostgreSQL keeps roles and databases as independent objects, so the matching names are a readability convention, not a requirement; the application could just as well live in a database called `predictions` owned by a role called `api`.

SQLAlchemy: Tables as Python Classes

SQLALCHEMY IS A PYTHON LIBRARY that maps database tables to Python classes. You define a Prediction class with fields like prediction, model_version, and created_at; SQLAlchemy generates the SQL to create the table, insert rows, and query data. This is called an ORM, an object-relational mapper, and is standard in Python web development.

PYTHON IS ERGONOMICS. Two reasons, why choosing to use an ORM: Writing Python is faster than writing string SQL, and the editor knows the field names. The second is safety: SQLAlchemy uses parameterised queries by default, so user input never ends up concatenated into a SQL statement. SQL injection becomes a footgun the framework refuses to hand you.

AN ENGINE AND A SESSION are the runtime objects that connect the model to the database. create_engine takes the connection string models.py assembles from DB_PASSWORD and opens a pool of connections; sessionmaker returns a SessionLocal factory the handlers call once per request; Base.metadata.create_all(engine) creates the table on first run.

SQLAlchemy:

<https://www.sqlalchemy.org/>

SQLAlchemy Tutorial:

<https://docs.sqlalchemy.org/en/20/tutorial/>

THE CONNECTION

STRING is a single URL:

postgresql://user:pass@host/dbname.
It encodes everything the driver needs to reach the database. models.py assembles it from the DB_PASSWORD secret in .env, so the secret stays in one place and a leaked connection string is still a leaked database.

Exercise: The Prediction Model

INSTALL THE DRIVER AND THE ORM ON dev into the virtual environment, then pin the new packages to requirements.txt so prod picks them up on the next deploy.

WRITE app/models.py on dev with both the schema for one prediction and the session factory the handlers will use: The class label, the confidence, and the model version, plus an engine that assembles its connection string from DB_PASSWORD and a SessionLocal bound to it. DB_PASSWORD is already in .env from the provisioning exercise, and create_engine only opens a connection when a query runs, so defining the engine here is safe:

```
from sqlalchemy import (Column, Integer, String,
                        Float, DateTime, create_engine)
from sqlalchemy.orm import declarative_base, sessionmaker
from datetime import datetime
import os
from dotenv import load_dotenv
```

```
load_dotenv()
```

```
Base = declarative_base()
```

```
class Prediction(Base):
    __tablename__ = "predictions"
    id = Column(Integer, primary_key=True)
    prediction = Column(String, nullable=False)
    confidence = Column(Float, nullable=False)
    model_version = Column(String, nullable=False)
    created_at = Column(DateTime,
                        default=datetime.utcnow)
```

```
DB_PASSWORD = os.getenv("DB_PASSWORD")
DATABASE_URL = (
    "postgresql://pixelwise:"
    f"{DB_PASSWORD}@localhost/pixelwise")
engine = create_engine(DATABASE_URL)
SessionLocal = sessionmaker(bind=engine)
```

NO DATABASE_URL LINE GOES IN .env. The engine assembles its connection string from DB_PASSWORD, which you already wrote into

INSTALL.

```
cd ~/pixelwise
source .venv/bin/activate
pip install sqlalchemy
psycopg2-binary
pip freeze > requirements.txt
```

THREE COLUMNS PLUS BOOKKEEPING. prediction, confidence, and model_version carry the answer; id and created_at are bookkeeping the database fills in for you. Every later question, distribution per digit, average confidence per model version, error rate over time, lives in this table.

ONE ENGINE, ONE FACTORY.

create_engine holds the connection pool and sessionmaker turns it into SessionLocal, the callable each handler invokes to open a database session. The engine is lazy: It connects only when a session first runs a query, so importing this module never needs a live database.

THREE PARSERS, ONE FILE. .env is read by bash in setup-server.sh, by python-dotenv in the app, and by systemd through its EnvironmentFile

.env during the provisioning exercise, so the password stays in exactly one place. Writing a DATABASE_URL that interpolates \${DB_PASSWORD} inside .env looks tidier but is a trap: Bash and python-dotenv expand the reference, yet systemd reads its EnvironmentFile without expanding anything, so the service would try to authenticate with the literal variable text and Postgres would reject it. Assembling the URL in models.py keeps every reader of .env in agreement, no matter who starts the process.

CREATE THE TABLE ONCE from a small initialisation script saved as init_db.py at the repository root on dev. The script imports the engine that app/models.py already configured and calls Base.metadata.create_all to materialise every model class into a table:

```
from app.models import Base, engine

Base.metadata.create_all(engine)
```

COMMIT AND DEPLOY. Stage the new and changed files on dev, push, and on each VM install the pinned dependencies and run python init_db.py to create the schema. Then confirm with psql that the predictions table is listed:

```
psql -U pixelwise -d pixelwise -h localhost -c "\dt"
```

INITIALISE THE SCHEMA. by running
python init_db.py on dev.

COMMIT AND DEPLOY.

```
On dev:
git add app/models.py init_db.py
requirements.txt .env.example
git commit -m "Add Prediction
model"
git push
source .venv/bin/activate
python init_db.py
On prod:
cd /opt/pixelwise && git pull
source .venv/bin/activate
pip install -r requirements.txt
python init_db.py
```

OPTIONAL: ONE-SHOT DEPLOY. Make bash setup-server.sh also install pinned dependencies and create the schema by appending two stanzas after the Postgres provisioning block:

```
if [ -d "$SCRIPT_DIR/.venv" ] &&
[ -f
"$SCRIPT_DIR/requirements.txt"
]; then
source
"$SCRIPT_DIR/.venv/bin/activate"
pip install -r
"$SCRIPT_DIR/requirements.txt"
fi
if [ -f "$SCRIPT_DIR/init_db.py"
]; then
(cd "$SCRIPT_DIR" && python
init_db.py)
fi
```

Exercise: Wiring the API to the Database

THE DATABASE IS READY. Let's make the API populate it with real data.

EDIT `app/main.py` ON dev so `POST /classify` writes a row before returning the response, and `GET /results` returns real rows instead of the current scaffold. The change lands in three pieces: Two new imports at the top of the file, an addition to the classify handler, and a real body for results.

```
from app.models import Prediction, SessionLocal

@app.post("/classify",
          response_model=ClassifyResponse,
          dependencies=[Depends(verify_api_key)])
@limiter.limit("30/minute")
def classify(request: Request,
            req: ClassifyRequest):
    arr = np.array(req.pixels,
                   dtype=np.uint8)[np.newaxis]
    result = classify_batch(arr)[0]

    db = SessionLocal()
    db.add(Prediction(
        prediction=result["prediction"],
        confidence=result["confidence"],
        model_version="v1"))
    db.commit()
    db.close()

    return result

@app.get("/results")
def results():
    db = SessionLocal()
    rows = (db.query(Prediction)
            .order_by(Prediction.created_at.desc())
            .limit(20).all())
    db.close()
    return {"results": [
        {"id": r.id,
         "prediction": r.prediction,
```

TWO NEW IMPORTS. `Prediction` is the row class and `SessionLocal` the session factory, both defined in `app/models.py`. Add the line at the top of `main.py` next to the existing imports.

CLASSIFY NOW WRITES A ROW. The decorators and signature from Block 5 stay untouched. The body picks up one new intermediate, `result`, so the same dict can be persisted and returned, then a five-line session block that opens a `SessionLocal`, adds the new `Prediction`, commits, and closes.

A session is the unit of database work: Open it at the start of the handler, close it at the end. FastAPI documents a more elegant variant through `Depends`; the manual version here is the simplest one that shows what is happening. The returned result dict is unchanged, so the response body stays identical to the contract we defined in Block 5.

RESULTS RUNS ONE QUERY. It pulls the twenty newest rows from predictions, ordered by `created_at` descending, and shapes each row into the JSON object the frontend will expect in the next block. The Block 5 stub that returned an empty list is gone.

```

    "confidence": r.confidence,
    "model_version": r.model_version,
    "created_at": r.created_at.isoformat()}
for r in rows]}

```

RESTART THE SERVICE. On dev the running uvicorn --reload picks up the new code on save. On prod restart through systemd so the unit re-reads app/main.py:

```
sudo systemctl restart pixelwise # prod only
```

SMOKE TEST THE WIRING with one unauthenticated GET /results call. An empty table is exactly what you want here: A successful response with an empty list proves the engine opened, the session works, the query ran, and the JSON shape is right. Any of those broken and the call returns a 500, not an empty list:

```
# on dev, against the local uvicorn
curl -s http://localhost:8000/results \
  | python3 -m json.tool
```

```
# on dev, against prod via the API gateway
curl -s http://192.168.56.11:8000/results \
  | python3 -m json.tool
```

COMMIT AND TAG v0.5.

```

git add app/main.py app/models.py
      init_db.py requirements.txt
      .env.example
git commit -m "Add PostgreSQL
persistence"
git tag v0.5
git push --follow-tags

```

Stage on dev, commit with a message that names the change, tag the result, and push both the commit and the tag in one shot.

OPTIONAL: FEED IT PAYLOAD. So far our flow has never been hit by real digit drawings, and on that note testing really came short. Build a 28-by-28 zero payload once, perhaps wrapped in a probe_init.py, then use it to test the API and the database.

Optional Exercise: First Migration with Alembic

SCHEMAS ALWAYS CHANGE. You design the table, deploy it, store 500 predictions. Then you realise you forgot a column. You cannot drop the table and recreate it; those 500 rows are real data. You need a tool that alters the live schema without losing them. This exercise is optional: The mainline of the course does not depend on it, but running it once gives you the feel of the migration cycle on a real database. The example below adds a `client_ip` column for debugging, but any new column works.

INITIALISE ALEMBIC ON DEV. Alembic is the migration tool for SQLAlchemy; the cycle is `init`, `revision --autogenerate`, `upgrade head`, run in that order. All file edits in this exercise happen on dev. `alembic init` writes its scaffold into the current directory, so run everything from the repository root:

```
cd ~/pixelwise
source .venv/bin/activate
pip install alembic
pip freeze > requirements.txt
alembic init alembic
```

The `pip freeze` step captures the new package in `requirements.txt` so prod picks it up on the next deploy.

Edit `alembic/env.py` to import `Base` and `engine` from `app.models` so `autogenerate` sees your tables and runs against the same connection string the application builds, rather than the `sqlalchemy.url` placeholder in `alembic.ini`.

ADD A NEW COLUMN to the model. In `~/pixelwise/app/models.py` on dev, add `client_ip = Column(String)` between the `confidence` and `model_version` lines.

GENERATE THE MIGRATION from the repository root on dev:

```
cd ~/pixelwise
alembic revision --autogenerate \
    -m "add client_ip column"
```

WIRE `client_ip` INTO THE HANDLER in `~/pixelwise/app/main.py` on dev. `Request` is already imported and `request: Request` is already a parameter on `classify` from the Block 5 rate limiter, so the only change is one keyword argument in the `Prediction(...)` call:

Alembic: <https://alembic.sqlalchemy.org/>

THE METAPHOR. A migration is a recipe that describes a change to the schema and how to roll it back. The migration tool keeps an internal log of which migrations have run, applies new ones in order, and refuses to run the same one twice. The result is that the schema in the database always matches the migrations folder in Git.

READ THE GENERATED CONFIG. `alembic init alembic` writes a folder of starter files. Skim them once before you change anything; the structure is small and the mental model becomes obvious by reading the templates.

READ THE MIGRATION BEFORE YOU RUN IT. The generated script lives in `alembic/versions/`. Open it, confirm the only operation is `add_column("client_ip")`, and only then run `upgrade head`. A wrong migration on a populated table is much harder to undo than to prevent.


```
db.add(Prediction(
    prediction=result["prediction"],
    confidence=result["confidence"],
    client_ip=request.client.host,
    model_version="v1"))
```

COMMIT AND DEPLOY. On prod, order matters: Install the new pinned dependency and apply the migration first, then run `setup-server.sh`. `setup-server.sh` restarts the service with code that references the new column, so if the migration has not run yet, the handler will try to INSERT a column that does not exist in the database. From dev, stage the changes and push:

```
cd ~/pixelwise
git add app/main.py app/models.py alembic/ \
    alembic.ini requirements.txt
git commit -m "Add client_ip column via Alembic"
git push
```

Then on prod, pull, install the new dependency, run the migration, and only then restart the service:

```
cd /opt/pixelwise
git pull
source .venv/bin/activate
pip install -r requirements.txt
alembic upgrade head
bash setup-server.sh
```

VERIFY ON prod WITH `psql`:

```
psql -U pixelwise -d pixelwise -h localhost \
    -c "SELECT prediction, client_ip
        FROM predictions
        ORDER BY created_at DESC LIMIT 5;"
```

Old rows show NULL for `client_ip`; new predictions store the caller's address now that the field is wired in. That gap is the visible signature of a migration: The column exists everywhere, but its values populate only from the moment the new code shipped.

Self-Reflection and Recap

SELF-REFLECTION questions to guide your thoughts during and after the exercises:

- Why do we store predictions in a database rather than just returning them in the API response?
- What is the difference between dropping and recreating a table versus running a migration?
- If you ran the optional Alembic exercise, why do old rows show NULL for the new column while later rows carry a value?
- What happens if you store the database password directly in your Python code instead of in `.env`?
- Why is a dedicated database user with minimal privileges better than connecting as the superuser?
- How does the ORM make SQL injection harder, and where does the protection break down?
- Why does it pay to think the schema through up front rather than adding columns later?
- If you did the optional Alembic section: what is the difference between dropping and recreating a table versus running a migration, and why did the new column show NULL on every existing row?

RECAP of key concepts:

- Relational databases store data in typed tables; SQL is the language that queries and changes them.
- SQLite is the database for this course: a single file, no server, no users, no passwords. PostgreSQL is the production next step we stop short of here.
- SQLAlchemy maps tables to Python classes and uses parameterised queries by default, so the model code looks the same against either backend.
- The one database secret, `DB_PASSWORD`, lives in `.env`; `models.py` builds the connection string from it, so the same code runs on dev, on prod, and under `systemd` alike.
- Optionally, Alembic generates and applies migrations so schemas can evolve without losing rows.

TEASER. We classify, we store, we query, but still no user can touch our backend.

BRIDGE. The service classifies, stores, and queries, but the only way in is `curl` with an API key on a non-standard port. No browser, no URL, no human-friendly surface. This needs to change next we put `nginx` in front of the FastAPI process, serves a small drawing frontend over plain HTTP and HTTPS, and route browser traffic into the same `POST /classify` and `GET /results` handlers you just wired up. The database stays where it is; the prediction it stores will start arriving from a canvas in a browser rather than a shell command.

Reverse Proxy & Frontend


2026-06-07 · lively boysenberry Wiesel

FORM FOLLOWS FUNCTION

The Why

THE ONLY WAY IN TODAY is a `curl` one-liner on port 8000 with a secret header attached. There is no human-friendly interface. The backend is the engine of the service; the frontend is what builds an experience around it and makes it accessible to people. Engineers tend to undervalue the frontend, and pay for it by not being understood.

THIS BLOCK closes that gap. We put Nginx in front of uvicorn, serve a small static frontend out of `/var/www/pixelwise/`, move the public surface off the application's port 8000 onto the standard web port 80, and route browser traffic through the same `POST /classify` and `GET /results` handlers we wired up in Block 6.



WHY NOT EXPOSE uvicorn TO THE INTERNET DIRECTLY? Put the process that classifies digits on the public network and every malformed request, port scan, and slow-connection flood hits your application code, where one crash takes the whole site down. A reverse proxy sits in front instead: It accepts each request, calls uvicorn on port 8000 itself, and returns the response, so the client never touches

Figure 10:

The browser reaches our reverse proxy Nginx over HTTP on port 80. Nginx serves the static frontend and forwards `/api/` requests inward to the Block 6 service on `localhost:8000`. The dashed box is the prod VM, where only Nginx faces the network. The optional HTTPS exercise upgrades the browser connection to TLS on port 443.

the application. Nginx absorbs the rough traffic at the edge and forwards only clean requests inward, and we can restart, scale, or swap the upstream without the public URL ever moving.

Hands On Experience

THE PRODUCT IS INVISIBLE until you give it a front door. From here on, renting a cheap server is the only step between your project and the world. The gap between a codebase on your laptop and a live service on the internet is huge. Build, share, learn, and have fun, and do not let anyone make your world small and narrow.

THIS SEVENTH BLOCK introduces the public surface for PixelWise. By the end you will have

- installed and configured Nginx as a reverse proxy on prod,
- served a static HTML and JavaScript drawing frontend from `/var/www/pixelwise/`,
- and reached the live application in a browser.

An optional section at the end of the block carries the edge further, terminating TLS so the site answers on `https://`.

Optional Exercise: Catch Up to v0.5

CATCHING UP VIA TAG. If you are not already at the end of Block 6, rewind both VMs to v0.5: PostgreSQL provisioned, the pixelwise role and database created, the SQLAlchemy Prediction model wired into the FastAPI handlers, and POST /classify writing rows that GET /results reads back. Both VMs follow the same sequence; run dev first, then prod.

ON dev. Run the margin sequence. setup-server.sh installs PostgreSQL and provisions the pixelwise role and database; python init_db.py creates the predictions table.

ON prod. Same sequence after the ssh and cd. Verify from dev: curl -s http://192.168.56.11:8000/results should return {"results":[]}. Both machines are now at the end-of-Block 6 state, ready for Block 7.

TAG MAP. v0.5 marks the end of Block 6, v0.6 will mark the end of this block. Browse the full set at <https://github.com/schutera/pixelwise/tags>.

SEQUENCE ON dev.

```
cd ~/pixelwise
git fetch origin --tags
git reset --hard v0.5
cp .env.example .env
nano .env
source .venv/bin/activate
pip install -r requirements.txt
bash setup-server.sh
python init_db.py

Set DB_PASSWORD in .env and keep the
SECRET_API_KEY from Block 5. If .env/
is missing, run python3 -m venv .venv
first.
```

SEQUENCE ON prod.

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git fetch origin --tags
git reset --hard v0.5
cp .env.example .env
nano .env
source .venv/bin/activate
pip install -r requirements.txt
bash setup-server.sh
python init_db.py

Use the same DB_PASSWORD as on dev.
```

Reverse Proxies and Nginx

A REVERSE PROXY SITS AT THE EDGE. The client opens a connection to the proxy, the proxy opens a connection to the upstream, and the response flows back along the same path.

FORWARD VERSUS REVERSE. A forward proxy sits in front of the client and reaches out to many servers on the client's behalf, the way a corporate web filter does. A reverse proxy sits in front of the server and serves many clients on the server's behalf. Same word, opposite direction.

NGINX is the reverse proxy this course uses.

It is fast, well documented, and the most common piece of software you will find sitting in front of Python and Node services in production. The configuration is a tree of nested blocks: A top-level http context, server blocks inside it, and location blocks inside each server. Each server block declares a listen port and a hostname; each location declares how to handle requests whose paths match a pattern.

THREE ROLES IN ONE PROCESS. In the PixelWise stack Nginx wears three hats at once:

- As a static file server, it ships HTML, CSS, and JavaScript from `/var/www/pixelwise/` to the browser as fast as the kernel can read them.
- As a reverse proxy, it forwards every request whose path begins with `/api/` to `uvicorn` on `localhost:8000`.
- As a TLS terminator, it holds the certificate and the private key, so the public connection is encrypted while the internal connection to `uvicorn` stays plain HTTP.

`nginx -t` is the cheapest debugging tool in the chapter. It parses the configuration without applying it, prints `syntax is ok` when the file is well formed, and points at the line and column of the first problem when it is not. Run it before every `systemctl reload nginx` and the daemon will never refuse to start because of a missing semicolon or a typo in a directive name.

Nginx:

<https://nginx.org/>
Nginx documentation:
<https://nginx.org/en/docs/>

WHY A HIGH PORT? On Linux, ports below 1024 are privileged; binding requires either root or a capability granted by root. `uvicorn` runs as the unprivileged produser, so it listens on port 8000 where no special permission is needed. Nginx is started by `systemd` as root, binds 80 and 443, then drops its worker processes to an unprivileged user.

Exercise: Install Nginx and the First Proxy Pass

INSTALL NGINX ON prod. On a fresh prod VM, install the package and confirm the daemon is up and listening on port 80:

```
ssh produser@192.168.56.11
sudo apt update
sudo apt install -y nginx
systemctl status nginx --no-pager
```

The first three lines should report Active: active (running). Visit <http://192.168.56.11> from your host browser; the default Nginx welcome page proves the daemon is reachable on port 80 of the VM over the host-only network.

AUTHOR THE SITE CONFIGURATION ON dev. The Nginx configuration is part of how the application runs, so it belongs next to the code, the systemd unit, and the setup script. Open the new file at `deploy/pixelwise.nginx` in the repository on dev:

```
cd ~/pixelwise
nano deploy/pixelwise.nginx
```

The file below is the first cut, plain HTTP on port 80, the frontend served from disk, every request under `/api/` forwarded to uvicorn:

```
# deploy/pixelwise.nginx
server {
    listen 80;
    server_name 192.168.56.11;

    # Serve the static frontend
    location / {
        root /var/www/pixelwise;
        index index.html;
        try_files $uri $uri/ =404;
    }

    # Forward API calls to uvicorn
    location /api/ {
        proxy_pass http://127.0.0.1:8000/;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
    }
}
```

THE DEFAULT SITE. A fresh install ships an Nginx Welcome page from `/etc/nginx/sites-enabled/default`. We will leave it in place until our own site is ready, then disable it in one step so there is no confusion about which configuration is serving requests.

```

    }
}

```

COMMIT ON dev, PULL ON prod. The configuration belongs in the repository so a fresh VM can pick it up from a clean clone instead of from somebody's terminal history:

```

cd ~/pixelwise
git add deploy/pixelwise.nginx
git commit -m "Add Nginx reverse proxy config"
git push

```

Then on prod, pull the file into place:

```

ssh produser@192.168.56.11
cd /opt/pixelwise
git pull

```

INSTALL THE SITE ON prod. Nginx looks for active sites under /etc/nginx/sites-enabled/. The conventional pattern is to keep the real files under sites-available/ and symlink the active ones into sites-enabled/, so disabling a site is one rm away:

```

sudo cp deploy/pixelwise.nginx \
    /etc/nginx/sites-available/pixelwise
sudo ln -sf /etc/nginx/sites-available/pixelwise \
    /etc/nginx/sites-enabled/pixelwise
sudo rm -f /etc/nginx/sites-enabled/default
sudo nginx -t
sudo systemctl reload nginx

```

CREATE THE FRONTEND DIRECTORY. The location / block points at /var/www/pixelwise/ as the document root. That directory does not exist yet on prod, and the next exercise will fill it. For now, drop a placeholder so the proxy pass has something to serve when the path is /:

```

sudo mkdir -p /var/www/pixelwise
echo "PixelWise frontend stub" | \
    sudo tee /var/www/pixelwise/index.html

```

DID THE PROXY TAKE? From dev, hit the proxied API and the static path back to back:

```

curl -s http://192.168.56.11/
curl -s http://192.168.56.11/api/health

```

TRAILING SLASH MATTERS.

proxy_pass http://127.0.0.1:8000/ strips the matched /api/ prefix before forwarding. Without the trailing slash, /api/classify would reach uvicorn as /api/classify and FastAPI would answer 404 because the route is registered as /classify. One character, opposite outcomes.

try_files IS THE STATIC FALLBACK.

For a single-page frontend like ours, the directive tries the requested path, then a directory listing, then returns 404. A more elaborate frontend with client-side routing would fall back to index.html instead so deep links survive a reload. Our app has one page, so the 404 fallback is correct.

RELOAD, NOT RESTART. systemctl

reload nginx tells the master process to re-read the configuration and gracefully replace its workers without dropping any in-flight connection. restart would kill the master and start it again, briefly closing the listening sockets. Reload is the default for configuration changes; restart is for the rare case where reload is not enough.

The first line should print `PixelWise frontend stub`. The second should print the same JSON `uvicorn` would have given you on port 8000, except now `Nginx` is the one that took the request. Two different programs answered two paths, and from the outside they look like one service on one port.

```
DID IT TAKE ON prod? curl -s
http://192.168.56.11/ should
print the placeholder text. curl -s
http://192.168.56.11/api/health
should print
{"status": "ok", "model_version": "v1"}.
The first answer came from Nginx
reading a file on disk; the second
travelled from Nginx to uvicorn on port
8000 and back. The browser sees one
origin, http://192.168.56.11, even
though two very different processes
produced the two responses.
```

The Static Frontend

HTML FOR STRUCTURE, CSS FOR LAYOUT, JAVASCRIPT FOR BEHAVIOUR. Those are the three pieces a browser needs to render an interactive page, and the three files PixelWise ships.

- HyperText Markup Language, HTML, describes the document tree.
- Cascading Style Sheets, CSS, describe how it looks.
- JavaScript, JS, describes what happens when the user clicks, draws, or types.

The browser is the runtime, the network is the transport, and the server is a file system that happens to speak HTTP.

THE PAGE IS A SINGLE CANVAS, TWO BUTTONS, AND A RESULTS PANEL. The drawing area is a `<canvas>` element shown at 280 by 280 pixels, a tenfold magnification of the 28 by 28 grid the model actually reads. The JavaScript paints straight into that grid, so the user draws at the model's own resolution and watches the strokes pixelate as they land. A `Classify` button extracts the pixel array and posts it. A results panel renders the predicted digit, the confidence, and a scrolling history pulled from `GET /api/results`. That is the whole user surface.

THE `fetch` CALL is the bridge between the browser and the backend. `fetch` is the browser's modern HTTP client, a Promise-based function that takes a URL and an options object and returns a response the JavaScript can read as JSON, text, or bytes. A `POST` to `/api/classify` carries the pixel array in the body and the API key in a header, and waits for the JSON answer. The exercise below builds this into `app.js`, where the `classify` function issues exactly this request.

THE RELATIVE PATH IS THE HEADLINE. `fetch("/api/classify", ...)` names no hostname or port. It resolves against the URL the page was loaded from, so the same JavaScript runs unchanged in the classroom on `http://192.168.56.11` and in production on `https://pixelwise.yourdomain.com`. The frontend has no knowledge of where the backend lives, because the backend lives wherever the frontend was served from.

WHY NO FRAMEWORK? React, Vue, and Svelte solve problems that show up at scale: Component reuse across dozens of views, state management across complex interactions, routing across many pages. PixelWise has one page, one canvas, one results panel. A framework would add a Node.js toolchain, a bundler, a learning curve orthogonal to this course, and an extra layer hiding the bytes that actually travel between browser and server. Vanilla JS keeps the HTTP layer visible.

SAME-ORIGIN BY DESIGN. The page is served from `http://192.168.56.11/` and posts to `/api/classify` on the same host. The browser treats both as the same origin because the scheme, host, and port all match.

OPEN BY DESIGN. The `X-API-Key` header sits on `POST /api/classify` because that endpoint is gated. `GET /api/results` and `GET /api/health` carry no header because they are unprotected: The browser renders the results panel and probes the service without juggling a secret in client-side JavaScript. That is acceptable for the demo here; in a real product `/api/results` would move behind a session cookie or a short-lived token from a proper login flow.

Exercise: Build and Deploy the Frontend

CREATE THE FRONTEND DIRECTORY ON dev. The three files live in the repository under frontend/ so they travel with the code, get reviewed in pull requests, and reach prod through the same Git path as the Python:

```
cd ~/pixelwise
mkdir -p frontend
nano frontend/index.html
```

The HTML is small. One canvas, two buttons, a result line, a recent-predictions list, one stylesheet link, and one script tag at the bottom so the document tree is built before the JavaScript looks for elements in it:

```
<!doctype html>
<html lang="en">
<head>
  <meta charset="utf-8">
  <meta name="viewport"
    content="width=device-width, initial-scale=1">
  <title>PixelWise</title>
  <link rel="stylesheet" href="style.css">
</head>
<body>
  <h1>PixelWise</h1>
  <p>Draw a digit, then click Classify.</p>
  <canvas id="pad" width="280" height="280"></canvas>
  <div class="controls">
    <button id="classify">Classify</button>
    <button id="clear">Clear</button>
  </div>
  <div id="result"></div>
  <h2>Recent predictions</h2>
  <ul id="history"></ul>
  <script src="app.js"></script>
</body>
</html>
```

AUTHOR THE STYLESHEET next to the HTML. The styling is minimal: A centred column, a visible canvas border, a button row, a result line, and a monospace history list. The point is readability, not visual design:

WHY A 280-PIXEL CANVAS? The model expects 28 by 28 grayscale pixels. Drawing at that size on a modern display would feel like trying to write on a postage stamp. So the visible canvas is that 28 by 28 grid magnified ten times: The user gets a comfortable pad, the model gets the input shape it was trained on, and nothing is hidden, because the chunky blocks on screen are exactly the pixels that travel to the API.

```
nano frontend/style.css
```

```
body {
  font-family: system-ui, sans-serif;
  max-width: 480px;
  margin: 2rem auto;
  padding: 0 1rem;
}
canvas {
  border: 1px solid #333;
  background: #fff;
  touch-action: none;
  display: block;
}
.controls { margin: 0.5rem 0 1rem 0; }
button {
  padding: 0.5rem 1rem;
  margin-right: 0.5rem;
  font-size: 1rem;
}
#result {
  font-size: 1.25rem;
  min-height: 1.5rem;
  margin-bottom: 1rem;
}
#history li { font-family: monospace; }
```

AUTHOR THE SCRIPT. `frontend/app.js` wires the canvas to the mouse, builds the pixel array, posts it, renders the prediction, and refreshes the history list. Open the file:

```
nano frontend/app.js
```

```
// frontend/app.js
const API_KEY = "REPLACE_ME"; // overwritten on deploy

const N = 28; // the model's input grid
const SCALE = 10; // on-screen pixels per grid cell (280 / 28)

const pad = document.getElementById("pad");
const view = pad.getContext("2d");
view.imageSmoothingEnabled = false; // draw crisp blocks

// The real drawing happens on a hidden 28x28 grid. The
// visible canvas is that grid magnified ten times, so
```

PixelWise

Draw a digit, then click Classify.



Recent predictions

- 6 0.71 2026-06-07T17:29:24.406413
- 4 0.46 2026-06-07T17:29:18.909609
- 1 1.00 2026-06-07T17:29:00.494572
- 8 0.71 2026-06-07T17:29:02.487546
- 7 1.00 2026-06-07T17:28:55.746267
- 4 1.00 2026-06-07T17:28:47.875584
- 4 1.00 2026-06-07T17:28:17.818776
- 7 0.56 2026-06-07T17:28:13.388762
- 8 1.00 2026-06-07T17:28:09.416996
- 2 1.00 2026-06-07T17:28:03.400223
- 2 0.84 2026-06-07T17:27:59.887430
- 1 1.00 2026-06-07T17:27:23.439512
- 5 0.74 2026-06-07T16:03:53.509456
- 8 0.81 2026-06-07T16:00:14.576455
- 1 0.81 2026-06-07T16:00:07.712073
- 1 1.00 2026-06-07T15:59:44.936974

Figure 11: The finished frontend in the browser: The canvas, the Classify and Clear buttons, the result line, and the recent-predictions list. Nginx serves these files and forwards the `/api/` calls to FastAPI on `localhost:8000`.

```

// the user paints at the model's own resolution.
const grid = document.createElement("canvas");
grid.width = N; grid.height = N;
const gctx = grid.getContext("2d");
gctx.lineWidth = 2.5;
gctx.lineCap = "round"; gctx.lineJoin = "round";

let drawing = false;

function render() {
  // Magnify the grid onto the canvas, smoothing off.
  view.drawImage(grid, 0, 0, pad.width, pad.height);
}

function clearPad() {
  gctx.fillStyle = "#fff";
  gctx.fillRect(0, 0, N, N);
  render();
}
clearPad();

// Mouse positions map onto the 28x28 grid via SCALE.
pad.onmousedown = e => {
  drawing = true; gctx.beginPath();
  gctx.moveTo(e.offsetX / SCALE, e.offsetY / SCALE);
};
pad.onmousemove = e => {
  if (!drawing) return;
  gctx.lineTo(e.offsetX / SCALE, e.offsetY / SCALE);
  gctx.stroke(); render();
};
pad.onmouseup = pad.onmouseleave = () => { drawing = false; };

function getPixels() {
  // The grid is already 28x28: read it and invert.
  const data = gctx.getImageData(0, 0, N, N).data;
  const pixels = [];
  for (let y = 0; y < N; y++) {
    const row = [];
    for (let x = 0; x < N; x++)
      row.push(255 - data[(y * N + x) * 4]);
    pixels.push(row);
  }
  return pixels;
}

```

```

}

async function classify() {
  const r = await fetch("/api/classify", {
    method: "POST",
    headers: {
      "Content-Type": "application/json",
      "X-API-Key": API_KEY
    },
    body: JSON.stringify({ pixels: getPixels() })
  });
  const out = document.getElementById("result");
  if (!r.ok) { out.textContent = "Error " + r.status; return; }
  const d = await r.json();
  out.textContent = `Prediction: ${d.prediction} ` +
    `(${(d.confidence * 100).toFixed(1)}%)`;
  refresh();
}

async function refresh() {
  const r = await fetch("/api/results");
  if (!r.ok) return;
  const ul = document.getElementById("history");
  ul.innerHTML = "";
  for (const row of (await r.json()).results) {
    const li = document.createElement("li");
    li.textContent = `${row.prediction} ` +
      `(${row.confidence.toFixed(2)} ${row.created_at})`;
    ul.appendChild(li);
  }
}

document.getElementById("classify").onclick = classify;
document.getElementById("clear").onclick = () => {
  clearPad();
  document.getElementById("result").textContent = "";
};
refresh();

```

COMMIT, PUSH, PULL. On dev, stage the three files and push:

```

cd ~/pixelwise
git add frontend/
git commit -m "Add static drawing frontend"

```

WHY INVERT THE CHANNEL? The MNIST training set stores digits as light strokes on a dark background, with zero meaning empty and high values meaning ink. The canvas draws black on white, the opposite convention. 255 - data[i] flips the polarity so the model sees the input distribution it was trained on. Skip the inversion and the classifier will report 8 for a blank canvas because the white background looks like a dense glyph.

THE API KEY IN CLIENT CODE. const API_KEY sits in plaintext JavaScript that any visitor can view, so it is not a real secret. That is fine for the classroom:


```
git push
```

On prod, pull and copy the frontend into the document root. The repository is a Git clone under /opt/pixelwise/, and the location / block in our Nginx config serves /var/www/pixelwise/, so copy the files across:

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git pull
sudo cp -r frontend/* /var/www/pixelwise/
```

EDIT THE API KEY IN PLACE. The script ships with `const API_KEY = "REPLACE_ME"`; the value that actually works lives in /opt/pixelwise/.env as `SECRET_API_KEY`. On prod, replace the placeholder so the deployed script carries a real key. A small sed command keeps it auditable:

```
KEY=$(grep ^SECRET_API_KEY /opt/pixelwise/.env \
| cut -d= -f2)
sudo sed -i \
"s/REPLACE_ME/$KEY/" \
/var/www/pixelwise/app.js
```

The same pattern moves into `setup-server.sh` in the next section, so a fresh provision picks up the right key without any manual editing.

DID THE FRONTEND TAKE? Open `http://192.168.56.11/` in your dev's browser. You should see the heading, the canvas, two buttons, and an empty recent-predictions list. Draw a clear digit, click `Classify`, and read the result line: The predicted digit and its confidence should appear, and the history list should grow by one entry that survives a page reload because it now lives in PostgreSQL.

DID IT TAKE? If the page loads but `Classify` produces `Error 401`, the API key in `app.js` does not match the one in `.env`. If the page loads but `Classify` produces `Error 404`, the `proxy_pass` trailing slash got lost in a save. If the page does not load at all, run `sudo nginx -t` on prod: A syntax error in the config block keeps the daemon from applying the new configuration.

Teach setup-server.sh to Provision the Edge

THE COMMANDS YOU JUST TYPED BELONG IN THE SCRIPT. A fresh prod VM should reach the same v0.6 state with no manual steps. On dev, add nginx to the apt install line and append one prod-gated, idempotent stanza after the systemd unit from Block 5:

```
# Install Nginx site and deploy the frontend on prod
if [ -f deploy/pixelwise.nginx ] && \
  command -v nginx >/dev/null 2>&1 && \
  id produser >/dev/null 2>&1; then

  sudo mkdir -p /var/www/pixelwise
  sudo cp -r frontend/* /var/www/pixelwise/

  KEY=$(grep ^SECRET_API_KEY /opt/pixelwise/.env \
    | cut -d= -f2)
  sudo sed -i "s/REPLACE_ME/$KEY/" \
    /var/www/pixelwise/app.js

  sudo cp deploy/pixelwise.nginx \
    /etc/nginx/sites-available/pixelwise
  sudo ln -sf /etc/nginx/sites-available/pixelwise \
    /etc/nginx/sites-enabled/pixelwise
  sudo rm -f /etc/nginx/sites-enabled/default
  sudo nginx -t && sudo systemctl reload nginx
fi
```

COMMIT, PUSH, RUN, TAG. Push the script, rerun it on prod to confirm it converges to the same state, then tag the block's end state:

```
# on dev
git add setup-server.sh
git commit -m "Provision Nginx and frontend in setup-server.sh"
git push

# on prod
ssh produser@192.168.56.11
cd /opt/pixelwise
git pull
bash setup-server.sh
```

PROD-ONLY AND IDEMPOTENT. The `id produser` guard keeps the edge off dev, where there is no public traffic to serve. Every step is safe to rerun: `ln -sf` refreshes the symlink, `cp` overwrites in place, and `sed` reads the live `SECRET_API_KEY` so the repository never holds the real key.

TAG v0.6.

```
git tag -a v0.6 -m "Nginx proxy
and frontend"
git push origin v0.6
```

The end state is the static frontend behind an Nginx reverse proxy, provisioned by `setup-server.sh`. An annotated tag pushed by name travels reliably; `--follow-tags` skips lightweight tags.

Optional TLS and HTTPS

THE CORE STACK ANSWERS A BROWSER OVER PLAIN HTTP. A request to `http://192.168.56.11/api/classify` carries the pixel array, the path, and every header, including `X-API-Key`, as readable bytes on the wire. On a host-only network with two VMs that is a theoretical risk; on a public network it is a credential leak waiting for the first packet capture. The fix is Transport Layer Security, TLS, which encrypts everything between browser and server with a key the two parties agree on at handshake time.

HTTPS IS HTTP OVER TLS. Same protocol, same methods, same headers, same status codes: A TLS handshake opens before the first HTTP byte, the rest of the connection is encrypted, and the browser shows a lock icon. Nginx holds the certificate and the private key, so the TLS layer lives in one place at the edge of the system.

SELF-SIGNED CERTIFICATES ARE TRAINING WHEELS. A self-signed certificate is one the server signed itself. It encrypts the channel just as well as one from a public authority, but the browser warns because it does not recognise the signer. For the classroom that warning is the lesson: The chain of trust between your browser and a random IP address is exactly zero. The classroom can run on it; a production deployment with users cannot, and swaps in a certificate from a real authority, leaving the rest of the configuration unchanged.

THE PRODUCTION REPLACEMENT IS LET'S ENCRYPT. Let's Encrypt is a free certificate authority. Its client `certbot` proves you control the domain by answering a challenge on port 80, installs a certificate every browser trusts, and renews it automatically through a `certbot.timer` systemd unit. With a real domain pointing at a public IP, provisioning is three commands:

```
sudo apt install -y certbot python3-certbot-nginx
sudo certbot --nginx \
    -d pixelwise.example.com \
    --redirect --agree-tos -m you@example.com
sudo systemctl status certbot.timer
```

`certbot --nginx` edits the site configuration in place, filling in the `ssl_certificate` paths and rewriting the port 80 redirect, so nothing in your `deploy/pixelwise.nginx` needs to know it ever ran self-signed first.

TLS IN THREE SENTENCES. The server presents a certificate that binds a hostname to a public key. The browser checks it against a trust store of known authorities. If it validates, the two sides agree on a session key with public-key cryptography, then encrypt the rest of the conversation symmetrically.

TLS overview: <https://datatracker.ietf.org/doc/html/rfc8446>

Let's Encrypt: <https://letsencrypt.org/>
 certbot: <https://certbot.eff.org/>

OUT OF SCOPE FOR THE VM LAB. Let's Encrypt needs a public, internet-routable hostname to issue a certificate, which the host-only network does not provide. The commands here map where the real path goes once you leave the classroom; the mechanics are identical, only the signer changes.

Self-Reflection and Recap

SELF-REFLECTION questions to guide your thoughts during and after the exercises:

- Why does uunicorn stay on port 8000 once Nginx is in place, and what does Nginx gain by owning port 80 alone?
- The frontend calls the API at the relative path `/api/` and ships the key in plain JavaScript any visitor can read. What does the relative path buy, and why is the visible key acceptable here?
- A browser shows 502 Bad Gateway. Do you read `/var/log/nginx/access.log` or `journalctl -u pixelwise` first, and why?
- `setup-server.sh` automates site configuration. What does the next `deploy` undo if you edit `/etc/nginx/` on prod by hand?
- If you ran the optional TLS section: A self-signed certificate makes the browser warn instead of connecting quietly. What does that warning still teach about the chain of trust, even on a network you fully control?

RECAP of key concepts:

- A reverse proxy is the public face of an internal service: Nginx faces the network, serves the static frontend, and forwards `/api/` requests to uunicorn on localhost.
- A static frontend of plain HTML, CSS, and JavaScript runs in any browser without a build step, and a same-origin relative fetch keeps the API contract simple.
- The site configuration lives in the repository and `setup-server.sh` installs it in one idempotent run, so a fresh VM reaches the same state with no manual edits.
- TLS terminates at the edge, where a self-signed certificate teaches the handshake and the role of Nginx as terminator before a real deployment swaps in a certificate from a public authority.

BRIDGE. The stack works, but the workflow does not. Every change reached prod by hand, a `git push` on dev, an `ssh` into prod, a `git pull`, a `bash setup-server.sh`, and a hope that nothing breaks. Yes, this is Handwerk - but even so it is okay to hand-off what's possible. The application stays exactly as it is; next we automate that path, so a push checks and ships itself.

CROSS-ORIGIN IS THE ADVANCED CASE. The relative fetch keeps the page and the API on one origin, so the browser's same-origin policy never intervenes. Serve them from different domains and the browser blocks the read until the server opts in with Cross-Origin Resource Sharing, CORS, headers. That is the first deployment puzzle beyond this block.

CORS reference: <https://developer.mozilla.org/en-US/docs/Web/HTTP/CORS>

TEASER. We ship by hand on every push, the server runs commands typed by a tired student at 11pm, one wrong `rm` away from a broken site. What could possibly go wrong?

Continuous Integration Continuous Deployment

2026-06-12 · valiant avocado Kranich

OPERATION VACATION.

The Why

EVERY DEPLOY is still a human ssh'ing into prod, running `git pull`, restarting the service, and wasting a coffee break or feature development capacity. The smallest tweak goes through the same heavy-weight path as a real feature, and the only gate between a push and production is the operator's attention.

THIS BLOCK CLOSES THAT GAP. A single bash script on prod fetches origin, runs the test suite, and restarts the service when the tests pass. A systemd timer fires the script every thirty seconds, so a green commit on main ships without a human touching a terminal. The end of the block introduces feature flags as the runtime switch that lets a service change behaviour without another deploy.




Figure 12: The classroom CI/CD cascade, all of it living on prod. The systemd timer is the trigger and fires every thirty seconds. The test stage fetches origin and runs `pytest` against the new commit. The deploy stage restarts the systemd unit, but only when the gate is green. A red test breaks the cascade before the restart.

YOU WROTE A FIX. Ran it locally on dev, it looks right, you push to main, ssh to prod, pull, restart, and the service is down. The fix passed your single manual check and shipped a regression that broke the service. No tests ran the existing endpoints, no guard sat between your push and prod, and the only person who would have noticed is you.

EVERY MANUAL STEP is a step a (tired) human will skip at some point, and every skipped step is a regression waiting to ship. Automation does not replace careful engineering: it encodes the steps careful engineers already do, so they happen every time, in the same order, at large-scale, without anyone remembering.

Hands On Experience

OUTSIDE THE CLASSROOM. Teams call this practice DevOps. The artefacts are different at scale, with Jenkins, GitLab CI, or GitHub Actions runners replacing the timer, and Kubernetes deployments replacing `systemctl restart`, but the shape is identical: one human push, an automated gate, one observable change to production.

THIS EIGHTH BLOCK introduces automation around the PixelWise pipeline. By the end you will have

- written a smoke test for `classify_batch` that gives pytest something to gate on,
- written an auto-deploy script on prod that fetches origin, runs pytest, and restarts the service on green,
- registered a systemd timer that fires the script every thirty seconds,
- and seen how a feature flag would switch behaviour at runtime without another deploy.

THE PIPELINE IS THE ENGINEERING SURFACE. Once a pipeline runs on every push, the test suite, the deploy script, and the timer become part of how the application is developed. A good pipeline is a good engineering surface: it makes it easy to build fast and reliably. A poor pipeline will grind you to a halt, as your code base gets more complex.

CATCH UP TO v0.6. If dev or prod drifted from the end of Block 7, align both machines with a hard reset to the tag:

```
cd ~/pixelwise
git fetch origin --tags
git reset --hard v0.6
bash setup-server.sh
```

The full set of release tags lives at <https://github.com/schutera/pixelwise/tags>.

What is CI/CD?

CONTINUOUS INTEGRATION, CI, is the practice of merging every change into a shared branch as soon as it passes an automated gate. The gate is a test suite that rejects broken behaviour, run on a clean machine so the result is reproducible. Green is safe to merge; red does not enter main.

CONTINUOUS DEPLOYMENT, CD, extends the chain by one step. The same gate that admits a commit to main also ships it to production, with no human between the merge and the restart. The unit of release shrinks from "a quarter of features" to "every green commit on main", which makes each release small enough to reason about and small enough to revert.

THREE PIECES shape every CI/CD pipeline, no matter how elaborate the tool around them. A trigger says when to run. A gate says whether to proceed. A deploy step says what to do when the gate is green.

THE INDUSTRY SHAPE. GitHub Actions, Jenkins, and GitLab CI are the cloud-hosted version of the same idea. A push triggers a workflow on a fresh VM the platform provisions for the run. The workflow installs dependencies, runs tests, and on green opens an SSH session to production and ships the code. The configuration lives in YAML, the runner lives somewhere else, and usually you pay for this beyond free tiers.

THE CLASSROOM SHAPE. A single bash script on prod does the work. It runs `git fetch origin`, compares the local commit hash against `origin/main`, and exits silently if they match. If main has advanced, it pulls, runs `pytest`, and restarts the `systemd` unit only when the tests pass. A `systemd` timer fires the script every thirty seconds, so any commit that lands on main reaches prod before you have switched back to the browser to check.

WHAT THIS IS NOT. This block does not build a GitHub Actions workflow, does not stand up a self-hosted runner, and does not wire a webhook from the forge into the VM. Those are real and worthwhile patterns, and the references in this section's margin notes describe them; the chapter teaches the substance underneath them. In the industry this substance is usually wrapped in a hosted tool, which you will easily adapt to after you understand the shape here.

DEPLOY GATE, NOT CI GATE. The `pytest` step in this chapter runs on prod after the push has already merged into main. It gates the restart, not the merge: a red test keeps the service on the previous green commit, but main already carries the break and any developer pulling main pulls the break too. A CI gate runs the same test on a pull request before the merge, so main stays green by construction. We omit this step for simplicity and to keep the scope manageable in this lecture block.

BARE-METAL FIRST. The classroom builds the pipeline from scratch on the same machine that hosts the service. The trigger is a `systemd` timer, the gate is `pytest` on prod, the deploy step is one bash script. The substance is identical to a GitHub Actions workflow; the difference is that every part of it is on disk in front of you, debuggable with `journalctl`, and free.

Exercise: The Gate or a Smoke Test

THE GATE NEEDS SOMETHING TO GATE ON. A smoke test proves the wiring works without claiming anything about correctness. The name comes from electronics: turn the power on and see whether smoke comes out. A smoke test that fails means the code does not even run; a smoke test that passes does not mean the code is right, only that it runs, something we can work with.

WHERE TESTS LIVE. The convention in Python is a `tests/` directory at the repository root, with files named `test_*.py` and functions named `test_*`. `pytest` discovers them automatically; no registration, no configuration. The auto-deploy script in the next exercise runs `python -m pytest tests/` from `/opt/pixelwise` and refuses to deploy on any non-zero exit code.

CREATE `tests/test_classifier.py` on dev with a smoke test for `classify_batch`. The test builds a batch of two blank images, passes them through the classifier, and asserts the response shape:

```
cd ~/pixelwise
mkdir -p tests
nano tests/test_classifier.py

# tests/test_classifier.py
import numpy as np
from app.classifier import classify_batch

def test_classify_batch_returns_one_per_image():
    images = np.zeros((2, 28, 28), dtype=np.uint8)
    results = classify_batch(images)
    assert len(results) == 2

def test_classify_batch_result_shape():
    images = np.zeros((1, 28, 28), dtype=np.uint8)
    result = classify_batch(images)[0]
    assert "prediction" in result
    assert "confidence" in result
    assert isinstance(result["confidence"], float)
```

INSTALL `pytest` AND RUN THE TEST LOCALLY ON dev before pushing, from the repository root so `pytest` resolves the app package

GO DEEPER LATER. The `pytest` documentation at <https://docs.pytest.org/> is the canonical reference. For a book-length treatment, Harry Percival's "Test-Driven Development with Python" is free online at <https://www.obeythetestinggoat.com/>. A second testing block belongs in a follow-on course.

TWO FUNCTIONS, ONE FILE. The first function checks the batch shape: a batch of two images in, two results out. The second function checks the result keys: every result is a dict with `prediction` and `confidence`. Together they prove the function loads, accepts the expected input, and returns the expected shape.

against the on-disk source tree:

```
cd ~/pixelwise
source .venv/bin/activate
pip install pytest
python -m pytest tests/
```

COMMIT AND PUSH the test. The next exercise builds the auto-deploy script that runs the same test on prod and gates the restart on it.

SAME GATE, DIFFERENT TRIGGER.

The same python `-m pytest tests/` invocation that you just did locally and we will gate deployment later is exactly what a CI runner would also execute on every pull request before merging into main.

DID IT TAKE ON dev?

`python -m pytest tests/` should print two passes. If pytest reports an import error like `ModuleNotFoundError: No module named 'app'`, you ran the command from the wrong directory; cd into `~/pixelwise` and try again. If `classify_batch` itself fails to import, the model file is missing on dev; rerun `bash setup-server.sh`.

```
STAGE AND SHIP. git add tests/
git commit -m "Add smoke
test for classify_batch"
git push
```

Auto-Deploy on prod

THE PIPELINE IS ONE BASH SCRIPT. The script lives in the repository and runs on prod. Its body is the trigger, the gate, and the deploy step laid out in order: fetch from origin, compare local against remote, and either exit or pull, install, test, and restart. A handful of operations, one shell file, version controlled like everything else in the project.

A COMMIT HASH COMPARISON IS THE TRIGGER. The script asks Git for the local commit hash on main and the remote commit hash after a fresh fetch. Equal hashes mean nothing changed since the last run and the script exits cheaply. Different hashes mean a new commit landed upstream and the rest of the script runs. The commit hash is the cheapest test for "is there work to do".

THE GATE IS pytest. After the pull, the script runs the test suite against the new commit. A non-zero exit code aborts the script before the restart. The previous service keeps running on the old commit, the operator sees the failure in the system log, and the next run will retry whatever the developer pushes to fix it.

THE TRIGGER IS A SYSTEMD TIMER. A service unit runs the script as a regular user; a timer unit activates the service every thirty seconds. The system survives reboots because the timer is wired into the boot target. Cron would do the same job with one line of a crontab; pick cron when systemd is not available, pick a systemd timer when it is. The substance is the same scheduler running the same script.

TWO THINGS THIS DESIGN INTENTIONALLY LACKS. The timer pulls, so there is no push notification from GitHub; a new commit waits for the next tick. The script is local, so there is no central dashboard of recent runs; the system journal is the dashboard. Both gaps trade visibility for simplicity, and the trade is the point: at this stage of the course, the goal is the substance underneath the dashboard, not the dashboard - you usually get that as free lunch anyways, but it will taste better now that you know how it was prepared.

ONE SCREEN, ONE HEAD. A pipeline that fits on one screen fits in one head. The deploy script is short enough to read top to bottom, the trigger is a single timer unit, and the gate is one pytest call. Everything else is an optimisation.

IDEMPOTENCE AND THE TIMER. The script is safe to run every thirty seconds forever. On every run where origin has not moved, it does the cheap fetch and exits. On every run where the tests fail, it leaves the running service untouched. The deploy is the side effect, not the schedule; a timer that fires once per second would do the same work and reach the same state, only louder.

PUSHING TO main IS SHIPPING. With auto-deploy live, every commit that lands on main restarts prod within seconds. A solo classroom workflow pushes straight to main because the operator is also the reviewer; a team workflow does not. The team shape is four steps:

- branch off main for the change, so main stays untouched while you work.
- commit on the branch as often as needed; nothing reaches prod yet.
- Open a PR so the diff, the CI run, and a reviewer meet before the merge.
- merge into main only on green CI and an approving review, at which point the pipeline ships it.

With auto-deploy on, the merge is the deploy; the PR is the last place to catch a regression cheap.

Exercise: Build and Activate the Pipeline

STEP 1. THE TRIGGER. On dev, author the systemd service unit and the timer unit. Both files live in the repository and ship to prod on the next pull; systemd on prod will read them once the provisioner in Step 3 copies them into `/etc/systemd/system/`. The service unit runs the script once, as produser, from the working directory `/opt/pixelwise`; the timer activates that service fifteen seconds after boot and then every thirty seconds:

```
cd ~/pixelwise
mkdir -p deploy/systemd
nano deploy/systemd/pixelwise-deploy.service

[Unit]
Description=PixelWise auto-deploy
After=network-online.target

[Service]
Type=oneshot
User=produser
WorkingDirectory=/opt/pixelwise
ExecStart=/opt/pixelwise/deploy/auto-deploy.sh

nano deploy/systemd/pixelwise-deploy.timer

[Unit]
Description=Run PixelWise auto-deploy

[Timer]
OnBootSec=15s
OnUnitActiveSec=30s
Unit=pixelwise-deploy.service

[Install]
WantedBy=timers.target
```

`Type=oneshot`. A service unit usually models a long-running daemon; the manager keeps the process alive and restarts it on exit. A oneshot unit models a job that runs to completion: systemd waits for the script to exit, records the exit code, and goes back to sleep. That is exactly what an auto-deploy script is.

`OnBootSec` vs `OnUnitActiveSec`.
`OnBootSec=15s` fires the first run fifteen seconds after boot;
`OnUnitActiveSec=30s` fires each subsequent run thirty seconds after the previous one finished. The pair gives you a quick first check after every reboot and a tight cadence in between, so a push to main reaches prod before you have switched windows. systemd accepts plain second values here, and a tick that does nothing but `git fetch` costs the VM nothing; production deployments tend to be slower because the gate, not the timer, sets the floor.

STEP 2. THE DEPLOY STEP. Still on dev, author `deploy/auto-deploy.sh`.

The file ships through Git to prod, where the timer from Step 1 will execute it; that is why the script's first line is `cd /opt/pixelwise`, the path it will sit at on prod, not the `~/pixelwise` you are editing it in. The script is the whole pipeline: fetch origin, compare local commit hash against origin/main, conditionally pull, install pinned dependencies, gate on pytest, restart the service on green:

```
mkdir -p deploy
nano deploy/auto-deploy.sh

#!/bin/bash
# deploy/auto-deploy.sh
# Authored on dev, executed on prod by the timer.
# Pull origin/main, gate on pytest, restart
# the service on green.
set -euo pipefail

cd /opt/pixelwise

git fetch origin

LOCAL=$(git rev-parse HEAD)
REMOTE=$(git rev-parse origin/main)
if [ "$LOCAL" = "$REMOTE" ]; then
    exit 0
fi

echo "New commit on main: $REMOTE"
git pull origin main

source .venv/bin/activate
pip install -r requirements.txt > /dev/null

if ! python -m pytest tests/; then
    echo "Tests failed, refusing to deploy."
    exit 1
fi

sudo systemctl restart pixelwise
echo "Deployed: $REMOTE"
```

Mark the script executable. The executable bit travels with the file in Git, so prod gets it on pull without a second `chmod`:

```
chmod +x deploy/auto-deploy.sh
```

`set -euo pipefail`. `-e` exits on any non-zero command, `-u` errors on unset variables, `-o pipefail` fails the whole pipeline if any stage fails. A silent failure inside a scheduled deploy is the worst failure mode a pipeline can have.

Why `sudo systemctl restart`? `produser` owns the code but not the service unit; restarting a system-level unit needs root. `setup-server.sh` drops a one-line sudoers rule into `/etc/sudoers.d/pixelwise` that grants `produser` a passwordless sudo entry for exactly this one command and nothing else:

```
produser ALL=(root) NOPASSWD: \
    /usr/bin/systemctl restart
    pixelwise
```

The script can do its one privileged step without holding broad privilege. If the rule is missing, the journal will show a password prompt and `exit 1` on the next deploy; re-running `setup-server.sh` reinstalls it.

STEP 3. ACTIVATE THE PIPELINE. Still on dev, teach `setup-server.sh` to install both unit files and enable the timer. Append a prod-only stanza so a fresh prod VM picks up the service unit, the timer unit, and the enabled state without any manual steps:

```
nano setup-server.sh

# Install the auto-deploy systemd timer on prod
if [ -f "$SCRIPT_DIR/deploy/systemd/pixelwise-deploy.timer" ] \
  && command -v systemctl >/dev/null 2>&1 \
  && id produser >/dev/null 2>&1; then
  sudo cp "$SCRIPT_DIR/deploy/systemd/pixelwise-deploy.service" \
    /etc/systemd/system/pixelwise-deploy.service
  sudo cp "$SCRIPT_DIR/deploy/systemd/pixelwise-deploy.timer" \
    /etc/systemd/system/pixelwise-deploy.timer
  sudo systemctl daemon-reload
  sudo systemctl enable --now pixelwise-deploy.timer
fi
```

Stage the script, both unit files, and the updated `setup-server.sh`, then push to origin. Switch to prod for the last manual deploy the chapter needs. SSH in, pull the new commit, and re-run the provisioner so the timer and the service unit land in `/etc/systemd/system/` and the timer starts ticking. After this run the timer takes over and no later step in the chapter touches prod by hand:

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git pull
bash setup-server.sh
```

Watch the pipeline live. On prod, in a second shell, tail the journal for the deploy unit. Every run prints either a single `exit 0` line on a no-op fetch, the new commit hash on a real deploy, or the pytest output on a red run:

```
journalctl -u pixelwise-deploy -f
```

Push a trivial commit from dev to see the end-to-end loop. A whitespace change to the README is enough to advance the commit hash on main. The next tick of the timer picks it up, runs the tests, restarts the service, and prints the new commit hash into the journal:

```
cd ~/pixelwise
echo "" >> README.md
git commit -am "Trivial commit to trigger auto-deploy"
git push
```

```
STAGE AND SHIP. git add deploy/ \
  setup-server.sh
git commit -m "Add auto-
  deploy script and timer"
git push
```

DID THE TIMER INSTALL?

```
On prod:
systemctl list-timers | grep
pixelwise
```

Should print one line with the next firing time, the previous run, and the timer unit name. If the line is missing, `systemctl status pixelwise-deploy.timer` usually shows the reason: a typo in the unit file, the service unit not found, or the timer not enabled.

READ THE JOURNAL LIKE A LOG.

```
-u narrows the journal to one unit.
-f follows it like tail -f, printing
every new line as the timer fires. A
pipeline you cannot read is a pipeline
you cannot trust; journalctl is the
dashboard the timer comes with for
free.
```

DID THE COMMIT LAND ON main?

After the next tick, the journal on prod should carry your new hash:

```
journalctl -u pixelwise-deploy
-n 5
```

A Deployed: `<hash>` line that matches the hash you just pushed means the pipeline did its job; no such line, or an older hash, is where the chain broke. Of course you can also just check the change in the readme.

Optional Exercise: Notify Discord on Deploy

A WEBHOOK IS AN HTTP POST TO A KNOWN URL. The pipeline runs in silence; the only signal of a deploy is a fresh commit hash in journalctl. A Discord webhook makes it audible: one curl call appends to a chat channel every time the script restarts the service or refuses to. Create one in the target channel under Integrations, Webhooks, "New Webhook", then "Copy Webhook URL".

ADD THE URL TO .env ON prod alongside the other runtime secrets, then edit auto-deploy.sh on prod to source .env near the top and to curl the webhook in the pytest failure branch and after the systemctl restart:

```
# /opt/pixelwise/.env
DISCORD_WEBHOOK_URL="https://discord.com/api/..."

# after `cd /opt/pixelwise`
if [ -f .env ]; then
    set -a; source .env; set +a
fi

# inside the pytest failure branch
curl -sf -H "Content-Type: application/json" \
    -d '{"content\":"Deploy red: $REMOTE\"}' \
    "$DISCORD_WEBHOOK_URL" > /dev/null || true

# after `sudo systemctl restart pixelwise`
curl -sf -H "Content-Type: application/json" \
    -d '{"content\":"Deployed: $REMOTE\"}' \
    "$DISCORD_WEBHOOK_URL" > /dev/null || true
```

Run the script once by hand to see the next message land in the channel:

```
sudo -u produser bash \
    /opt/pixelwise/deploy/auto-deploy.sh
```

THE URL IS THE SECRET. A Discord webhook URL ends in a long opaque token; anyone with the URL can post to the channel. Treat it like a password: keep it off GitHub and out of screenshots. The variable does not land in .env.example; each operator pastes their own URL into .env on the machine that needs it.

WHY || true? set -e aborts the script on any non-zero exit. A flaky Discord call should not roll back a successful deploy or hide a failed test; || true swallows the notification's exit so the rest of the script keeps going. Notifications are useful, not load-bearing.

NOT PART OF v0.7. The canonical auto-deploy.sh at v0.7 does not include the notification. Push it to your own main if you want to keep it; just know it is your local addition and the reference tag stays clean.

Optional: Feature Flags

A FEATURE FLAG IS A RUNTIME SWITCH. The code path exists in the binary, but a check at the start of the path decides whether to take it. The check reads a boolean from `.env`, an environment variable, or a row in a config table. Flipping the flag changes behaviour without changing code, without a deploy, and without a build.

A DEPLOY SHIPS CODE; A FLAG SHIPS BEHAVIOUR. Decoupling the two is the single most useful primitive in modern service operation: a broken feature turns off in seconds without a rollback, a risky migration runs dark for a day before going live, and a code path can land on main, pass CI, sit in production unused, and switch on only when the dashboards are ready to watch what it does. A commented-out path cannot do any of that: it is invisible to the test suite, the type checker, and any reviewer searching for usages, where a flag-gated path is real code that compiles, runs in CI, and lights up the moment the flag flips.

USE CASES IN THE WILD cover more shapes of release than the on/off picture suggests:

- Progressive rollout flips a path on for a fraction of traffic first.
- Dark launch deploys code before the user-visible surface is ready.
- A kill switch keeps a wrong path gated until the fix ships.
- An A/B experiment runs two arms in parallel and keeps the winner.
- A targeted flag fires only for one team's accounts so the new path runs in real traffic before a wider audience sees it.

HOW A FLAG IS READ matters more than what it contains. A flag read once at import time survives until the next restart; a flag read on every request turns on the moment the operator edits `.env` and the next request comes in. Small services read on every request; large services cache the read with a short refresh to keep the hot path fast.

THE SHAPE OF THE FLAG. A simple flag is a boolean. A multi-arm flag is a string from a known set, like `v1` vs `v2`. A percentage flag rolls a feature out to a fraction of traffic. A targeted flag fires for one user, one cohort, one region. The boolean is the floor; the others compose from it.

```
if NEW_UI: render_v2()
if MODEL == "v2": ...
if random() < 0.1: ...
if user in COHORT: ...
```

DARK LAUNCH IN DETAIL. The code ships to production, runs against real traffic, and stays invisible to users because the button, page, or endpoint that would expose it is gated off. A new search backend runs on every query, returns its answers into a sink, and logs latency and error rates next to the live backend it shadows. A new pricing engine recomputes every cart in the background and writes its number to a log column the customer never sees. The point is to learn from production load before any user can notice the change: the day the surface goes live the path has already been exercised at full scale, and the team argues from a week of graphs rather than a guess.

OUTSIDE THE CLASSROOM. Production teams use hosted flag systems like LaunchDarkly, Unleash, or Statsig that store flags out of band and push changes to running services in seconds. A small project uses `.env` because it is the lightest implementation of the same pattern; the mechanism is identical.

Tag v0.7

THE PIPELINE IS THE ARTEFACT. The end of every block is a tagged, runnable state. Stage every file the block added or changed, commit, push, tag, and push the tag. Future catch-ups reset to this tag and arrive at the same place you are now:

```
cd ~/pixelwise
git add tests/ \
    deploy/auto-deploy.sh \
    deploy/systemd/pixelwise-deploy.service \
    deploy/systemd/pixelwise-deploy.timer \
    setup-server.sh
git commit -m "End of Block 8: auto-deploy pipeline"
git tag -a v0.7 \
    -m "End of Block 8: auto-deploy script and timer"
git push --follow-tags
```

THE PIPELINE NOW RUNS WITHOUT YOU. Every commit on main is tested. Every green test ships. The service is automated.

CATCHING UP TO v0.7. If a future block finds either machine drifted from this state, align both with a hard reset to the tag and a re-run of the provisioner on prod. On dev the tag carries the test, the script, and the unit files into the working tree:

```
cd ~/pixelwise
git fetch origin --tags
git reset --hard v0.7
```

On prod the same reset is followed by the provisioner so the timer and the deploy service land in /etc/systemd/ and the next tick of the timer picks the tag up:

```
ssh produser@192.168.56.11
cd /opt/pixelwise
git fetch origin --tags
git reset --hard v0.7
bash setup-server.sh
```

The full set of release tags lives at <https://github.com/schutera/pixelwise/tags>.

WHAT LANDS IN v0.7.
tests/test_classifier.py
with the smoke test;
deploy/auto-deploy.sh as the pipeline;
deploy/systemd/pixelwise-deploy.service
and
deploy/systemd/pixelwise-deploy.timer
as the trigger; an extended
setup-server.sh that installs both
units. The optional Discord notification
and the feature flag theory do not
change the repository's tagged state.

WATCH THE PIPELINE LAND ITS OWN TAG. The push that creates v0.7 is a push to main, so the next tick of the timer picks it up. The script pulls the tagged commit on prod, runs the tests, restarts the service, and writes the new commit hash into the journal. The first time a chapter ships its own deploy is a moment worth pausing for.

WHY setup-server.sh AGAIN. A bare git reset restores the unit files inside the repository on prod, but systemd reads from /etc/systemd/system/. The Block 8 additions to the provisioner copy the timer and the service into that directory, run systemctl daemon-reload, and enable the timer. Re-running the provisioner is what closes the gap between repo state and systemd state.

Self-Reflection and Recap

SELF-REFLECTION questions to guide your thoughts:

- A test catches a regression on a feature branch; a user catches it in prod. Where in the pipeline is the cheapest place to catch the next one?
- What kinds of regression does the smoke test against `classify_batch` catch, and what kinds does it miss?
- The timer fires every thirty seconds. What property of `auto-deploy.sh` makes it safe to fire on the same commit hash again, and what breaks if that property goes away?
- Where does the classroom auto-deploy script differ from a GitHub Actions workflow that does the same job, and where is the substance identical?
- A feature flag stays `false` for three months and the code path never runs. What makes the flag-gated version safer to leave in the binary than the same code commented out?

RECAP of key concepts:

- Continuous Integration runs tests on every change against a known state; Continuous Deployment ships every green commit to main without a human in the loop.
- The classroom pipeline is three primitives: a systemd timer as the trigger, `python -m pytest tests/` as the gate, and one bash script as the deploy step.
- The smoke test in `tests/test_classifier.py` doubles as the gate: a non-zero exit from `pytest` stops the deploy before `systemctl restart`.
- `systemctl list-timers` shows the schedule and `journalctl -u pixelwise-deploy` reads its history; swap the timer for a GitHub Actions cron and the script body barely changes.
- Feature flags gate code that lives in the binary, so flipping behaviour is a config edit and a restart instead of a deploy.

MILESTONE. You now own a service that ships itself. A commit to main flows through tests, lands on prod without a human in the loop, and tag `v0.7` marks the pipeline itself as the artefact.

YOU MADE IT ZAUBERLEHRLING. Und nun sollen seine Geister Auch nach meinem Willen leben. Seine Wort' und Werke Merkt ich und den Brauch, Und mit Geistesstärke Tu ich Wunder auch.

Ship on something that runs

2026-06-09 · lively cherimoya Mufflon

HIGH FIVE.

The Why

The Projektbericht lets you extend a block of this course on your own snapshot. HighFive offers the other path: A real full stack application, running live at highfive.schutera.com, with its source on my GitHub and an open invitation to ship into it. Where the report proves you can build in isolation, a contribution proves you can build inside a system that other people already depend on.

HIGHFIVE IS AN OPEN MONITORING PIPELINE for wild bee nesting activity, built as a citizen science kit. Camera modules running on solar power photograph bee hotels, an analysis pipeline turns those images into counts and trends, and an interactive dashboard with a map puts the results in front of anyone who cares about pollinators.

IT IS THE FULL STACK YOU JUST LEARNED, end to end. The frontend is React with Vite and TypeScript. The backend is Node.js with Express alongside Python services on Flask. Data lives in DuckDB. Firmware for the ESP32-CAM camera modules is written in C++. Every layer from this course, from the Linux box up to the served frontend, has a home somewhere in this repository.

PICK SOMETHING AND MAKE IT BETTER. Sharpen the image pipeline, harden a backend service, tidy the dashboard, document a rough edge, or write the test that was missing. Open an issue describing what you want to do, agree the scope with me, then send a pull request. By prior agreement a merged contribution can form the basis of your Projektbericht: Same evaluation, real impact, and your name in the history of a project that outlives the semester. And yes, other

FIND HIGHFIVE HERE. The deployment runs at highfive.schutera.com, and the source lives at github.com/schutera/highfive.



Read the arc42 architecture document in the repository before you touch code; it explains how the pieces fit together.

BEFORE YOU OPEN A PULL REQUEST. HighFive is source available rather than open source: Free for any noncommercial use under the PolyForm Noncommercial License 1.0.0, with commercial use under a separate license. Contributions require you to accept the Contributor License Agreement in the repository. For commercial questions, write to schutera@dhbw-ravensburg.de.

students did ship into this repo in previous semesters, so you will be in good company.

Index

- [.env](#), [39](#), [50](#)
- [.env.example](#), [50](#)
- [.gitignore](#), [29](#), [69](#)
- [.pkl](#), [57](#), [60](#)
- [/etc](#), [17](#)
- [/opt](#), [17](#)
- [__pycache__](#), [69](#)
- [12-factor app](#), [38](#), [50](#)

- [abstraction](#), [54](#)
- [Alembic](#), [91](#), [104](#)
 - [autogenerate](#), [104](#)
- [alembic.ini](#), [104](#)
- [Ansible](#), [44](#)
- [API](#), [73](#), [75](#)
 - [authentication](#), [116](#), [120](#)
 - [design](#), [83](#)
- [APIKeyHeader](#), [86](#)
- [apt](#), [39](#)
- [assertion](#), [67](#), [68](#)
- [assessment](#), [9](#)
- [authentication](#), [84](#)
- [automation](#), [125](#), [130](#), [132](#), [134](#)

- [Bash](#), [16](#), [44](#)
- [bash](#), [130](#), [132](#), [134](#)
 - [set](#), [132](#)
- [batch inference](#), [67](#)

- [canvas](#), [107](#), [116](#), [117](#)
- [certbot](#), [123](#)
- [certificate](#), [123](#)
- [chmod](#), [16](#)
- [chown](#), [16](#)
- [CI](#), [129](#)
- [CI/CD](#), [124](#), [127](#), [136](#)
- [citizen science](#), [137](#), [139](#)
- [classifier](#), [55](#)
- [code review](#), [130](#)

- [conda](#), [46](#)
- [containers](#), [54](#)
- [continuous deployment](#), [127](#)
- [continuous integration](#), [36](#), [127](#)
- [contribution](#), [137](#)
- [CORS](#), [107](#), [124](#)
- [course work](#), [9](#)
- [CSS](#), [107](#), [116](#)
- [curl](#), [80](#), [89](#), [103](#)
- [CVE](#), [48](#), [49](#)

- [data hygiene](#), [26](#)
- [database](#), [91](#)
- [DATABASE_URL](#), [97](#), [99](#)
- [dependencies](#), [38](#)
- [deploy](#), [125](#), [130](#), [131](#)
- [deployment](#), [73](#)
- [dev VM](#), [11](#)
- [Discord](#), [134](#)
- [distribution](#), [16](#)
- [Docker](#), [54](#)
- [Dockerfile](#), [54](#)
- [DVC](#), [26](#)

- [Ed25519](#), [18](#)
- [environment variables](#), [50](#)
- [exercises](#), [14](#), [17](#), [20](#), [25](#), [27](#), [31](#), [33](#), [34](#), [37](#), [41](#), [45](#), [47](#), [49](#), [51](#), [52](#), [59](#), [61](#), [65](#), [68](#), [77](#), [79](#), [82](#), [85](#), [88](#), [95](#), [97](#), [100](#), [102](#), [104](#), [111](#), [113](#), [117](#), [128](#), [131](#), [134](#)

- [FastAPI](#), [73](#), [80](#)
- [feature flag](#), [124](#), [135](#)
 - [dark launch](#), [135](#)
- [fetch](#), [116](#)
- [fork](#), [35](#)
- [frontend](#), [107](#), [116](#), [117](#)

- [Git](#), [21](#)
 - [-u flag](#), [35](#)
 - [blob](#), [26](#)
 - [branch](#), [26](#), [130](#)
 - [branching](#), [32](#)
 - [clone](#), [35](#)
 - [commit](#), [26](#), [129](#), [133](#)
 - [commit hash](#), [130](#)
 - [data model](#), [26](#)
 - [default branch](#), [33](#)
 - [diff](#), [47](#)
 - [fetch](#), [35](#)
 - [fork](#), [35](#), [43](#), [71](#), [77](#)
 - [HEAD](#), [26](#)
 - [history](#), [30](#)
 - [HTTPS clone](#), [37](#), [43](#)
 - [install](#), [27](#)
 - [LFS](#), [26](#)
 - [merge conflict](#), [32](#)
 - [merging](#), [32](#)
 - [pull](#), [35](#), [52](#)
 - [push](#), [34](#), [35](#)
 - [remotes](#), [35](#)
 - [repository](#), [29](#)
 - [reset](#), [43](#)
 - [staging](#), [29](#)
 - [staging area](#), [29](#)
 - [status](#), [33](#)
 - [tag](#), [41](#), [59](#), [69](#), [77](#), [95](#), [111](#), [122](#), [126](#), [136](#)
 - [tags](#), [34](#)
 - [tree](#), [26](#)
 - [working directory](#), [29](#)
- [GitHub](#), [21](#), [35](#)
 - [username](#), [37](#)
- [GitHub Actions](#), [124](#), [127](#)
- [GitLab CI](#), [127](#)
- [unicorn](#), [87](#)

- HighFive, 137
- HTML, 107, 116
- HTTP, 73, 75, 78, 107, 109, 123
 - 429, 86
 - headers, 84
 - methods, 78, 79
 - status codes, 78
- httpbin, 79
- HTTPS, 107, 123
- Hugging Face, 26, 60
- hypervisor, 13
- idempotence, 44, 61, 88, 97, 122, 130
- inference, 55
- inference interface, 67
- infrastructure as code, 88
- interface contracts, 24
- isolation, 46
- issues, 35
- JavaScript, 107, 116, 118
- Jenkins, 127
- joblib, 48, 55, 60, 63
- journalctl, 87, 133
- JSON, 79
- KeePass, 18
- keyboard layout, 14
- least privilege, 16, 96
- Let's Encrypt, 123
- license, 2
- Linux, 10, 16
- lock file, 53
- LogisticRegression, 61
- migration, 104
- MNIST, 11, 55, 65, 120
- model artefact, 57
- model card, 63
- model contract, 63
- model integration, 55
- model versioning, 60, 61, 71
- MODEL_PATH, 67
- MODELCARD.md, 63
- nano, 19, 33
- Nginx, 107, 109, 112–114
- NumPy, 63
- OpenAPI, 78
- OpenSSH, 18
- operations, 126
- Oracle VirtualBox Manager, 10
- ORM, 91, 99
- os.getenv, 50
- path operation, 80
- persistence, 93
- pickle, 60
- pip, 38, 48
- pip audit, 48, 49
- pipeline, 126
- PixelWise, 11, 39, 57, 75, 93, 110
- poetry, 53
- ports, 112
- PostgreSQL, 91, 96
- postgresql-client-common, 97
- primary key, 96
- prod, 133
- project report, 9
- project structure, 38
- Projektbericht, 9
- provisioning, 44, 45, 61, 62
- proxy, 112
- proxy_pass, 114
- psql, 96, 98
- pull request, 35, 127, 130
- PuTTY, 18
- Pydantic, 73, 80
- PyPI, 48
- pyproject.toml, 53
- pytest, 127, 128, 130
- python-dotenv, 48, 50
- rate limiting, 84
- recap, 21, 55, 72, 90, 106, 124, 137
- recovery, 25, 41, 59, 77, 95, 111, 126, 136
- ReDoc, 78
- reflection, 21, 55, 72, 90, 106, 124, 137
- relational database, 93, 96
- Remote-SSH, 18
- repository layout, 45
- reproducibility, 39
- requirements.txt, 39, 48, 49
- REST, 73, 78
- reverse proxy, 107, 109, 112
- same-origin, 116
- schema, 100
 - initialisation, 101
- scikit-learn, 48, 63
- secrets, 30, 134
- security, 72, 90, 137
- semantic versioning, 34, 48
- server, 11
- service, 76
- set -euo pipefail, 44
- setup-server.sh, 39, 44, 61, 88, 97, 122, 133
- shebang, 44
- shell prompt, 47
- slowapi, 84
- snapshot, 13, 25
- SQL, 91, 96
 - GROUP BY, 96
 - JOIN, 96
- SQL injection, 99
- SQLAlchemy, 91, 99
 - session, 100, 102
- SQLite, 91, 96, 97
- SSH, 10, 18, 124
 - absolute path, 27
 - config, 27, 42
 - file permissions, 27
 - IdentitiesOnly, 27
 - key, 27, 42, 43
 - keygen, 27
 - known_hosts, 42
 - login, 37
 - private key, 28
 - public key, 28
- ssh-agent, 18, 42
- stateless, 78
- sudo, 27, 132
- Swagger UI, 78, 83, 86
- systemctl, 19, 87, 89
- systemd, 19, 73, 75, 87, 124, 127, 132, 136
 - EnvironmentFile, 100
 - timer, 130, 131
 - Type, 131
 - unit, 131
- testing, 128
 - smoke, 128
- tig, 34
- TLS, 107, 123
- touch typing, 17
- train.py, 60
- train/serve skew, 57, 63
- transitive dependencies, 53

try_files, 114
type-1 hypervisor, 13
type-2 hypervisor, 13

Ubuntu, 16
 terminal, 27
Unix philosophy, 16
URL design, 116
uv, 53

uvicorn, 73, 75, 83, 87, 109

validation, 58
venv, 46
verification, 58, 66, 114, 115, 121
version control, 21, 23
Virtual Machines, 10, 13
virtualenv, 38, 46
VS Code, 18

webhook, 134
Weights & Biases, 26
whoami, 16
wild bees, 139

X-API-Key, 85, 116, 121